



Indexing 3

Instructor: Beom Heyn Kim

beomheyunkim@hanyang.ac.kr

Department of Computer Science



Overview

- B+-Tree Index Files II
- Assignments



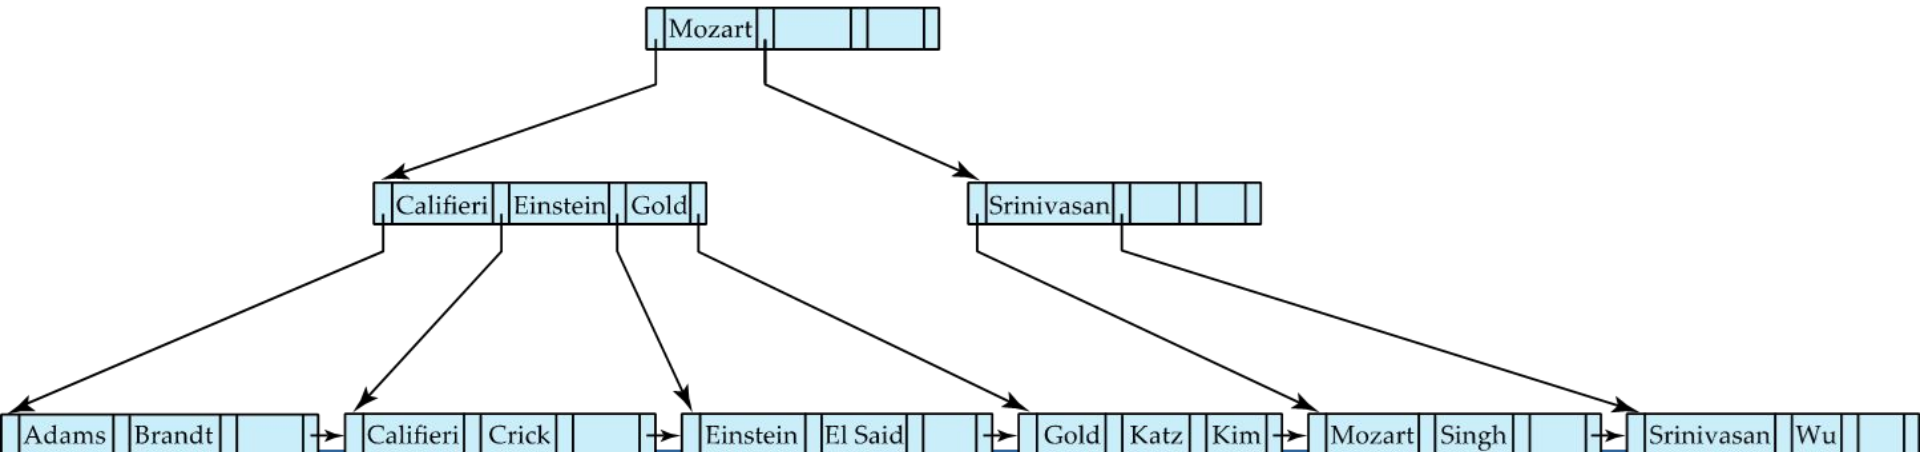
Queries on B+-Trees

function *find*(*v*)

1. *C* = *root*
2. **while** (*C* is not a leaf node)
 1. Let *i* be least number s.t. $V \leq K_i$
 2. **if** there is no such number *i* **then**
 3. Set *C* = *last non-null pointer in C*
 4. **else if** ($v = C.K_i$) Set *C* = P_{i+1}
 5. **else set** *C* = $C.P_i$
3. **if** for some *i*, $K_i = V$ **then** return $C.P_i$
4. **else** return null /* no record with search-key value *v* exists. */

Try:

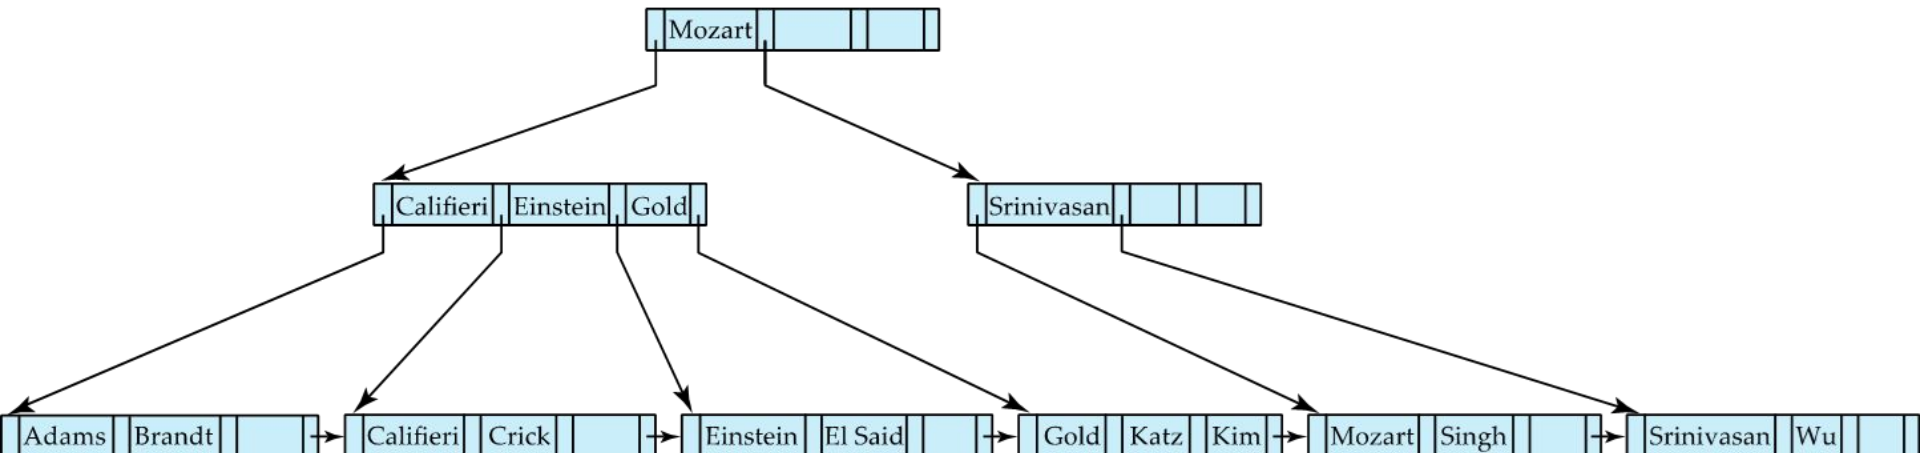
1. find(Crick)
2. find(Wu)
3. find(Lee)





Queries on B+-Trees (Cont.)

- **Range queries** find all records with search key values in a given range
 - See book for details of **function** *findRange(lb, ub)* which returns set of all such records
 - Real implementations usually provide an iterator interface to fetch matching records one at a time, using a *next()* function





Queries on B+-Trees (Cont.)

- If there are K search-key values in the file, the height of the tree is no more than $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$.
- A node is generally the same size as a disk block, typically 4 kilobytes
 - and n is typically around 100 (40 bytes per index entry).
- With 1 million search key values and $n = 100$
 - at most $\log_{50}(1,000,000) = 4$ nodes are accessed in a lookup traversal from root to leaf.
- Contrast this with a balanced binary tree with 1 million search key values — around 20 nodes are accessed in a lookup
 - above difference is significant since every node access may need a disk I/O, costing around 20 milliseconds



Non-Unique Keys

- If a search key a_i is not unique, create instead an index on a composite key (a_i, A_p) , which is unique
 - A_p could be a primary key, record ID, or any other attribute that guarantees uniqueness
- Search for $a_i = v$ can be implemented by a range search on composite key, with range $(v, -\infty)$ to $(v, +\infty)$
 - $-\infty$ denotes the smallest and $+\infty$ denotes the largest possible values of A_p
- But more I/O operations are needed to fetch the actual records
 - If the index is clustering, all accesses are sequential
 - If the index is non-clustering, each record access may need an I/O operation



Updates on B+-Trees

- Insertion and deletion are more complicated than lookup
- Nodes never become too large or too small.
 - Insertion may **split** a node
 - Deletion may **coalesce** nodes (if a node becomes too small (fewer than $\lceil n/2 \rceil$ pointers))
- Overview of Insertion
 - Using the same technique as for lookup from find, find the leaf node where the inserted value would appear
 - Insert an entry in the leaf node, maintaining search keys in order
- Overview of Deletion
 - Using the same technique as for lookup from find, find the entry/entries to be deleted
 - Remove the entry in the leaf node
 - Entries to the right of the deleted entry are shifted left by one position (no gap allowed between entries)



Updates on B+-Trees: Insertion


- Assume record already added to the file. Let
 - pr be pointer to the record, and let
 - v be the search key value of the record
- Find the leaf node in which the search-key value would appear
 - If there is room in the leaf node, insert (v, pr) pair in the leaf node
 - Otherwise, split the node (along with the new (v, pr) entry) as discussed in the next slide, and propagate updates to parent nodes.



Updates on B+-Trees: Insertion (Cont.)

- Splitting a leaf node:
 - take the n (search-key value, pointer) pairs (including the one being inserted) in sorted order.
 - Place the first $\lceil n/2 \rceil$ in the original node, and the rest in a new node.
 - let the new node be p , and let k be the least key value in p . Insert (k,p) in the parent of the node being split
 - If the parent is full, split it and **propagate** the split further up.
- Splitting of nodes proceeds upwards till a node that is not full is found.
 - In the worst case the root node may be split increasing the height of the tree by 1.

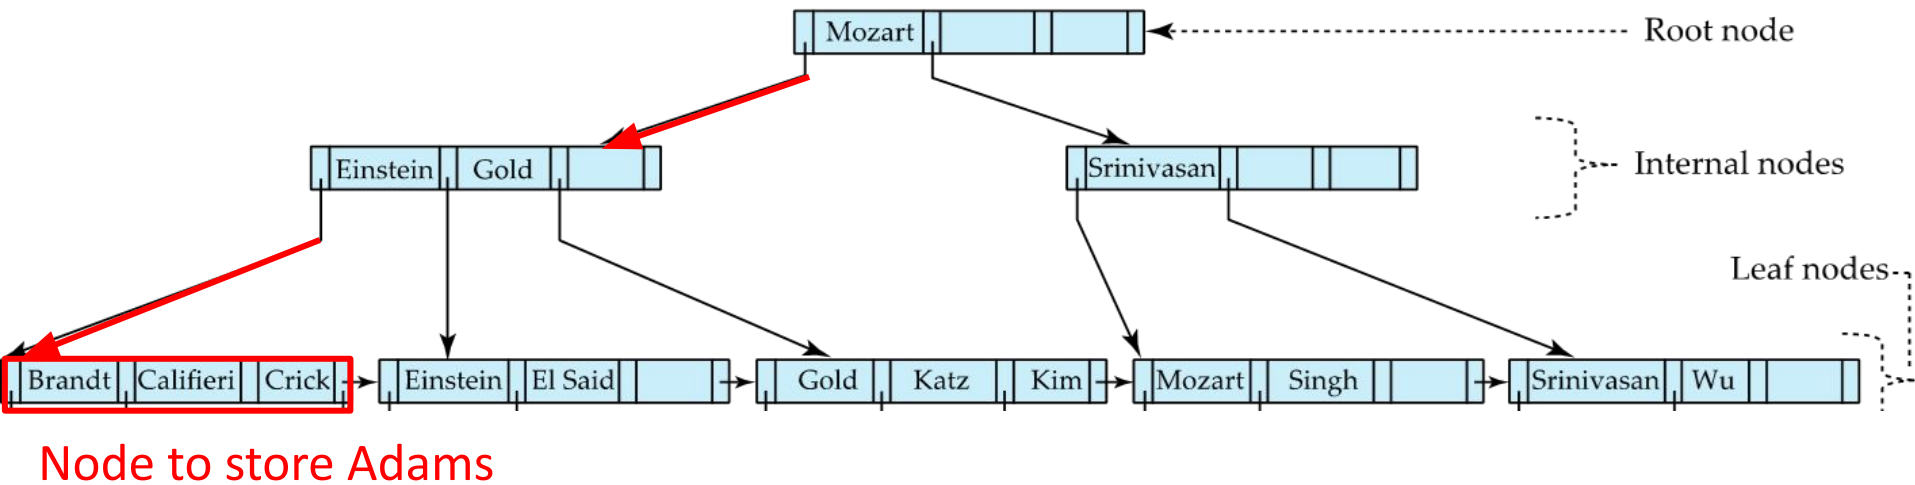
Splitting for
inserting D:





B+-Tree Insertion


B⁺-Tree before insertion of “Adams”






B+-Tree Insertion

B⁺-Tree before insertion of “Adams”

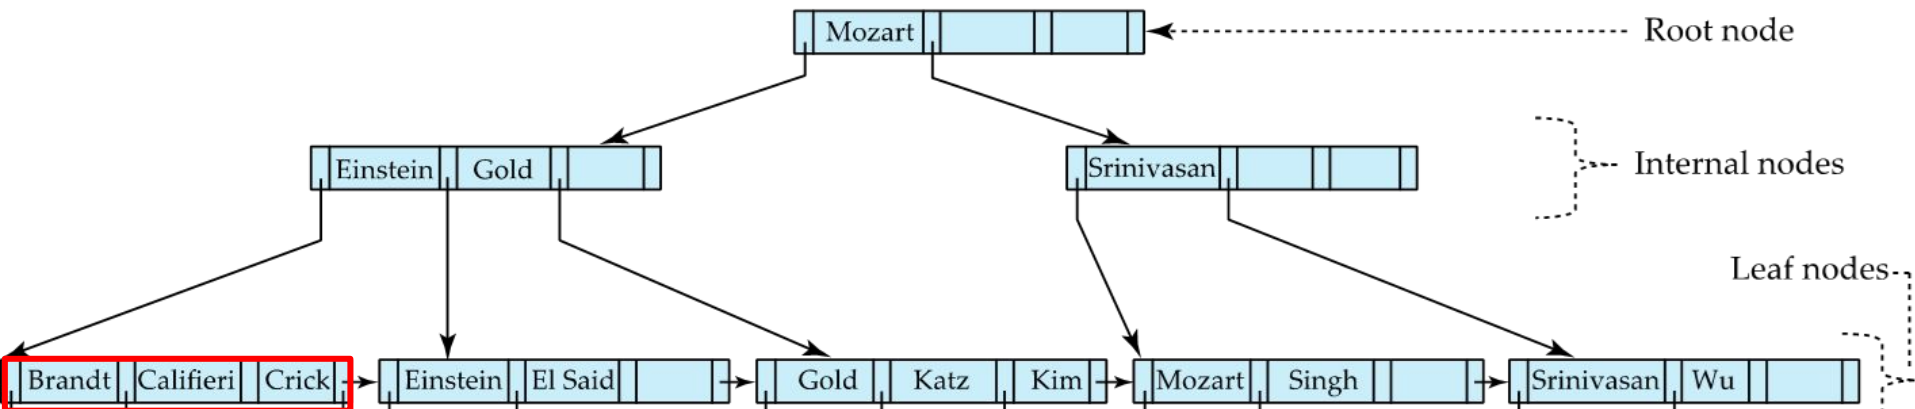


Need to split the node

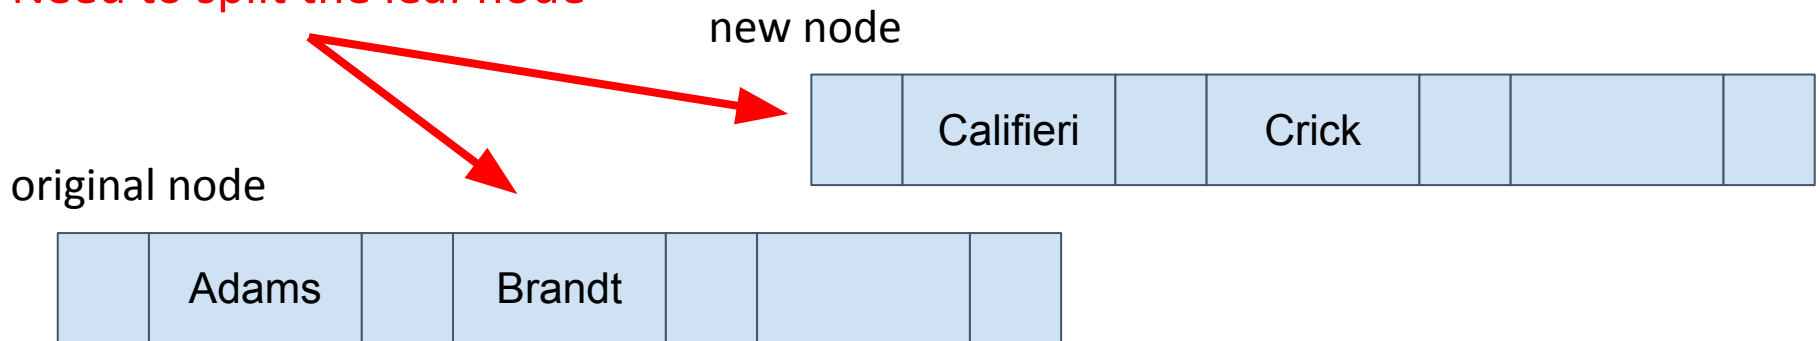


B+-Tree Insertion

B⁺-Tree before insertion of “Adams”



Need to split the leaf node

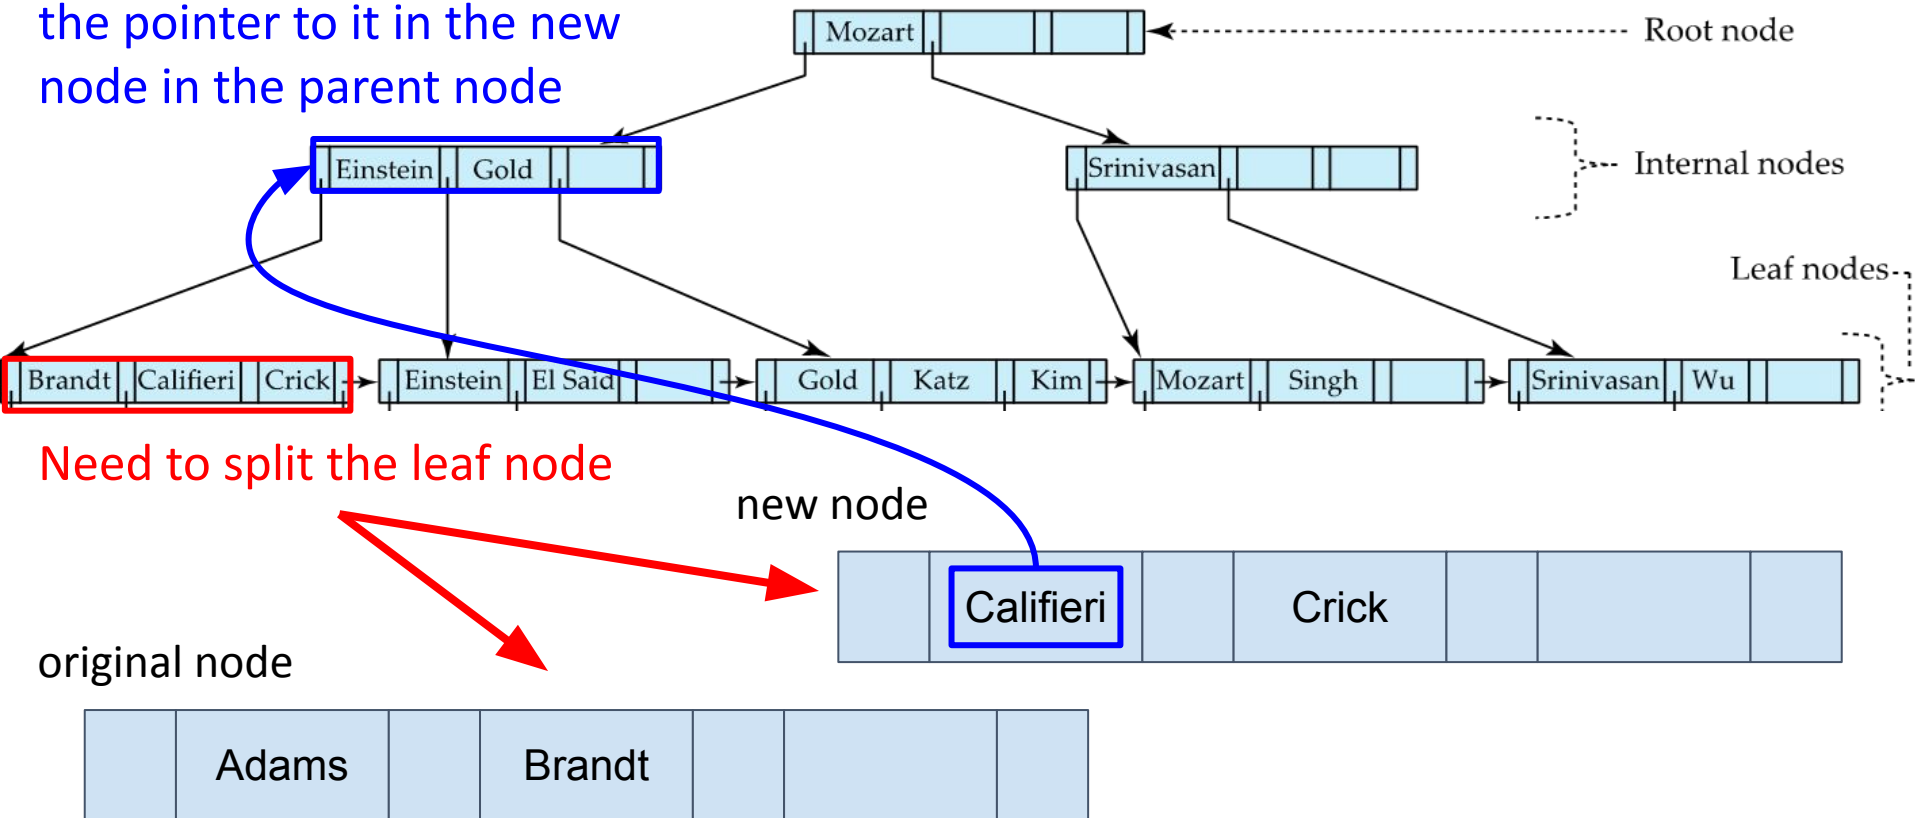




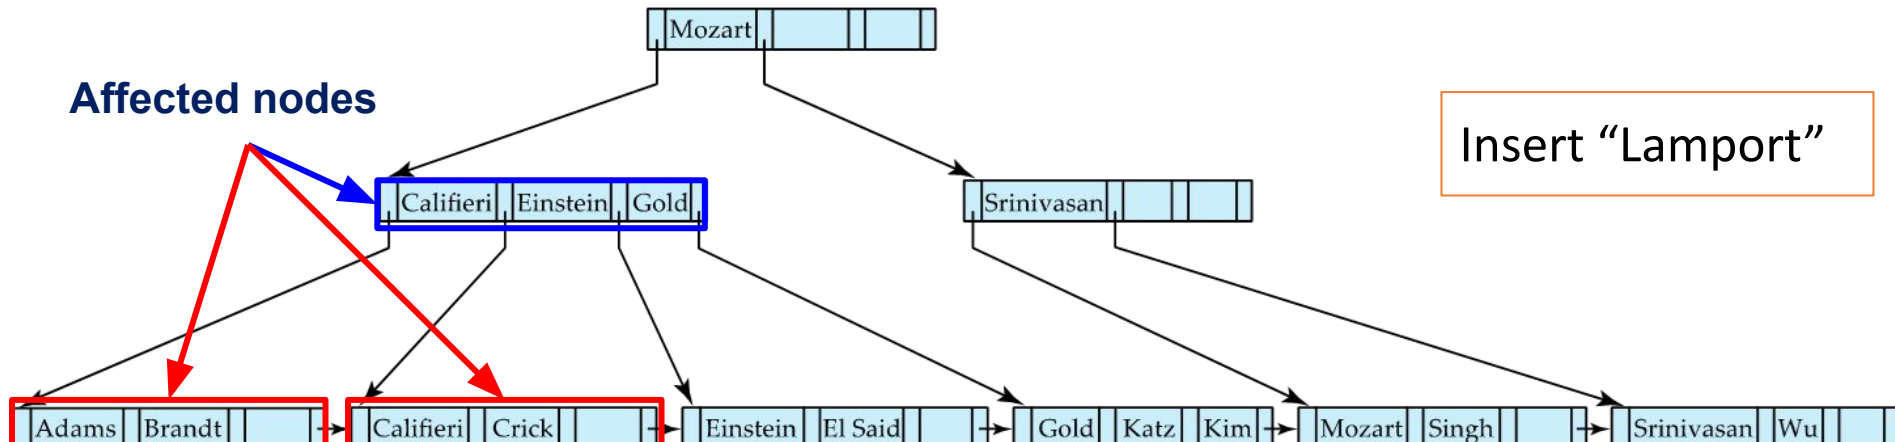
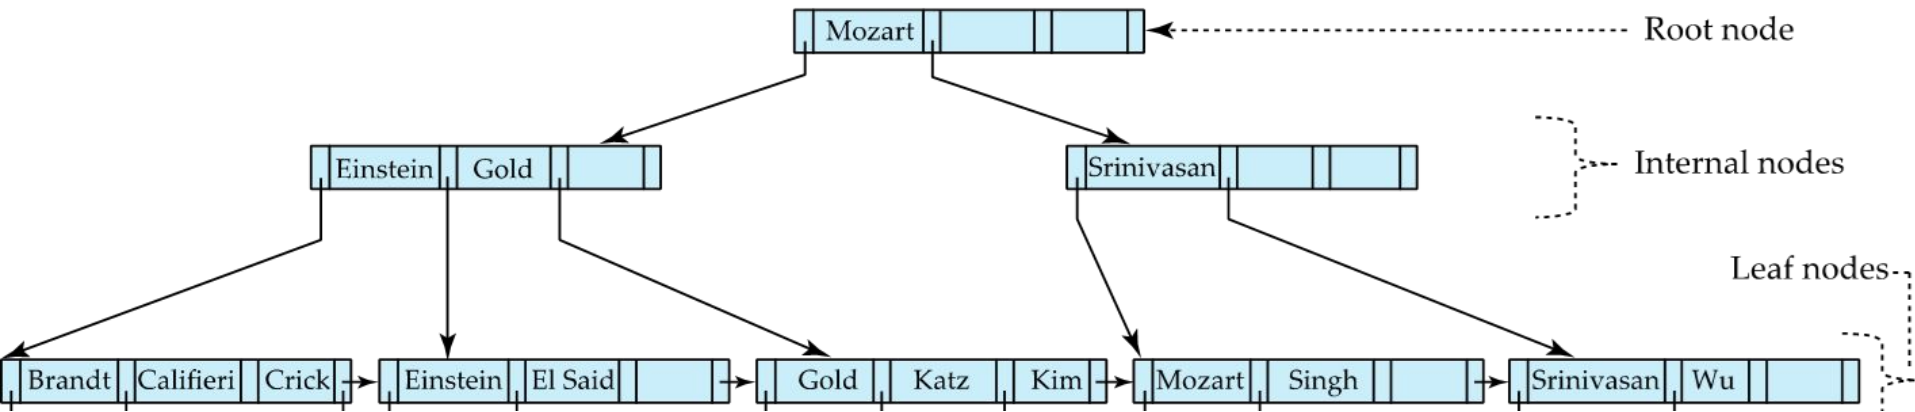
B+-Tree Insertion

B⁺-Tree before insertion of “Adams”

Insert the least search-key and
the pointer to it in the new
node in the parent node



B+-Tree Insertion



B⁺-Tree before and after insertion of "Adams"



Need to split
the non-leaf
node (next
slide)

parent node

Need to split the non-leaf node (next slide)


original node

new node

Node to store Lamports

Insertion in B+-Trees (Cont.)

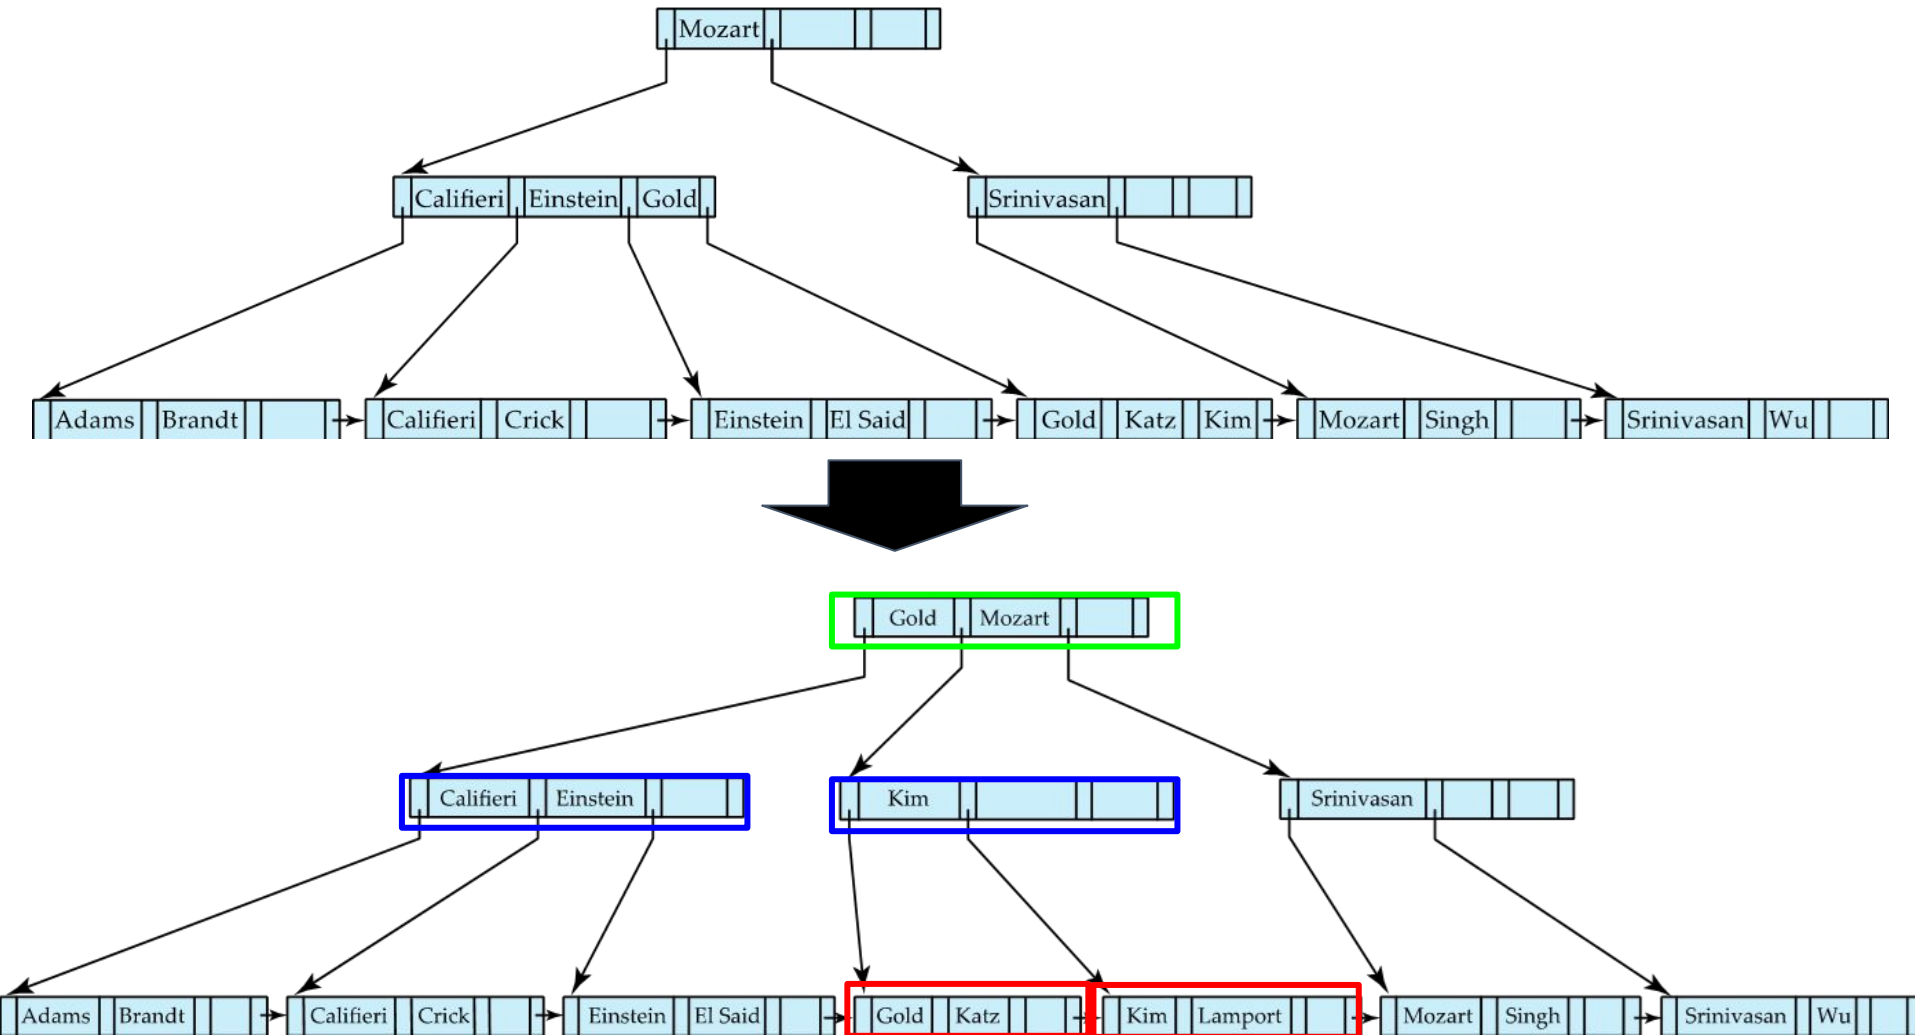
- Splitting a non-leaf node: when inserting (k,p) into an already full internal node N
 - Copy N to an in-memory area M with space for $n+1$ pointers and n keys
 - Insert (k,p) into M
 - Copy $P_1, K_1, \dots, K_{\lceil n/2 \rceil - 1}, P_{\lceil n/2 \rceil}$ from M back into node N
 - Copy $P_{\lceil n/2 \rceil + 1}, K_{\lceil n/2 \rceil + 1}, \dots, K_n, P_{n+1}$ from M into newly allocated node N'
 - Insert $(K_{\lceil n/2 \rceil}, N')$ into parent N
- Example



- **Read pseudocode in book!**



B+-Tree Insertion




B⁺-Tree before and after insertion of "Lampport"

Updates on B+-Trees: Deletion

- Assume record already deleted from file. Let V be the search key value of the record, and Pr be the pointer to the record.
 - Remove (Pr, V) from the leaf node
 - If the node has too few entries due to the removal, and the entries in the node and a sibling fit into a single node, then **merge siblings**:
 - Insert all the search-key values in the two nodes into a single node (the one on the left), and delete the other node.
 - Delete the pair (K_{i-1}, P_i) , where P_i is the pointer to the deleted node, from its parent, recursively using the above procedure.

(Leaf Node Example)

Merging siblings (Coalescing nodes) for deleting D:



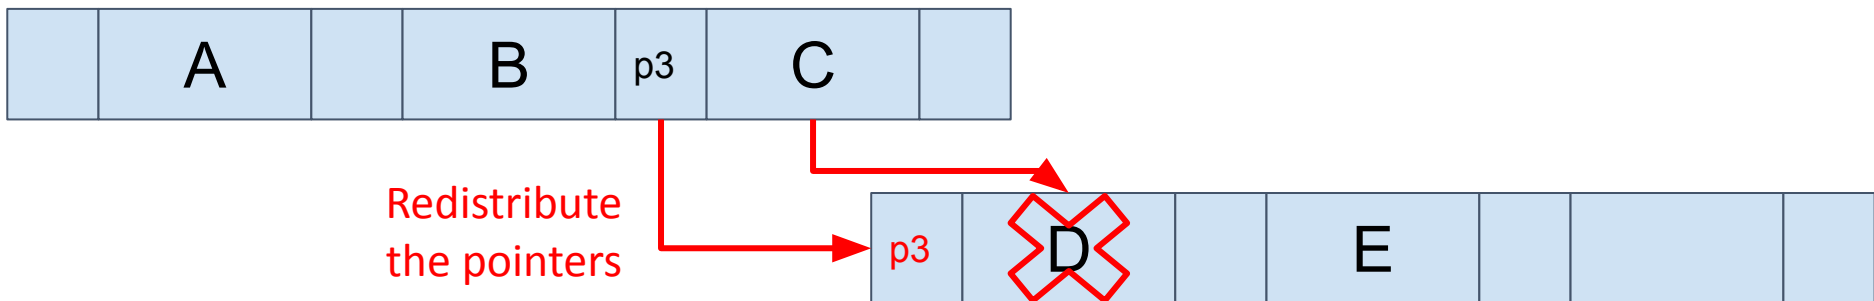


Updates on B+-Trees: Deletion

- Otherwise, if the node has too few entries due to the removal, but the entries in the node and a sibling do not fit into a single node, then **redistribute pointers**:
 - Redistribute the pointers between the node and a sibling such that both have more than the minimum number of entries.
 - Update the corresponding search-key value in the parent of the node.
- The node deletions may cascade upwards till a node which has $\lceil n/2 \rceil$ or more pointers is found.
- If the root node has only one pointer after deletion, it is deleted and the sole child becomes the root.
- **Read pseudocode in book!**

(Leaf Node Example)

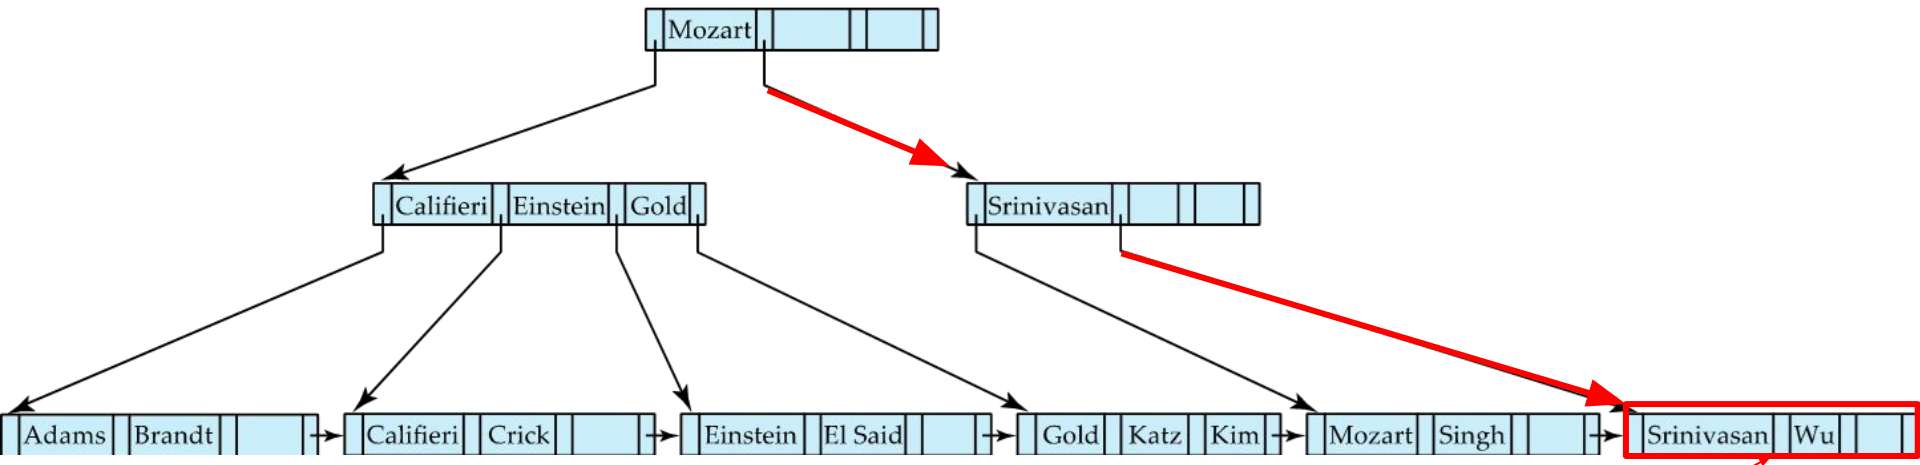
Redistributing entries
for deleting D:





Examples of B+-Tree Deletion

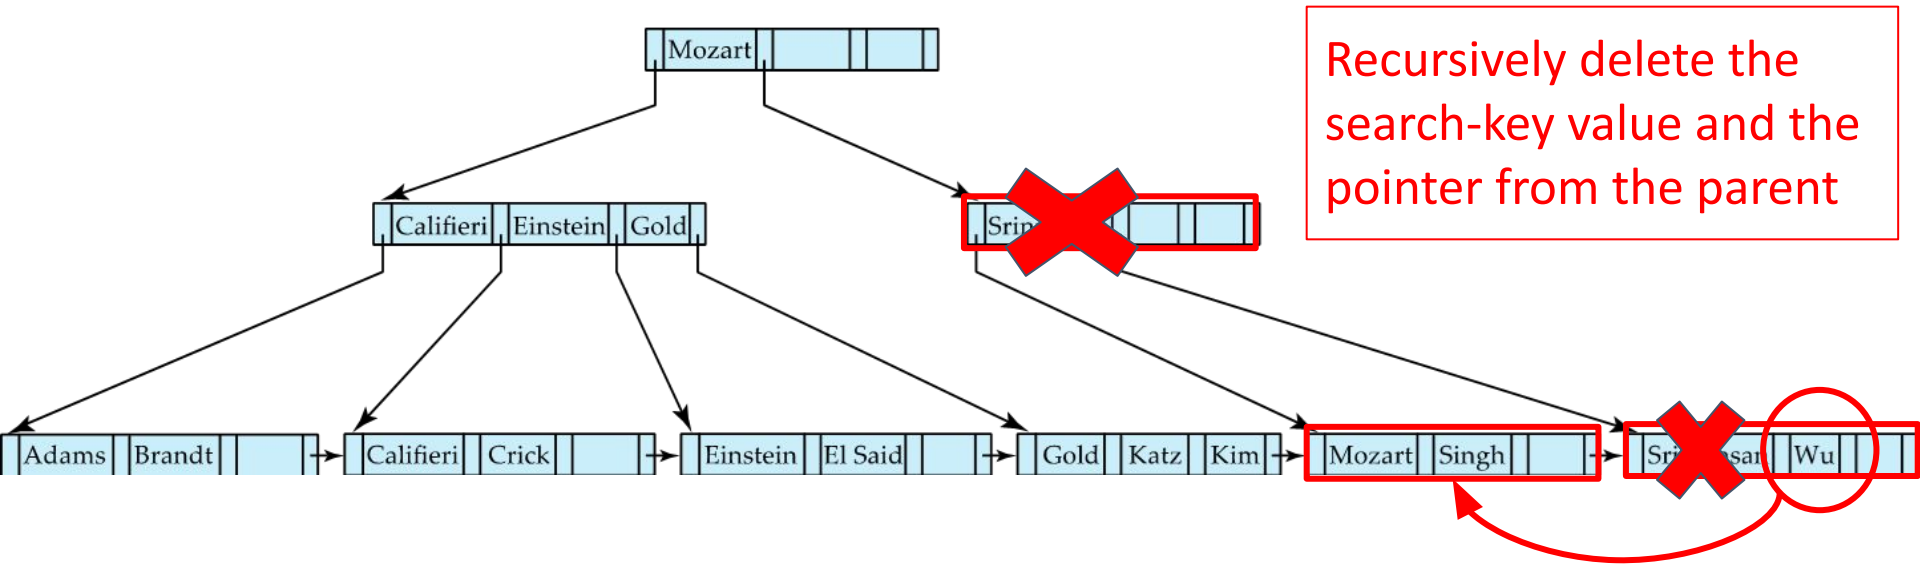
Before deleting “Srinivasan”



After deleting Srinivasan, the node becomes underfull

Examples of B+-Tree Deletion (Cont.)

Before deleting “Srinivasan”



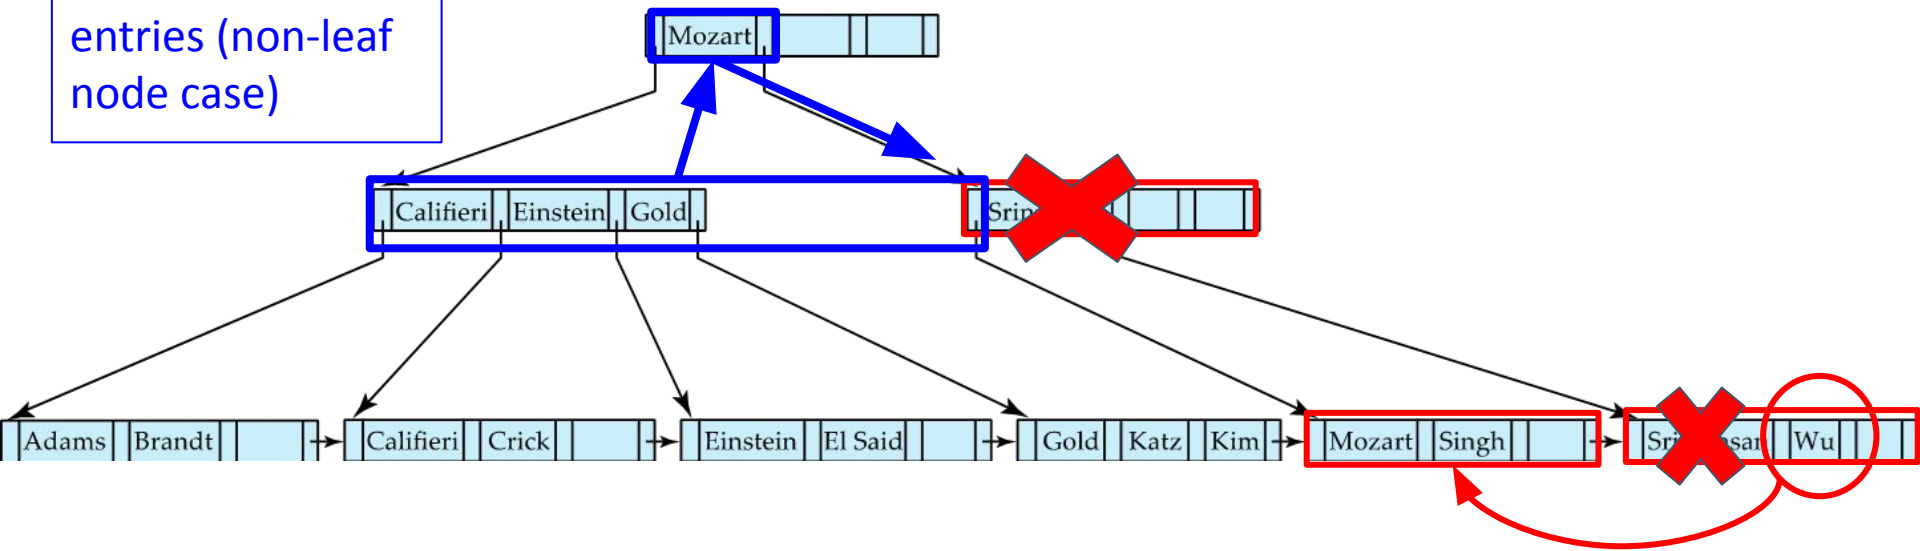
Recursively delete the search-key value and the pointer from the parent

Merge: Inserting Wu into the sibling node and delete the node used to store Srinivasan

Examples of B+-Tree Deletion (Cont.)

Before deleting “Srinivasan”

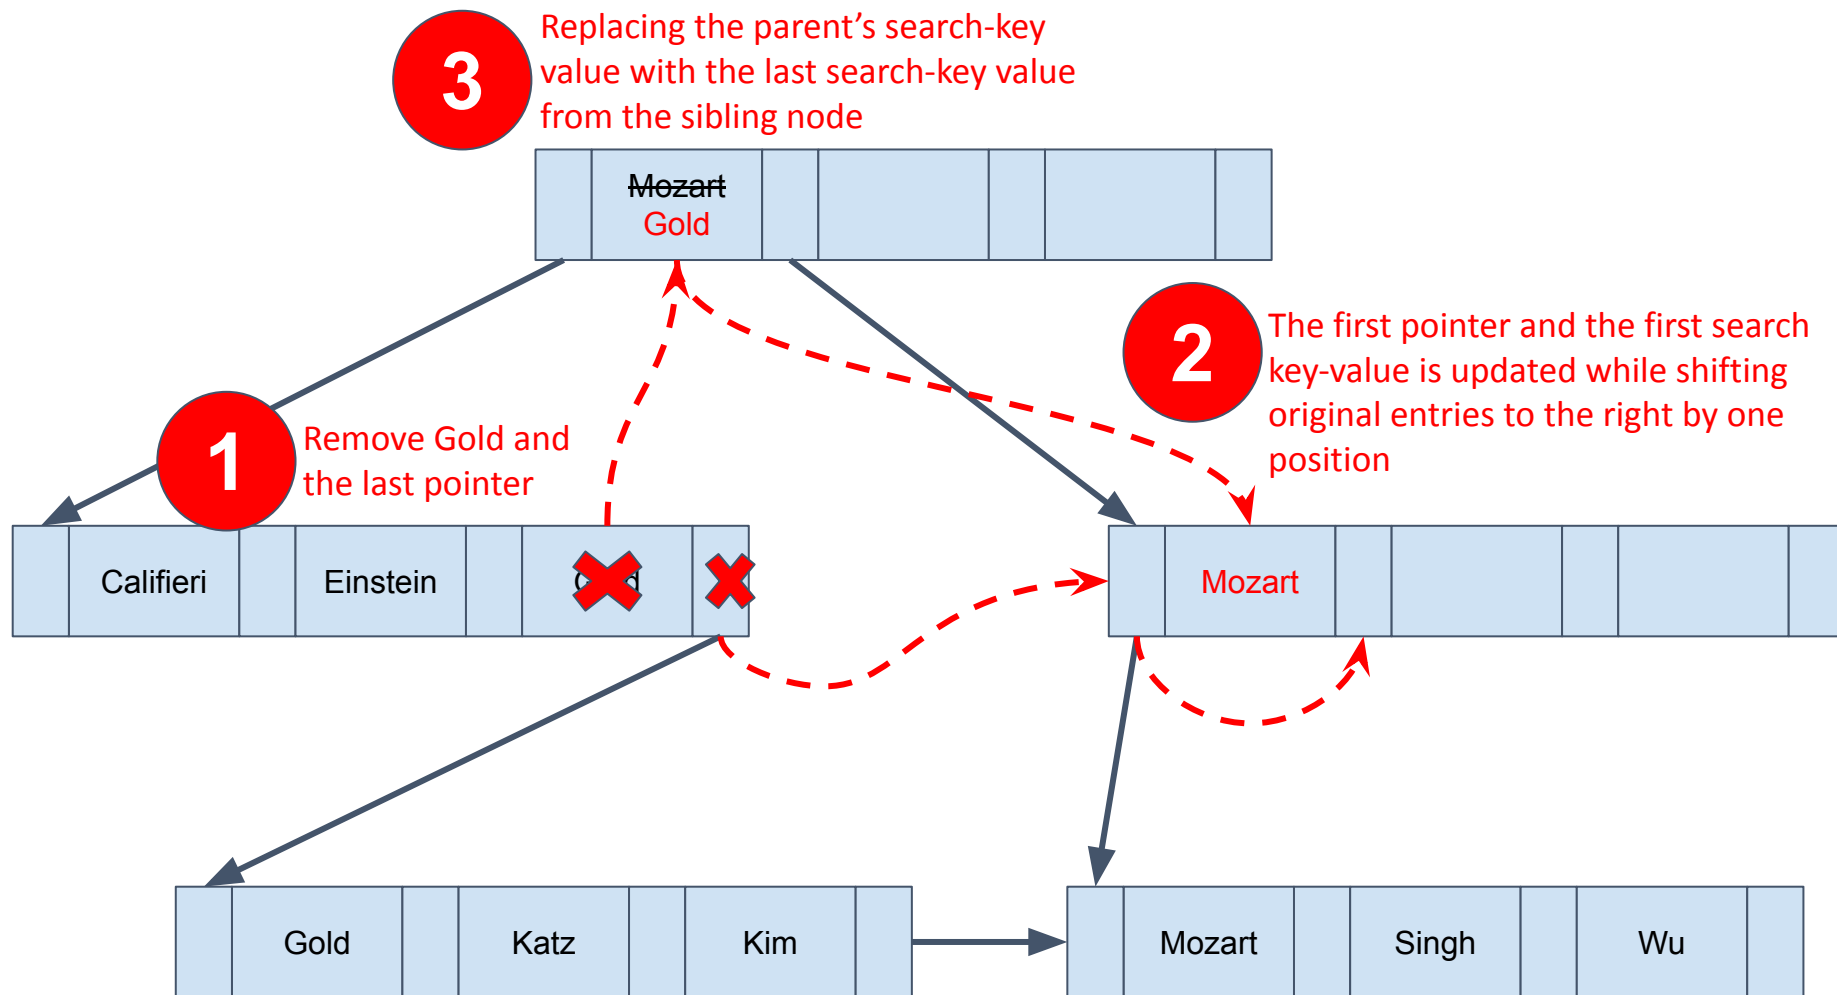
Redistribute
entries (non-leaf
node case)



- Deleting “Srinivasan” causes **merging** of underfull leaf nodes and redistribution of non-leaf nodes



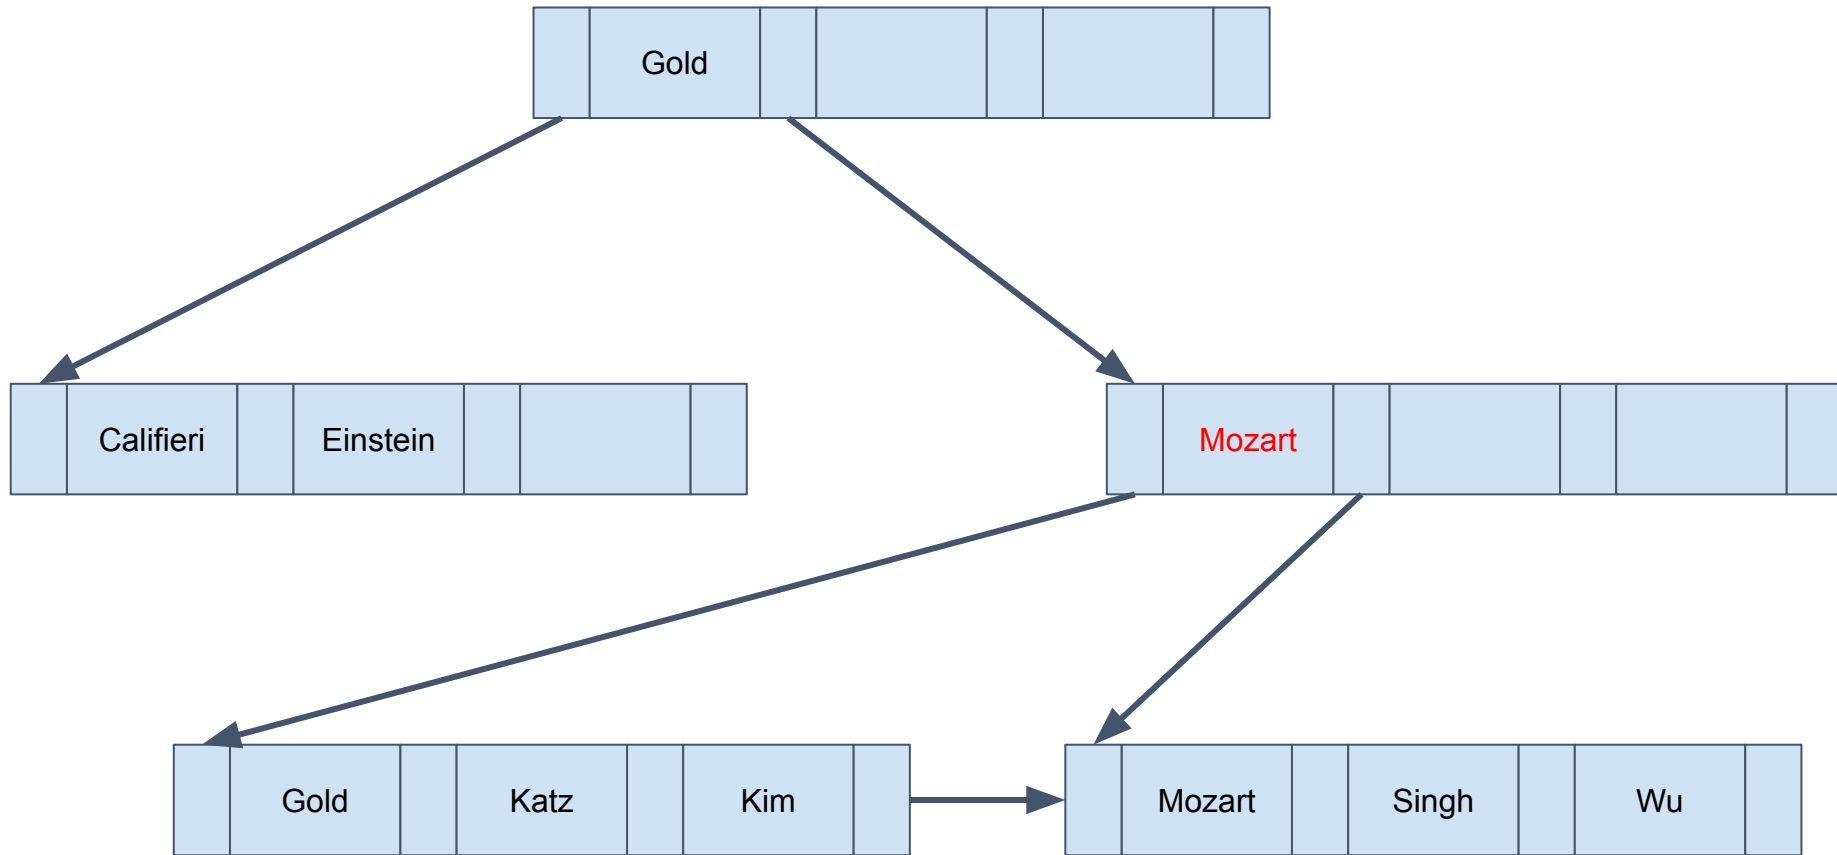
Examples of B+-Tree Deletion (Cont.)





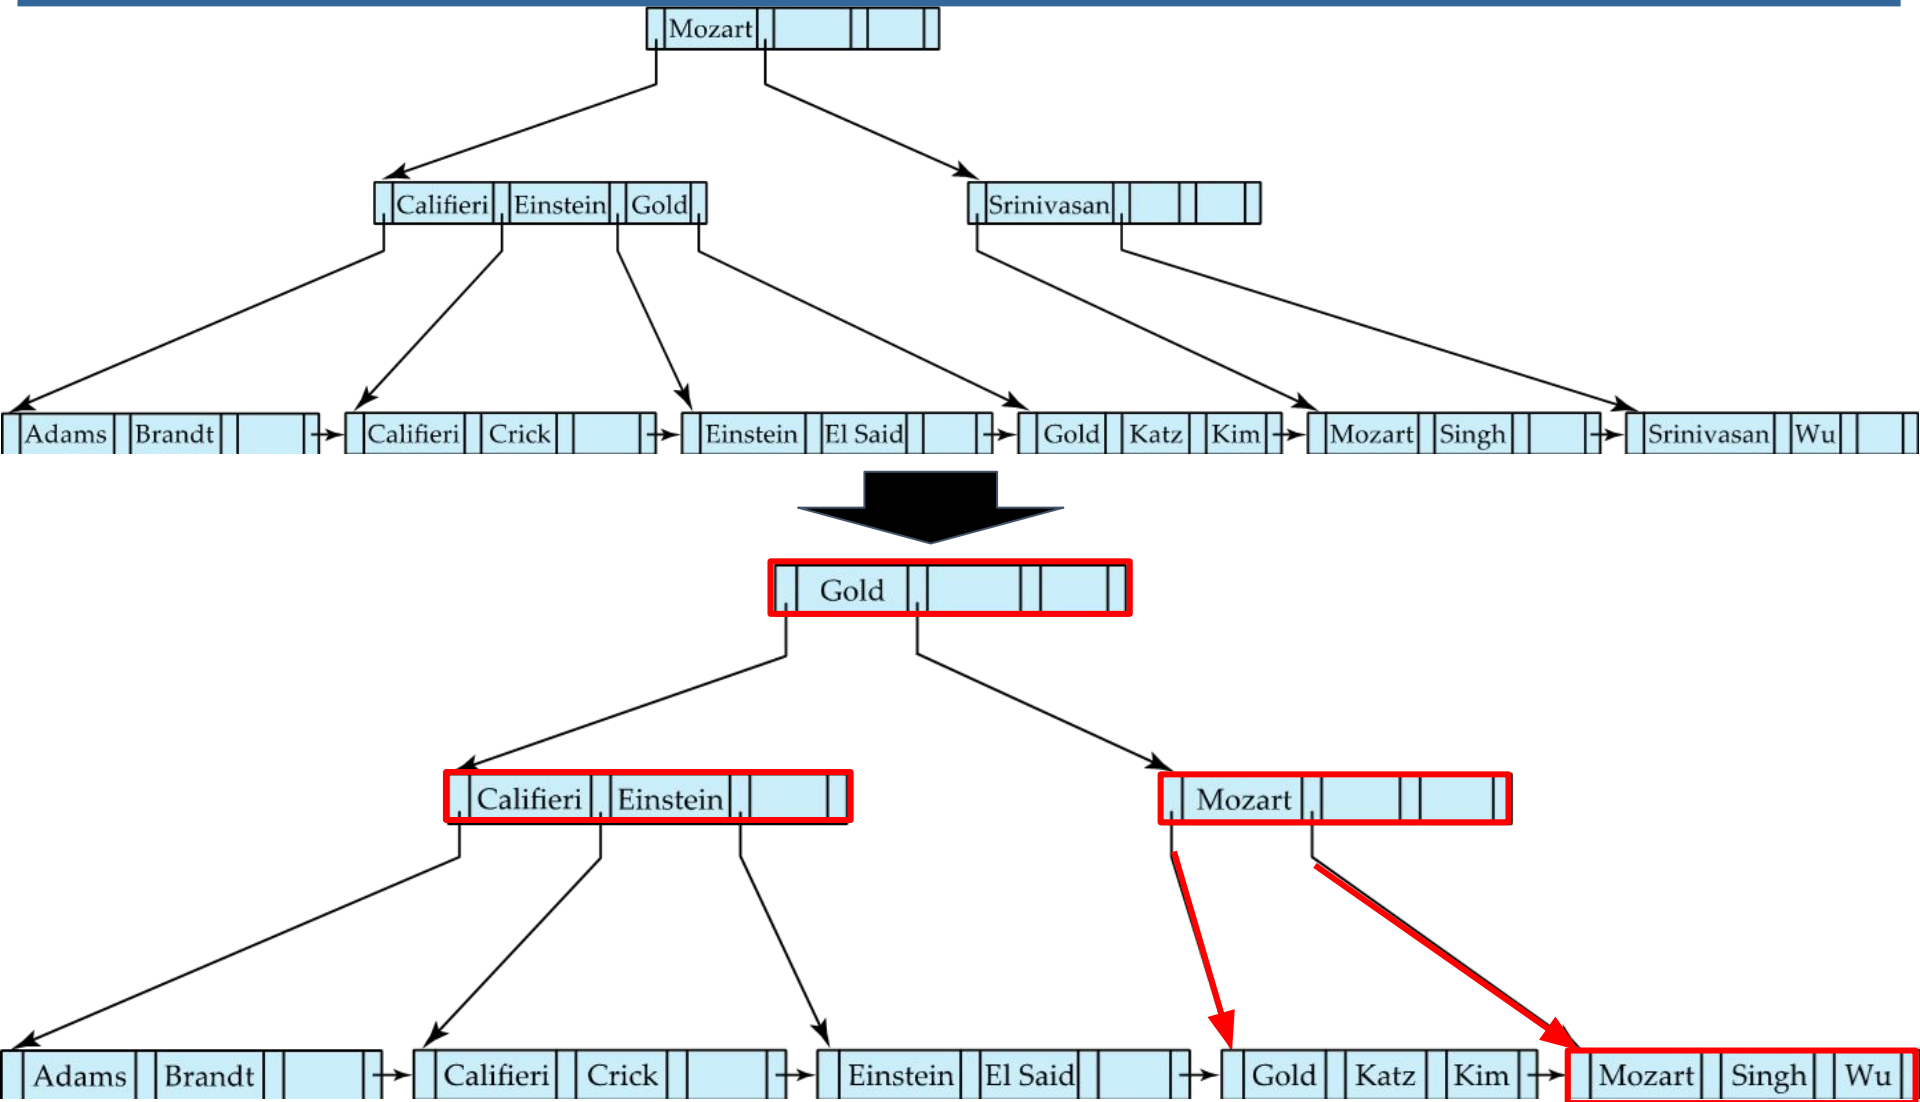
Examples of B+-Tree Deletion (Cont.)

Redistribution result after deleting “Srinivasa”





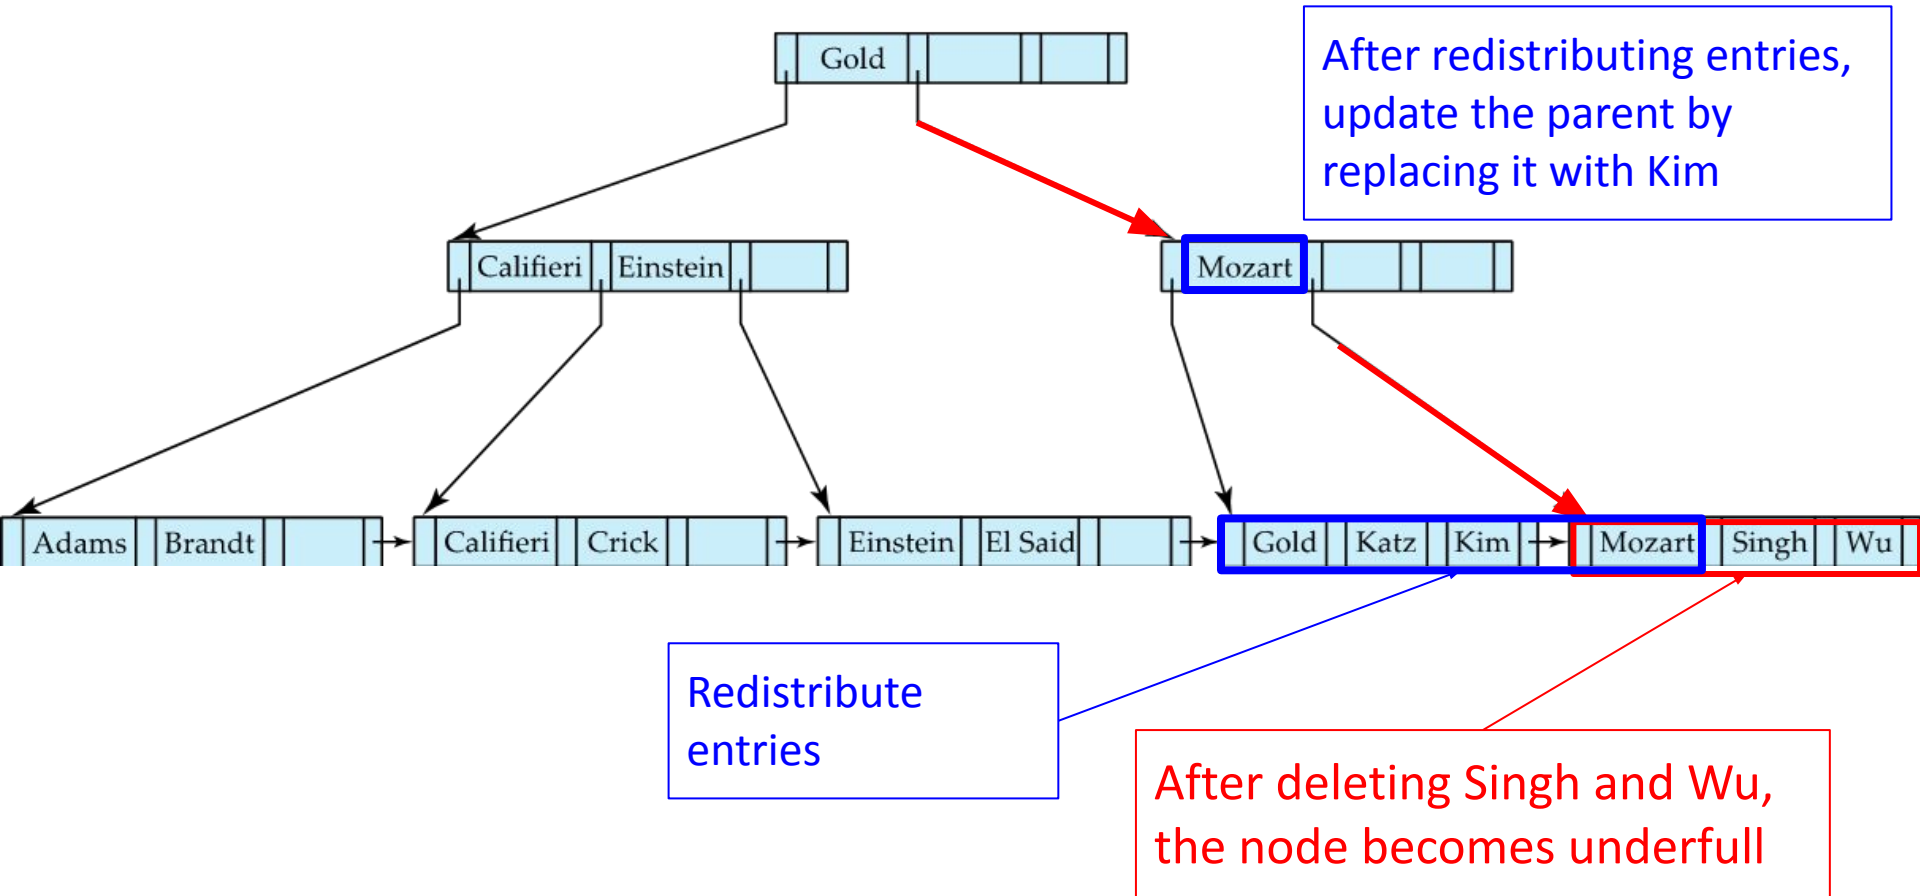
Examples of B+-Tree Deletion (Cont.)



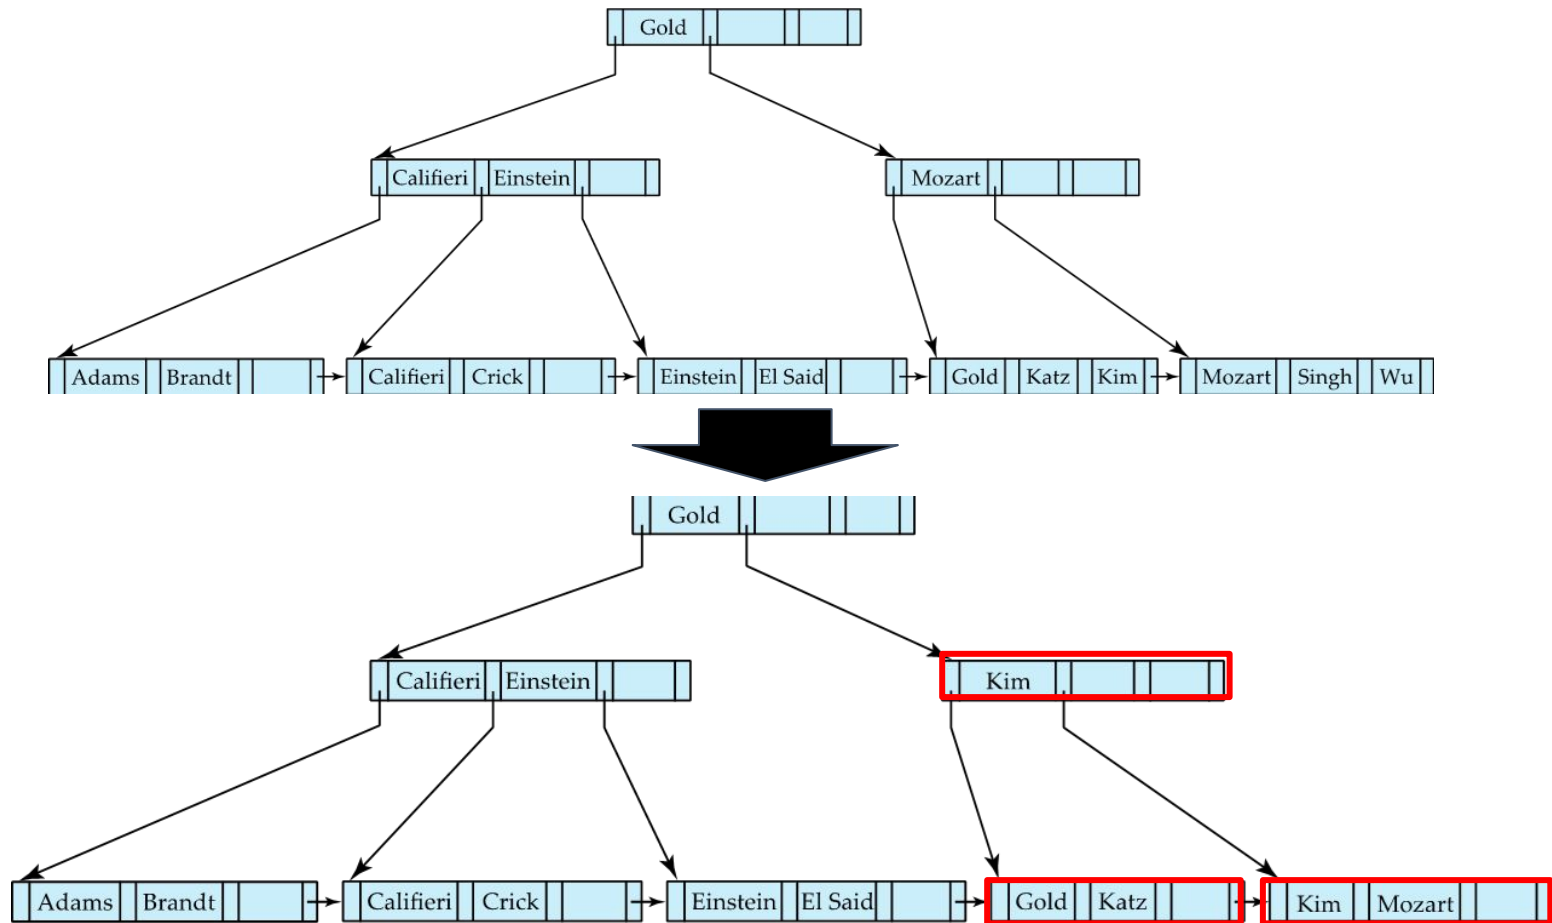
Before and after deleting “Srinivasan”

Examples of B+-Tree Deletion (Cont.)

Before deleting “Singh” and “Wu”



Examples of B+-Tree Deletion (Cont.)

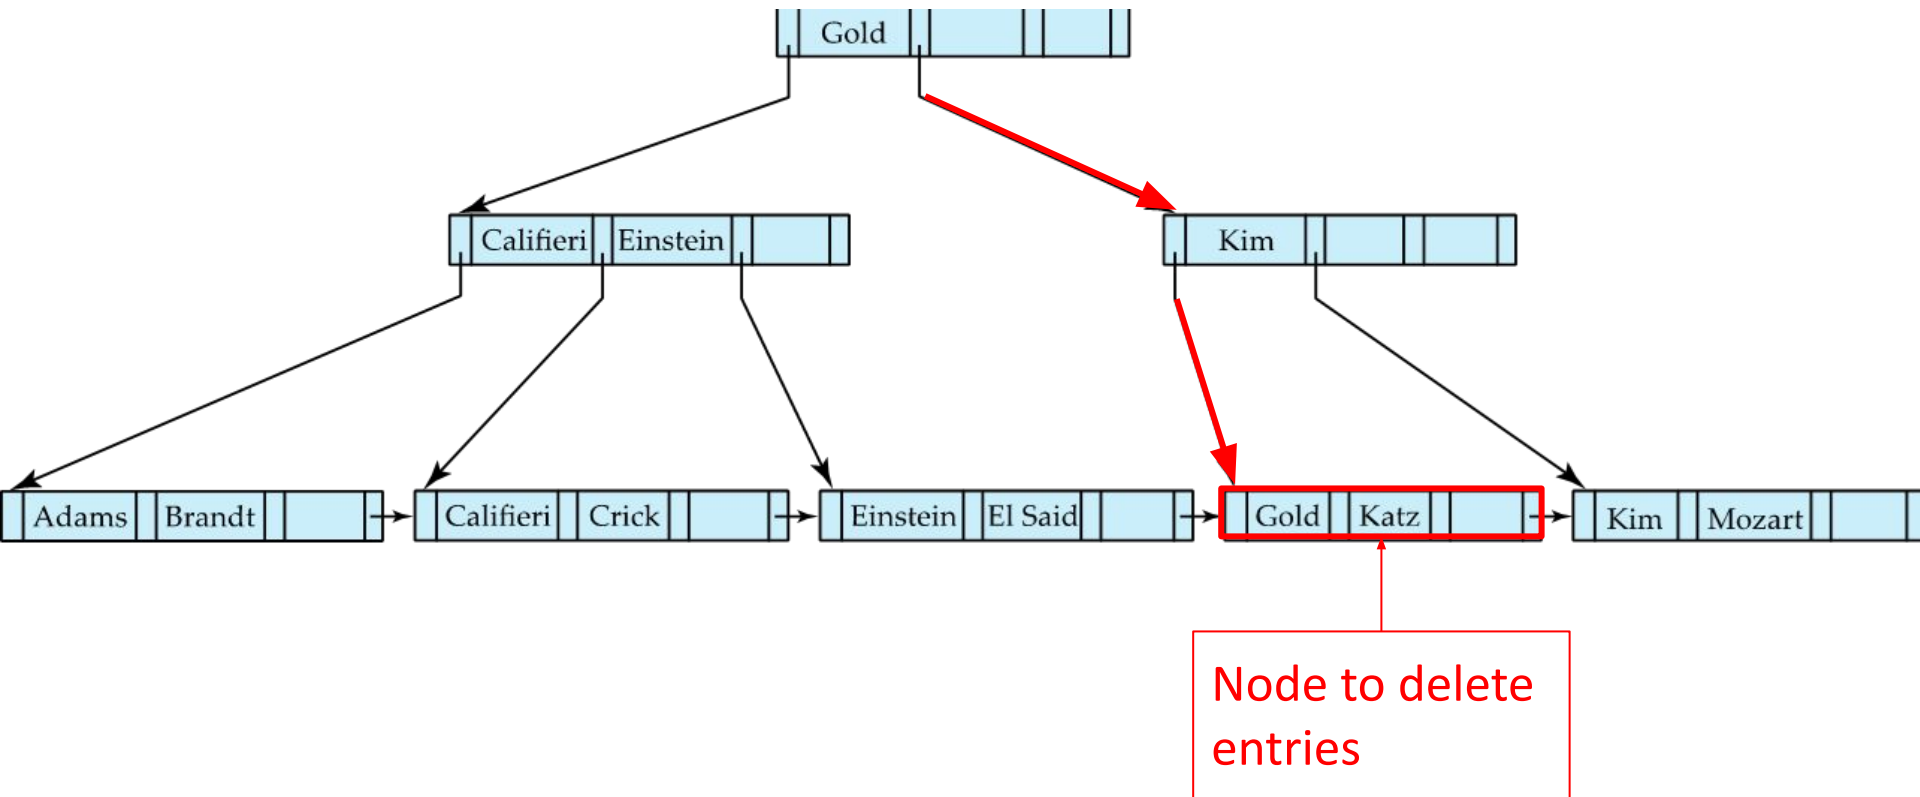


Before and after deleting “Singh” and “Wu”

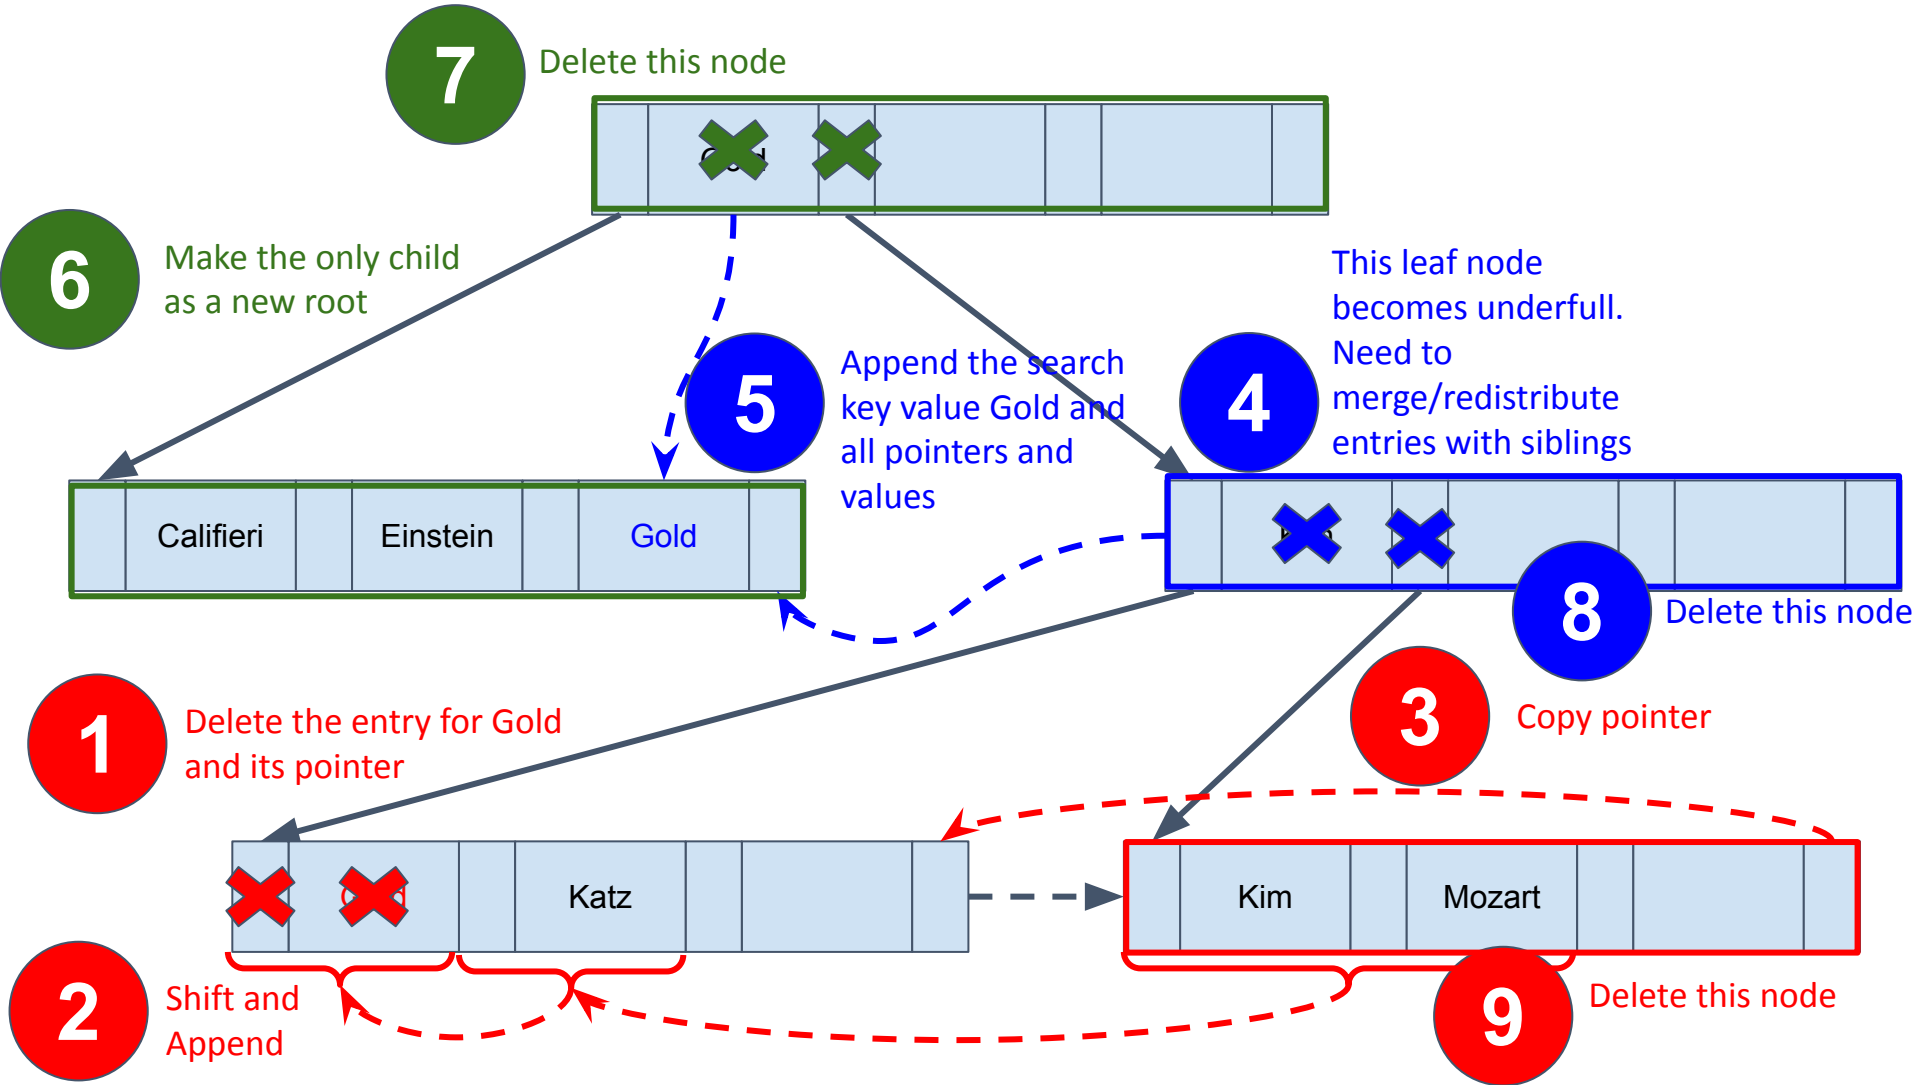
- Leaf containing Singh and Wu became underfull, and **borrowed a value Kim** from its left sibling
- Search-key value in the parent changes as a result

Example of B+-tree Deletion (Cont.)

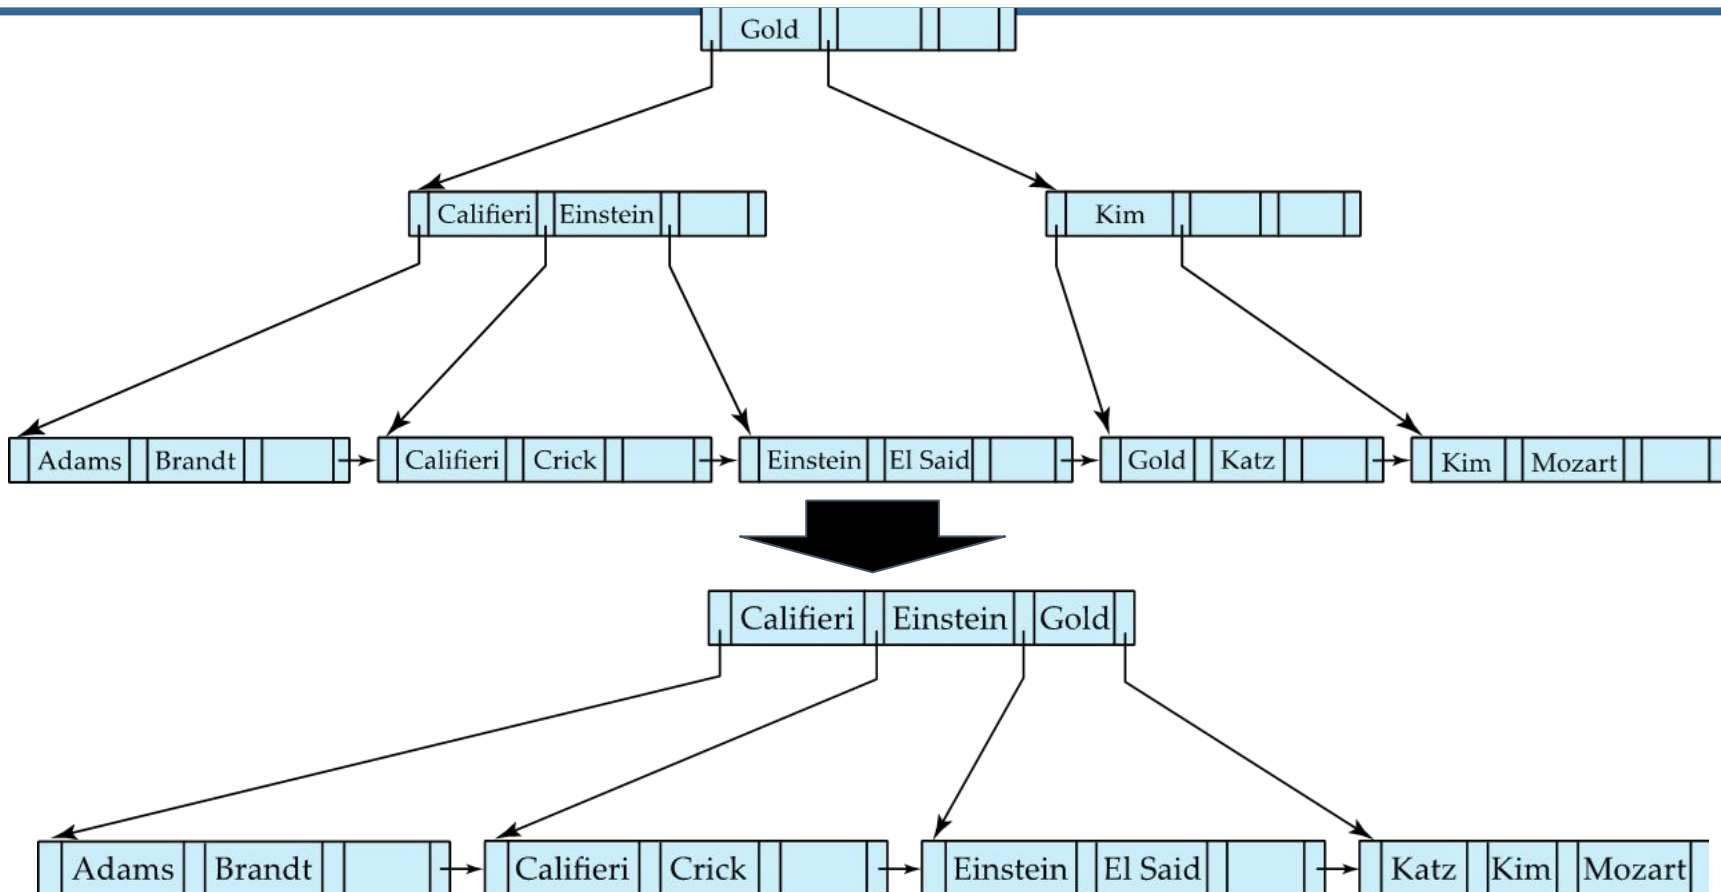
Before deletion of “Gold”



Example of B+-tree Deletion (Cont.)



Example of B+-tree Deletion (Cont.)



Before and after deletion of “Gold”

- Node with Gold and Katz became underfull, and was merged with its sibling
- Parent node becomes underfull, and is merged with its sibling
 - Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is deleted



Complexity of Updates

- Cost (in terms of number of I/O operations) of insertion and deletion of a single entry proportional to height of the tree
 - With K entries and maximum fanout of n , worst case complexity of insert/delete of an entry is $O(\log_{\lceil n/2 \rceil}(K))$
- In practice, number of I/O operations is less:
 - Internal nodes tend to be in buffer
 - Splits/merges are rare, most insert/delete operations only affect a leaf node
- Average node occupancy depends on insertion order
 - 2/3rds with random, $\frac{1}{2}$ with insertion in sorted order



Non-Unique Search Keys

- Alternatives to scheme described earlier
 - Store a bucket (or a list) of record pointers with a search-key value in the leaf node of the B+-Tree
 - Low space overhead – no need to store search key values multiple times
 - Variable sized bucket can cause difficulties and issues (e.g. may get larger than the node itself)
 - Store search key values once per record
 - Significantly complex to handle split and search on internal nodes – may be more than one leaf node containing the same search key value
 - High space overhead due to the repeated store of the same search-key value
- Duplicate search key introduces significant complication, so most database systems only handle unique search keys
 - Automatically add record-ides or other attributes to make nonunique search keys unique



Overview

- B+-Tree Index Files II
- Assignments



Assignments

- Reading: Ch14.3.2-14.3.5
- Practice Exercises: 14.1, 14.2, 14.3a, 14.4 (for 14.3a), 14.7

Solutions to the Practice Exercises:

<https://www.db-book.com/Practice-Exercises/index-solu.html>



The End
