



---

# Query Processing 3

**Instructor: Beom Heyn Kim**

[beomheyunkim@hanyang.ac.kr](mailto:beomheyunkim@hanyang.ac.kr)

Department of Computer Science

---



# Overview

---

- Join Operation
- Assignments



# Join Operation

- Join Operation:
  - Equi-join:  $r \bowtie_{r.A=s.B} s$
  - Natural join:  $r \bowtie s$
  - Theta join:  $r \bowtie_{\theta} s$  for an arbitrary join condition  $\theta$
- Several different algorithms to implement joins
  - Nested-loop join
  - Block nested-loop join
  - Indexed nested-loop join
  - Merge-join
  - Hash-join
- Choice based on cost estimate
- Running example:  $student \bowtie takes$ 
  - Assume:
    - Number of records of *student*: 5,000      *takes*: 10,000
    - Number of blocks of *student*: 100      *takes*: 400



# Nested-Loop Join

- To compute the theta join  $r \bowtie_{\theta} s$   
    **for each** tuple  $t_r$  **in**  $r$  **do begin**  
        **for each** tuple  $t_s$  **in**  $s$  **do begin**  
            test pair  $(t_r, t_s)$  to see if they satisfy the join condition  $\theta$   
            if they do, add  $t_r \cdot t_s$  to the result.  
        **end**  
    **end**
- $r$  is called the **outer relation** and  $s$  the **inner relation** of the join.
- Requires no indices and can be used with any kind of join condition.
- Expensive since it examines every pair of tuples in the two relations.

$t_r \cdot t_s$  denotes the tuple constructed by concatenating the attribute values of tuples  $t_r$  and  $t_s$



# Nested-Loop Join (Cont.)

- In the worst case, if there is enough memory only to hold one block of each relation, the estimated cost is  
 $n_r * b_s + b_r$  block transfers, plus  $n_r + b_r$  seeks
- If a relation fits entirely in memory, use that as the inner relation.
  - Reduces cost to  $b_r + b_s$  block transfers and 2 seeks
- Assuming the worst case memory availability, cost estimate is
  - with *student* as outer relation:
    - $5000 * 400 + 100 = 2,000,100$  block transfers,
    - $5000 + 100 = 5100$  seeks
  - with *takes* as the outer relation
    - $10000 * 100 + 400 = 1,000,400$  block transfers and 10,400 seeks
- If a relation (e.g. *student*) fits entirely in memory, the cost estimate will be 500 block transfers.
- Block nested-loops algorithm (next slide) is preferable.

$n_r$ : Number of tuples in relation  $r$



# Block Nested-Loop Join

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

```
for each block  $B_r$  of  $r$  do begin  
    for each block  $B_s$  of  $s$  do begin  
        for each tuple  $t_r$  in  $B_r$  do begin  
            for each tuple  $t_s$  in  $B_s$  do begin  
                Check if  $(t_r, t_s)$  satisfy the join condition  
                if they do, add  $t_r \cdot t_s$  to the result.  
            end  
        end  
    end  
end
```

- Worst case estimate (neither relation fits in memory):  $b_r * b_s + b_r$  block transfers and  $2 * b_r$  seeks
  - Each block in the inner relation  $s$  is read once for each block in the outer relation
- Best case (inner relation fits in memory):  $b_r + b_s$  block transfers and 2 seeks.
- A few more improvements are discussed in book
  - One of them is using index on inner relation if available (next slide)



# Indexed Nested-Loop Join

- Index lookups can replace file scans if
  - join is an equi-join or natural join and
  - an index is available on the inner relation's join attribute
    - Can construct an index just to compute a join.
- For each tuple  $t_r$  in the outer relation  $r$ , use the index to look up tuples in  $s$  that satisfy the join condition with tuple  $t_r$ .
- Worst case: buffer has space for only one page of  $r$ , and, for each tuple in  $r$ , we perform an index lookup on  $s$ .
- Cost of the join:  $b_r (t_T + t_S) + n_r * c$ 
  - Where  $c$  is the cost of traversing index and fetching all matching  $s$  tuples for one tuple of  $r$
  - $c$  can be estimated as cost of a single selection on  $s$  using the join condition.
- If indices are available on join attributes of both  $r$  and  $s$ , use the relation with fewer tuples as the outer relation.



# Merge-Join

1. Sort both relations on their join attribute (if not already sorted on the join attributes).
2. Merge the sorted relations to join them
  - a. Join step is similar to the merge stage of the sort-merge algorithm.
  - b. Main difference is handling of duplicate values in join attribute — every pair with same value on join attribute must be matched
  - c. Merge-Join Algorithm Summary (detailed algorithm in book):

$pr \rightarrow$

$a1$	$a2$
a	3
b	1
d	8
d	13
f	7
m	5
q	6

$r$

$ps \rightarrow$

$a1$	$a3$
a	A
b	G
c	L
d	N
m	B

$s$

1. Keep moving  $ps$  down, collect all tuples from  $s$  that have the same join attribute value and put them into a set,  $S_s$
2. Keep moving  $pr$  down, while the  $pr$  currently points to a tuple in  $r$  that has the value for join attribute which is less than those of tuples in  $S_s$
3. For each tuple in  $r$  that has the join attribute value, perform concatenation of it with each tuple in  $S_s$
4. repeat from 1 until  $ps$  and  $pr$  become null





# Merge-Join (Cont.)

- Can be used only for equi-joins and natural joins
- Each block needs to be read only once (assuming all tuples for any given value of the join attributes fit in memory)

- Thus, the cost of merge join is:

$b_r + b_s$  block transfers and  $\lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil$  seeks  
and the cost of sorting if relations are unsorted.

$b_b$ : Number of buffer blocks allocated to each relation

- **hybrid merge-join**: If one relation is sorted, and the other has a secondary B<sup>+</sup>-tree index on the join attribute
  - Merge the sorted relation with the leaf entries of the B<sup>+</sup>-tree .
  - Sort the result on the addresses of the unsorted relation's tuples
  - Scan the unsorted relation in physical address order and merge with previous result, to replace addresses by the actual tuples
    - Sequential scan is more efficient than random lookup



# Overview

---

- Join Operation
- Assignments



# Assignments

---

- Reading: Ch15.5.1, 15.5.2, 15.5.3, 15.5.4
- Practice Exercises: n/a

Solutions to the Practice Exercises:

<https://www.db-book.com/Practice-Exercises/index-solu.html>



---

**The End**

---