



Transactions 3

Instructor: Beom Heyn Kim

beomheyunkim@hanyang.ac.kr

Department of Computer Science



Overview

- Transaction Isolation and Atomicity
- Transaction Isolation Levels
- Implementation of Isolation Levels
- Transactions as SQL Statements
- Assignments



Recoverable Schedules

- Transactions may fail in the middle of execution.
- Consider the following schedule:

T_6	T_7
read(A)	
write(A)	
	read(A)
	commit
read(B)	

T_7 is dependent on T_6 because read(A) saw write(A)

Partial schedule: no commit or abort for T_6

- This is an example of *nonrecoverable* schedule.
 - T_6 fails before it commits
 - Because T_6 fails, it should be aborted.
 - However, T_7 already saw T_6 committed, so T_7 cannot abort.
 - Thus, it is not possible to recover from the failure.
- **Recoverable schedule:** if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .



Cascading Rollbacks

- In some cases, a single transaction failure leads to a series of transaction rollbacks (**cascading rollback**)
- Consider the following schedule:

T_8	T_9	T_{10}
read(A)		
read(B)		
write(A)		
	read(A)	
	write(A)	
		read(A)
abort		

- If T_8 fails, T_9 and T_{10} must also be rolled back.
 - Not desirable, because it can take significant amount of work.



Cascadeless Schedules

- **Cascadeless schedules** — schedules where cascading rollbacks cannot occur
 - For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



Overview

- Transaction Isolation and Atomicity
- Transaction Isolation Levels
- Implementation of Isolation Levels
- Transactions as SQL Statements
- Assignments



Weak Levels of Consistency

- Serializability is useful concept to maintain consistency of database while executing transactions concurrently.
 - Ensuring serializability may allow too little concurrency
- Some applications are willing to live with weaker levels of consistency, allowing schedules that are not serializable
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate



Isolation Levels

- The SQL standard specifies several isolation levels:
 - **Serializable** — default; ensures serializable execution.
 - However, some database systems implement this isolation level in a manner that may allow nonserializable executions in some cases. (e.g., snapshot isolation by oracle or postgres prior to version 9)
 - **Repeatable read** — only committed data can be read and further requires that repeated reads of same data item by a transaction must return same value.
 - However, a transaction may not be serializable with respect to other transactions
 - **Read committed** — only committed data can be read, but does not require repeatable reads; successive reads of record may return different (but committed) values.
 - **Read uncommitted** — even uncommitted records may be read.

All isolation levels above disallow dirty writes, which is writes to a data item that has already been written by another transaction that has not yet committed or aborted



Serializability vs Performance

- An application designer may decide to use weaker isolation level for better system performance.
 - Serializability may force transactions to wait for other transactions or to abort due to deadlocks (Ch18)
- In the real-world, most major database systems run at a weaker isolation level by default, which can be explicitly set at database level or at start of transaction
 - Oracles, Postgres: Read Committed
 - MySQL: Repeatable Read
- Also, some systems actually implement a weaker level of isolation for the serializable level
 - Snapshot isolation is called "serializable" mode in Oracle and PostgreSQL versions prior to 9.1, which may cause confusion with the "real serializability" mode.
- Many means to implement isolation levels
 - Helpful to understand some details of implementation to avoid or minimize the chance of inconsistency



Overview

- Transaction Isolation and Atomicity
- Transaction Isolation Levels
- Implementation of Isolation Levels
- Transactions as SQL Statements
- Assignments



Concurrency Control

- Concurrency-control policies ensure that only acceptable schedules are generated
 - An example concurrency-control policy:
 - A transaction acquires a lock on the entire database before it starts and releases the lock after it has committed.
 - No other transaction is allowed to acquire the lock and must wait
 - The example above will generate serial schedules only, because only one transaction can execute at a time
 - Are serial schedules recoverable/cascadeless? (Yes)
 - The example above provides a poor degree of concurrency
- The goal of concurrency-control policies is to provide a high degree of concurrency while only generate schedules that are:
 - either **conflict or view serializable**
 - **recoverable and preferably cascadeless**



Implementation of Isolation Levels

- Concurrency-Control Mechanisms to implement isolation levels:
 - Locking
 - Lock may provide mutual exclusion on accessing database or data items
 - Lock for the entire database leads to poor degree of concurrency
 - Lock for each data item should be hold for long enough
 - Two-phase locking protocol can ensure serializability (Ch18)
 - Timestamps
 - Assigns timestamp for each transaction, typically when a transaction begins
 - Data items store two timestamps
 - Read timestamp
 - Write timestamp
 - Timestamps are used to detect out of order accesses
 - Multiple versions of each data item
 - Allow transactions to read from a “snapshot” of the database (**snapshot isolation** for an example)



Implementation of Isolation Levels

- Snapshot isolation overview:
 - Each transaction is given its own version (or snapshot) of the database when it begins
 - Read sees the data in the snapshot
 - Update is made to the snapshot – later, applied to “real” database
 - Transactions updating some data item can commit if there is no other transaction that has modified the same data items.
 - Otherwise, transactions cannot commit but abort.
- In snapshot isolation, read never need to wait unlike locking and read-only transaction cannot be aborted
 - Significant performance improvement compared to locking if most transactions are read-only
- Limitation is too much isolation
 - Concurrent updates in transactions may not be seen by each other
 - Can result in inconsistent database state due to nonserializable execution (Ch18)



Overview

- Transaction Isolation and Atomicity
- Transaction Isolation Levels
- Implementation of Isolation Levels
- Transactions as SQL Statements
- Assignments



Transactions as SQL Statements

- Our simple model so far only considers read and write.
 - Here write allows to update the value of data items assumed to exist
 - Yet, in the real world, data items may be created (i.e. inserted) and deleted
 - Writes can be either insert, update, or delete
 - Also, reading some data items are not trivial by looking at SQL statements.
 - Consider the following SQL query (Transaction T1)

```
select ID, name  
from instructor  
where salary > 90000
```

Which data items get accessed is depending on the Where clause

Choi	70000
Einstein	100000
Wu	90000
Brandt	95000



Phantom Phenomenon

Choi	70000
Einstein	100000
Wu	90000
Brandt	95000

- Consider concurrent execution of the following transactions

- Transaction T1:

```
select ID, name  
from instructor
```

```
where salary > 90000
```

- Transaction T2:

```
insert into instructor values ('11111', 'James', 'Marketing', 100000);
```

- Do they conflict?

- Consider T1 → T2

- It seems not, because T1 accesses only tuples containing 'Choi', 'Einstein', 'Wu', and 'Brandt'

- Now, consider T2 → T1

- There exists conflict, because T1 accesses the tuple containing 'James' as well now.

- So, there exists actual conflict which may or may not be revealed depending on the timing.

- This is called **phantom phenomenon**.

James	100000
-------	--------

"Phantom" data



Transactions as SQL Statements

- Consider the following example:
 - T1 uses the index to find data items that satisfy the condition in the where clause
 - Before T1 finishes its search, T2 inserts a new data item
 - Before T2 adds the corresponding index entry for the new data item, T1 finishes the search
 - Although T2 has already written the “phantom” data before T1 finishes reading, T1 may not see it because the index entry is not updated on time.
- Only considering each data item is not enough
 - Information used to find tuples must be considered for concurrency control as well
 - Index-locking protocols maximizing concurrency, ensuring serializability in spite of inserts, deletes and predicates in queries are discussed in Ch18.4.3



Predicate Locking

- Consider followings:

- Transaction T1:

```
select ID, name  
from instructor  
where salary > 90000
```

- Transaction T3 (Wu's salary = 90000):

```
update instructor  
set salary = salary * 0.9  
where name = 'Wu'
```

Choi	70000
Einstein	100000
Wu	90000 81000
Brandt	95000

- If the conflict is determined only based on low-level data item accesses, conflict may or may not exist for the example above:
 - If query processing of T1 searches through each tuple in the table, there exists conflict between T1 and T3
 - If index is used to find data items satisfying the condition in T1's where clause, there is no conflict
- Alternative concurrency control approach may look at predicate: “**predicate locking**” (Ch18.4.3)



Overview

- Transaction Isolation and Atomicity
- Transaction Isolation Levels
- Implementation of Isolation Levels
- Transactions as SQL Statements
- Assignments



Assignments

- Reading: Ch 17.7, 17.8, 17.9, Note 17.2, 17.10
- Practice Exercises: 17.7, 17.8, 17.9, 17.10, 17.11

Solutions to the Practice Exercises:

<https://www.db-book.com/Practice-Exercises/index-solu.html>

For 17.9, solution has an error. T_2 withdraws \$200 from the saving account



The End
