



Concurrency Control 1

Instructor: Beom Heyn Kim

beomheyunkim@hanyang.ac.kr

Department of Computer Science



Overview

- Lock-based Protocols
- Assignments



Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
 1. **shared** (S) *mode*.
 - Data item Q can only be read.
 - S-lock is requested using **lock-S** instruction. E.g., lock-S(Q)
 2. **exclusive** (X) *mode*.
 - Data item Q can be both read as well as written.
 - X-lock is requested using **lock-X** instruction. E.g., lock-X(Q)
- locks can be unlocked by calling unlock instruction. E.g., unlock(Q)
- A transaction makes lock requests to concurrency-control manager.
 - Transaction can proceed only after request is **granted**.



Lock-Based Protocols (Cont.)

- A concurrency-control manager decides whether to grant or not a lock depending on lock compatibility.
- For the given lock modes A and B :
 - If a transaction can be granted a lock of mode A on Q immediately, in spite of the presence of the mode B lock, then mode A is **compatible** with mode B . (A and B can be either S or X)
- Compatibility relation between lock modes can be represented as a lock-compatibility matrix.
 - E.g. compatibility matrix between S -lock and X -lock

	S	X
S	true	false
X	false	false

Shared mode is compatible with shared mode, but not with exclusive mode

- A S -lock can be granted even if other transactions hold S -locks
- A S -lock cannot be granted if another transaction holds a X -lock
- A X -lock cannot be granted if another transaction holds a lock



Schedule With Lock Grants

Consider the following transactions T1 and T2:

T_1 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
unlock(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(A).

T_2 : lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock(B);
display($A + B$).

Schedule With Lock Grants (Cont.)

Consider the following schedule:

T_1	T_2	concurrency-control manager
lock-X(B)		
read(B)		grant-X(B, T_1)
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T_2)
	unlock(A)	
	lock-S(B)	
	read(B)	grant-S(B, T_2)
	unlock(B)	
	display($A + B$)	
lock-X(A)		
read(A)		grant-X(A, T_1)
$A := A + 50$		
write(A)		
unlock(A)		

- In this schedule, T_2 sees the inconsistent state
 - Reading state where \$50 has been subtracted from A but not added to B ; then displaying the sum.
- The reason of having this schedule is because unlock(B) for X-lock on B was done too early.
 - Thus, T_2 could read inconsistent state of B
- Unlock instructions may be delayed to the end of each transaction



Schedule With Lock Grants (Cont.)

Consider the following transactions T3 and T4:

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A).

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).

Problem in the previous slide is not possible.
However, there is another problem.



Schedule With Lock Grants (Cont.)

Consider the following transactions T3 and T4:

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A).

1

Like before,
T4 gets
executed
after
write(B)
instruction
in T3

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).

2

This instruction
blocks and
waits for T3 to
release the
X-lock on B

3

This instruction
blocks and waits
for T4 to release
the S-lock on A

Is it possible to make any further progress?



Deadlock

- Consider the partial schedule (more focused view on the example schedule in the previous slide which leads to a deadlock)

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock, one of T_3 or T_4 must be rolled back and its locks released. (More on deadlock handling in Ch18.2)



Locking Protocol

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
 - It indicates when a transaction may lock and unlock each data item.
- Locking protocols restricts possible schedules
 - All possible schedules enforced by a locking protocol are serializable schedules
 - We shall see some locking protocols allowing only conflict-serializable schedules.
- The potential for deadlock exists in most locking protocols.
 - Deadlocks are necessary devil associated with locking.
 - Deadlocks are not desirable but still better than inconsistent states.



Starvation

- **Starvation** is also possible if concurrency control manager is badly designed. For example:
 - A transaction T1 attempts to hold a X-lock on an item Q.
 - A transaction T2 already holds a S-lock on an item Q.
 - Before T2 releases the S-lock, T3 requests for a S-lock on Q which is immediately granted; T1 should wait not only T2 releases the S-lock but T3 releases the S-lock
 - There may be a long sequence of more transactions requesting S-lock on Q before all previous transactions release their S-locks.
 - The same transaction T1 cannot get a X-lock and starves.
- Concurrency control manager can be designed to prevent starvation.
 - Trivially, there should be no other transaction holding a lock in non-compatible mode
 - Additionally, only grant the lock when there is no other transaction that already made a request



Overview

- Lock-based Protocols
- Assignments



Assignments

- Reading: Ch 18.1.1, 18.1.2
- Practice Exercises: n/a



The End
