



---

# Concurrency Control 2

**Instructor: Beom Heyn Kim**

[beomheyunkim@hanyang.ac.kr](mailto:beomheyunkim@hanyang.ac.kr)

Department of Computer Science

---



# Overview


---

- Lock-based protocols 2
- Assignments

# Two-Phase Locking Protocol

Two-phase locking protocol has two phases:

- Phase 1: **Growing Phase**
  - Transaction may obtain locks
  - Transaction may not release locks
- Phase 2: **Shrinking Phase**
  - Transaction may release locks
  - Transaction may not obtain any new locks
- Two-phase locking protocol ensures conflict serializability.
  - It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).





# Two-Phase Locking Protocol (Cont.)

- T1 and T2 are NOT two phase

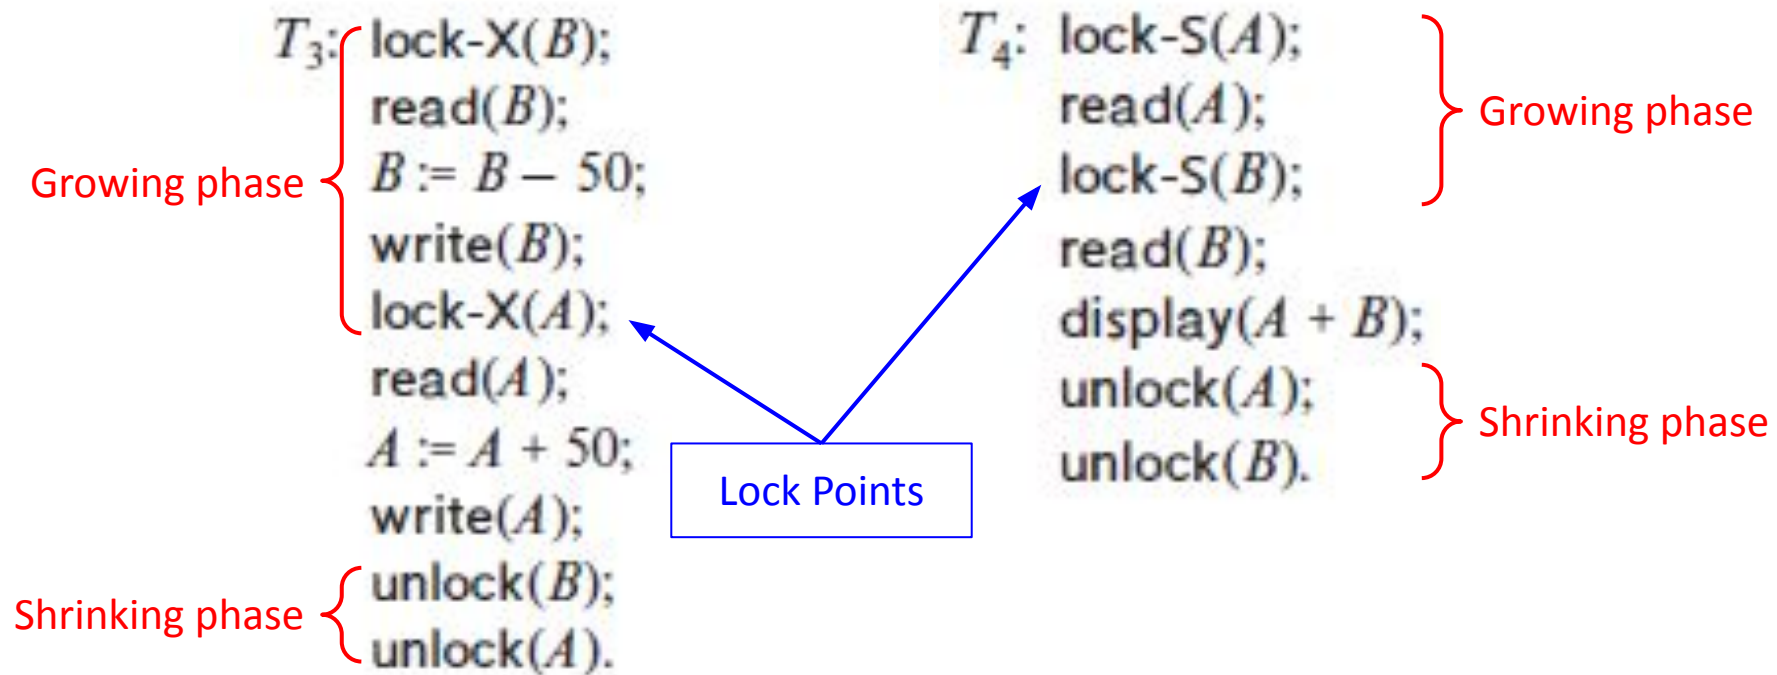
```
 $T_1$ : lock-X( $B$ );  
      read( $B$ );  
       $B := B - 50$ ;  
      write( $B$ );  
      unlock( $B$ );  
      lock-X( $A$ );  
      read( $A$ );  
       $A := A + 50$ ;  
      write( $A$ );  
      unlock( $A$ ).
```

```
 $T_2$ : lock-S( $A$ );  
      read( $A$ );  
      unlock( $A$ );  
      lock-S( $B$ );  
      read( $B$ );  
      unlock( $B$ );  
      display( $A + B$ ).
```



# Two-Phase Locking Protocol (Cont.)

- T3 and T4 are two phase



The two -phase locking protocol ensures conflict-serializability



# Two-Phase Locking Protocol (Cont.)

- Two-phase locking protocol is not free from deadlock.
  - Although  $T_3$  and  $T_4$  in the previous slide are two-phase, deadlock can occur. Recall:

$T_3$	$T_4$
lock-X( $B$ ) read( $B$ ) $B := B - 50$ write( $B$ )	
	lock-S( $A$ ) read( $A$ ) lock-S( $B$ )
lock-X( $A$ )	



# Two-Phase Locking Protocol (Cont.)

- The two-phase locking is not free from cascading rollback.

$T_5$	$T_6$	$T_7$
lock-X( $A$ ) read( $A$ ) lock-S( $B$ ) read( $B$ ) write( $A$ ) unlock( $A$ )	lock-X( $A$ ) read( $A$ ) write( $A$ ) unlock( $A$ )	lock-S( $A$ ) read( $A$ )



# Two-Phase Locking Protocol (Cont.)

- Extensions to basic two-phase locking to avoid cascading rollbacks:
  - **Strict two-phase locking**: a transaction must hold all its exclusive locks until it commits/aborts.
    - Ensures recoverability and avoids cascading rollbacks
  - **Rigorous two-phase locking**: a transaction must hold *all* locks until it commits/aborts.
    - Transactions can be serialized in the order in which they commit.
- Most databases implement rigorous two-phase locking, but refer to it as simply *two-phase locking*





# Lock Conversion

- Consider the following transactions:

```
T8: read(a1);  
    read(a2);  
    ...  
    read(an);  
    write(a1).
```

```
T9: read(a1);  
    read(a2);  
    display(a1 + a2).
```

- What types of locks needed to be acquired by T<sub>8</sub> and T<sub>9</sub>?
  - T<sub>8</sub> has write(a<sub>1</sub>), so it requests a lock-X
  - T<sub>9</sub> has only reads, so it requests a lock-S

```
lock-X(a1)  
T8: read(a1);  
    read(a2);  
    ...  
    read(an);  
    write(a1);  
unlock(a1)
```

T<sub>9</sub> cannot be executed between.  
Only serial execution is possible.  
Note: X-lock is only needed right before write(a<sub>1</sub>)



# Lock Conversion (Cont.)

- Lock conversion:
  - Exclusive locks can be **downgraded** to shared locks
  - Shared locks can be **upgraded** to exclusive locks
  - Upgrading can take place in only the growing phase
  - Downgrading can take place in only the shrinking phase
- More concurrency becomes possible:

$T_8$	$T_9$
lock-S( $a_1$ )	lock-S( $a_1$ )
read( $a_1$ );	read( $a_1$ );
lock-S( $a_2$ )	lock-S( $a_2$ )
read( $a_2$ );	read( $a_2$ );
lock-S( $a_3$ )	display( $a_1 + a_2$ );
lock-S( $a_4$ )	unlock( $a_1$ )
read( $a_n$ );	unlock( $a_2$ )
lock-S( $a_n$ )	
upgrade( $a_1$ )	
write( $a_1$ );	

- 2 phase locking with lock conversion ensures conflict-serializability



# Locking Protocols

- Given a locking protocol (such as 2PL – i.e. two-phase locking)
  - A schedule S is **legal** under a locking protocol if it can be generated by a set of transactions that follow the protocol
  - A protocol **ensures** serializability if all legal schedules under that protocol are serializable
    - e.g. 2PL protocol ensures serializability, so:
      - all legal schedules under a 2PL protocol are serializable
- However, 2PL protocol is not a necessary condition for serializability
  - There are conflict-serializable schedules that cannot be obtained if the two-phase locking protocol is used
    - There are other non-two-phase locking protocols ensuring conflict-serializability with additional requirements (e.g. graph-based protocols in 18.1.5)



# Automatic Acquisition of Locks

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:
  - if**  $T_i$  has a lock on  $D$
  - then**
    - read( $D$ )
  - else begin**
    - if necessary wait until no other transaction has a **lock-X** on  $D$
    - grant  $T_i$  a **lock-S** on  $D$ ;
    - read( $D$ )
  - end**



# Automatic Acquisition of Locks (Cont.)

- The operation **write**( $D$ ) is processed as:  
    **if**  $T_i$  has a **lock-X** on  $D$   
    **then**  
        write( $D$ )  
    **else begin**  
        if necessary, wait until no other transaction has any lock on  $D$   
        if  $T_i$  has a **lock-S** on  $D$   
        **then**  
            **upgrade** lock on  $D$  to **lock-X**  
        **else**  
            grant  $T_i$  a **lock-X** on  $D$   
        write( $D$ )  
    **end;**
- **All locks are released after commit or abort**




# Implementation of Locking

---

- A **lock manager** can be implemented as a separate process
- Transactions can send lock and unlock requests as messages
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
  - The requesting transaction waits until its request is answered
- The lock manager maintains an in-memory data-structure called a **lock table** to record granted locks and pending requests

# Lock Table



- Dark rectangles indicate granted locks, light colored ones indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
  - lock manager may keep a list of locks held by each transaction, to implement this efficiently



# Graph-Based Protocols


- Graph-based protocols are an alternative to 2PL protocol
- Impose a partial ordering  $\rightarrow$  on the set  $\mathbf{D} = \{d_1, d_2, \dots, d_h\}$  of all data items.
  - If  $d_i \rightarrow d_j$  then any transaction accessing both  $d_i$  and  $d_j$  must access  $d_i$  before accessing  $d_j$ .
  - Implies that the set  $\mathbf{D}$  may now be viewed as a directed acyclic graph, called a *database graph*.
- The *tree-protocol* is a simple kind of graph protocol.





# Tree Protocol

- Only exclusive locks are allowed.
- The first lock by  $T_i$  may be on any data item. Subsequently, a data  $Q$  can be locked by  $T_i$  only if the parent of  $Q$  is currently locked by  $T_i$ .
- Data items may be unlocked at any time.
- A data item that has been locked and unlocked by  $T_i$  cannot subsequently be relocked by  $T_i$ .





# Graph-Based Protocols (Cont.)

- The tree protocol ensures conflict serializability as well as freedom from deadlock.
- Unlocking may occur earlier in the tree-locking protocol than in the two-phase locking protocol.
  - Shorter waiting times, and increase in concurrency
  - Protocol is deadlock-free, no rollbacks are required
- Drawbacks
  - Protocol does not guarantee recoverability or cascade freedom
    - Need to introduce commit dependencies to ensure recoverability
  - Transactions may have to lock data items that they do not access.
    - increased locking overhead, and additional waiting time
    - potential decrease in concurrency
- Schedules not possible under two-phase locking are possible under the tree protocol, and vice versa.



# Overview

---

- Lock-based protocols 2
- Assignments



# Assignments

---

- Reading: Ch 18.1.3, 18.1.4, 18.1.5
- Practice Exercises: 18.1, 18.2, 18.5

Solutions to the Practice Exercises:

<https://www.db-book.com/Practice-Exercises/index-solu.html>



---

**The End**

---