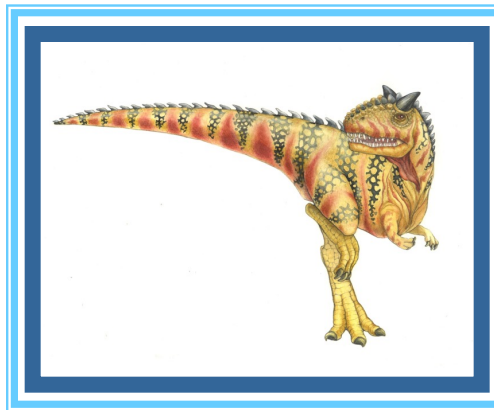


Chapter 4: Threads & Concurrency





Chapter 4: Threads

- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading (Thread Pool)
- Threading Issues
- Operating System Example





Objectives

- Identify the **basic components** of a thread, and **contrast** threads and processes
- Describe the **benefits** and **challenges** of designing multithreaded applications
- Illustrate different approaches to **implicit** threading including thread pools, fork-join, and Grand Central Dispatch
- Describe how the **Linux** operating systems represent threads





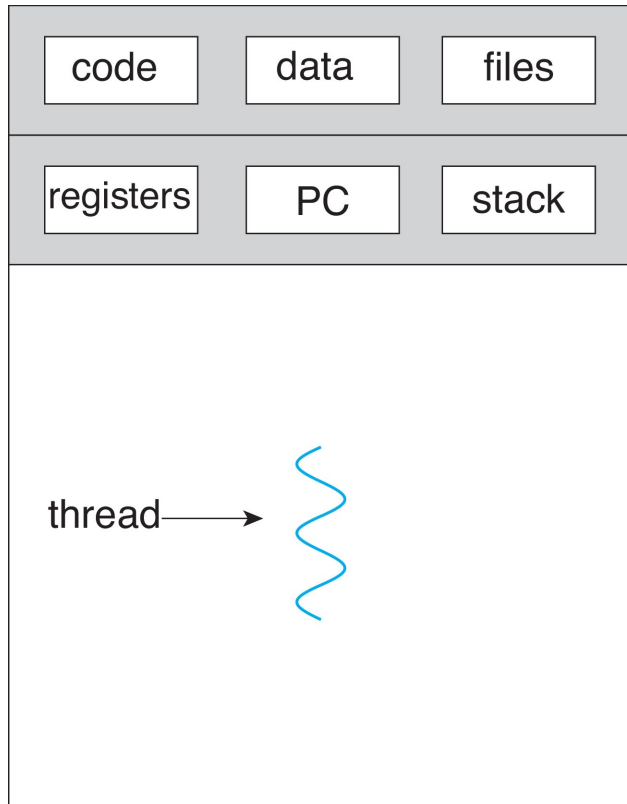
Motivation

- Most modern applications are **multithreaded**
- Threads run within application
- **Multiple tasks** with the application can be implemented by **separate threads**
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is **light-weight**
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

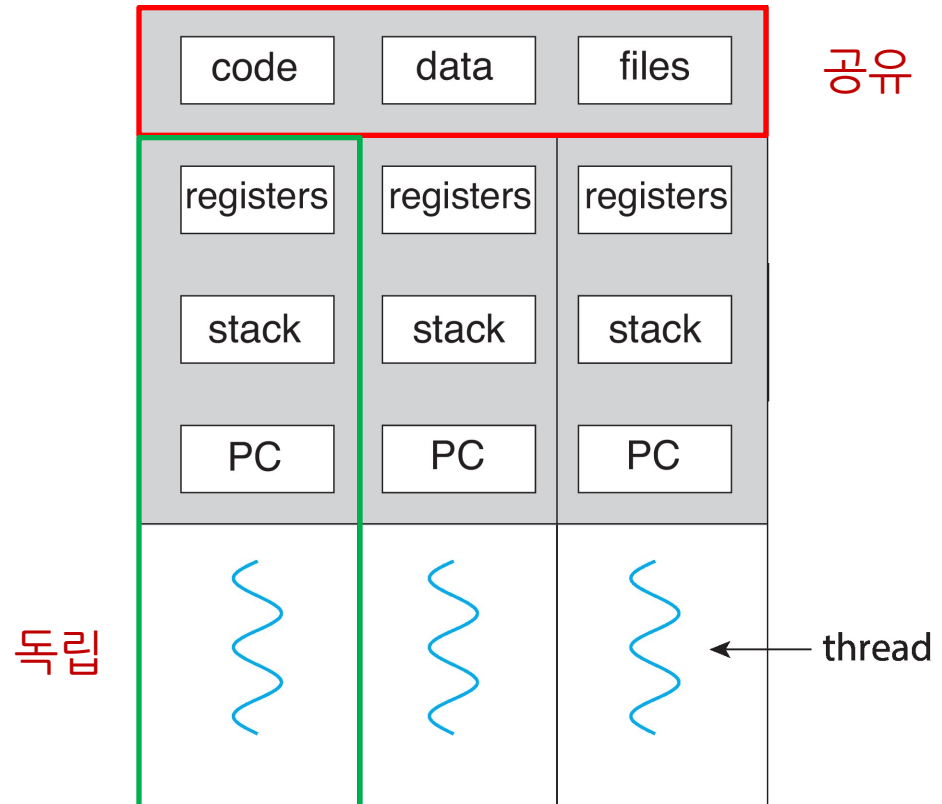




Single and Multithreaded Processes



single-threaded process

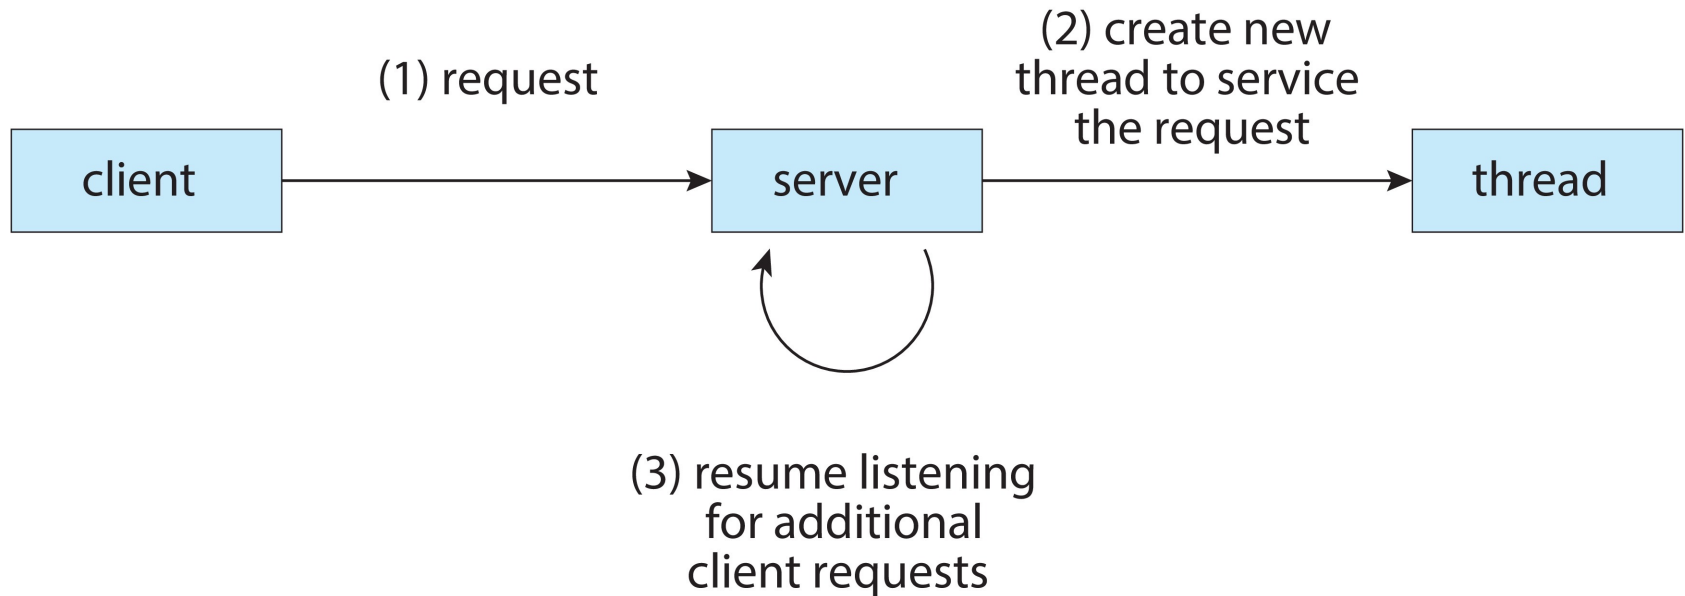


multithreaded process





Multithreaded Server Architecture





Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multicore architectures





Multicore Programming

- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:

- **Dividing activities**
- **Balance**
- **Data splitting**
- **Data dependency**
- **Testing and debugging**

병렬

- **Parallelism** implies a system can perform more than one task **simultaneously**

- **Concurrency** supports more than one task making **progress**

병행


- Single processor / core, **scheduler** providing concurrency






Concurrency vs. Parallelism

■ Concurrent execution on single-core system:



■ Parallelism on a multi-core system:





Multicore Programming

■ Types of parallelism

- **Data parallelism** – distributes **subsets** of the **same data** across multiple cores, **same operation** on each
- **Task parallelism** – distributing threads across cores, each thread performing **unique operation**


각 코어가 데이터를 나누어서 덧셈을 하는 경우 → data parallelism

같은 데이터에 대해 각 코어가 다른 통계 연산을 하는 경우 → task parallelism





Data and Task Parallelism





Amdahl's Law

- Identifies **performance gains** from **adding** additional **cores** to an application that has both serial and parallel components
- S is serial portion (%)
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$


Serial portion of an application has disproportionate effect on performance gained by adding additional cores

- But does the law take into account contemporary multicore systems?





Amdahl's Law






User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**
 - Windows threads
 - Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
 - Windows
 - Linux
 - Mac OS X
 - iOS
 - Android





User and Kernel Threads





Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many






Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block in a process
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads

장점: 원하는 만큼 사용자 스레드를 생성할 수 있음

주의: 커널 스레드가 운영체제 전체에 하나만 있다는 의미는 아님. 사용자 프로세스에 상응하는 커널 스레드가 하나만 있다는 의미임.






One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux

장점: 스레드가 블로킹되면
다른 스레드를 실행할 수 있음.

장점: 여러 개의 스레드를
다중코어에 매핑할 수 있음.



단점: 사용자 스레드와 커널 스레드가 1:1이기
때문에 사용자 스레드를 무한정 생성할 수 없음.






Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise **not very common**

장점: 사용자 스레드를 마음대로 생성
가능 + 멀티코어에서 병렬로 실행 가능

단점: 구현하기 어려움




커널 스레드의 수는 응용 또는
코어의 수에 따라 다를 수 있다





Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread





Thread Libraries

- **Thread library** provides programmer with **API** for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS (**results in a system call to the kernel**)

Asynchronous threading: 부모와 자식 스레드가 독립적으로 병행 실행

Synchronous threading: 부모 스레드는 모든 자식 스레드가 종료될 때까지 기다림

POSIX Pthreads: user level or kernel level

Windows: kernel level

Java threads: depending on OS





Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification, not implementation***
- **API specifies behavior of the thread library**, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)





Pthreads Create

- `#include <pthread.h>`
- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
void *(*start_routine)(void *), void* arg);`
 - `thread`: thread ID
 - `attr`: thread attributes
 - `start_routine`: address of start routine
 - `arg`: argument to the start routine





Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```





Pthreads Example (cont)

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```





Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

부모가 다수의 자식 스레드가
종료할 때까지 기다리는 경우





Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```





Windows Multithreaded C Program (Cont.)

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```

WaitForMultipleObjects(N, THandles,
TRUE, INFINITE);

갯수 → N
배열 → THandles
모두 기다림 → TRUE
시간 무제한 → INFINITE





Implicit Threading

- Growing in popularity as numbers of **threads increase**, program correctness more **difficult with explicit** threads
- Creation and management of threads done by **compilers** and **run-time libraries** rather than programmers
- Five methods explored
 - Thread Pools
 - Fork-Join
 - OpenMP (**Open Multi-Processing**)
 - Grand Central Dispatch (**Apple's thread pool**)
 - Intel Threading Building Blocks (**C++ template library**)

개발자는 병렬로 수행할 수 있는 task를 식별하는데 주력하고,
스레드는 컴파일러나 런타임 라이브러리가 담당





Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - ▶ i.e. Tasks could be scheduled to run periodically
- Windows API supports thread pools: Task 실행과 task 생성을 분리

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```


QueueUserWorkItem() 함수를 통해 pool에 있는 스레드가 PoolFunction을 실행하게 함






Fork-Join Parallelism

- Multiple threads (tasks) are **forked**, and then **joined**.





Fork-Join Parallelism





OpenMP

컴파일러 지시문

- Set of **compiler directives** and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel fork
Create as many threads as there are
cores

코어의 수만큼

join

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```

코어의 수만큼 메시지가 중복되서 출력됨.
모든 코어가 동일한 코드를 실행하기 때문에
출력은 겹쳐서 보일 수 있음





OpenMP

■ Run the for loop in parallel

코어의 수만큼 스레드를 생성하고
N개의 iteration을 스레드에 분할 매핑

```
#pragma omp parallel for  
for (i = 0; i < N; i++) {  
    c[i] = a[i] + b[i];  
}
```

- 이 예제에서는 data의 종속성이 없어서 병렬 실행이 가능함. 그러나 종속성이 있는 경우에는 주의해야 함
- 스레드의 수를 지정할 수 있음 `num_threads(N)`
- 매핑 순서에 변화를 줄 수 있음 `schedule(dynamic)`

Linux, Windows, macOS용 컴파일러 존재






Grand Central Dispatch

Thread pool의 크기를 자동 조절

하부는 POSIX pthread로 구현

- Apple technology for macOS and iOS operating systems
- Extensions to C, C++ and Objective-C languages, API, and run-time library
- Allows **identification of parallel sections** 코드의 어느 부분을 병렬로 실행시킬 수 있는지 개발자가 지정할 수 있다
- **Manages** most of the **details** of threading GCD는 스레딩의 세세한 것을 관리함
- Block is in “`{ }`”:  이 부분을 스레드에 할당해서 실행

```
^{ printf("I am a block"); }
```

- Blocks placed in **dispatch queue**
 - **Assigned** to **available thread** in thread pool when removed from queue

Block은 dispatch queue에 들어갔다가
사용 가능한 스레드가 있으면 할당함





Grand Central Dispatch

- Two types of dispatch queues: Process 내에 있기 때문에 private dispatch queue라 부름
- **serial** – blocks removed in **FIFO** order **one at a time**, queue is per process, called **main queue**
 - ▶ Programmers can create additional serial queues within program

- **concurrent** – removed in **FIFO** order but **several** may be removed **at a time**

- ▶ **Four system wide queues** divided by quality of service:

- | | | |
|---|------------------------------|-------------------|
| <div style="display: flex; align-items: center;"><div style="writing-mode: vertical-rl; transform: rotate(180deg); color: red; font-weight: bold;">빠른
응답</div><div style="color: blue; font-size: 2em; margin: 0 10px;">↑</div></div> | ○ QOS_CLASS_USER_INTERACTIVE | 신속한 응답이 요구될 때 |
| | ○ QOS_CLASS_USER_INITIATED | 적당한 시간의 응답이 요구될 때 |
| | ○ QOS_CLASS_USER_UTILITY | 긴 시간이 요구될 때 |
| | ○ QOS_CLASS_USER_BACKGROUND | 시간에 관계없는 task일 때 |

System-wide queue이기 때문에 global dispatch queue라 부름





Intel Threading Building Blocks (TBB)

- **Template library** for designing parallel C++ programs
- A serial version of a simple for loop

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

- The same for loop written using TBB with `parallel_for` statement:

```
parallel_for (size_t(0), n, [=](size_t i) {apply(v[i]);});
```

Windows	} Open-source or Commercial
Linux	
macOS	





Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage

내부에서 발생 ex) illegal
memory access, divide by zero

외부에서 발생 ex) Ctrl-C, timer





Semantics of `fork()` and `exec()`

- Does `fork()` duplicate ① only the calling thread or ② all threads?
 - Some UNIXes have two versions of `fork`
- `exec()` usually works as normal – replace the running process including all threads

`fork()` 후에 `exec()`가 바로 수행되는
경우 ①이 유리함

Linux는 ①을 지원함

“Don't use both threads and forks”





Signal Handling

- **Signals** are used in UNIX systems to **notify** a process that a particular **event** has occurred.
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can **override default**
 - For **single-threaded**, signal delivered to process





Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies → Ex) divide by zero
 - Deliver the signal to every thread in the process → Ex) Ctrl-C
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

Unix signal 생성 함수

```
kill(pid_t pid, int signal);  
pthread_kill(pthread_t tid, int signal);
```

대부분의 Unix 계열의 OS에서는 스레드마다 받는 또는 거부하는 signal을 설정할 수 있음. 따라서 다중 스레드인 경우 signal을 허용하는 첫 번째 스레드가 처리함





Thread Cancellation

- Terminating a thread **before it has finished**
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread **immediately** → 할당된 자원을 모두 free하지 못할 경우가 생길 수 있음
 - **Deferred cancellation** allows the target thread to **periodically check** if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);  
  
/* wait for the thread to terminate */  
pthread_join(tid, NULL);
```

↓
Cancellation point에서
스레드를 캔슬함.
일반적으로 read()와 같은
blocking system call이
cancellation point가 됨.





Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but **actual cancellation depends on thread state**

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation **disabled**, cancellation **remains pending** until thread enables it
- **Default type is deferred**
 - **Cancellation only occurs** when thread reaches **cancellation point**
 - ▶ I.e. `pthread_testcancel()` 인위적으로 cancellation point를 지정할 때 사용
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through **signals**

Q: 이 함수를 호출하면
다시 이 자리로 리턴될까?





Thread-Local Storage

- **Thread-local storage (TLS)** allows each thread to have its **own copy** of data
- **Useful** when you do not have control over the thread creation process (i.e., when using a **thread pool**)
- Different from local variables
 - Local variables **visible only during single function** invocation
 - TLS **visible across function** invocations
- Similar to **static** data
 - TLS is **unique** to each thread

스레드는 자신이 속한 프로세스의 데이터를 공유한다. 그러나 경우에 따라서는 자신만의 데이터가 필요할 수 있다. → TLS

스레드 pool을 사용하면 사용자는 스레드 생성에 관여하지 않기 때문에 TLS가 유용함.





Linux Threads

- Linux refers to them as **tasks** rather than **threads**
 - Thread creation is done through `clone()` system call
 - `clone()` allows a child task to **share the address space** of the parent task (process)
 - Flags control behavior
- fork() and pthread_create() eventually call clone().*

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

*fork()는 이 자료구조를 copy함 (복사로 처리)
clone()은 이 자료구조를 point함 (링크로 처리)*



End of Chapter 4

