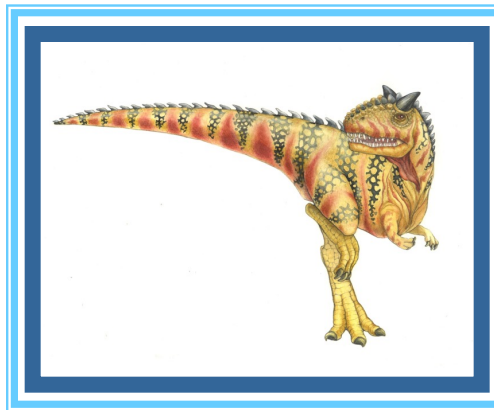


# Chapter 9: Main Memory

---





# Chapter 9: Memory Management

---

- Background
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Swapping
- Example: The Intel 32 and 64-bit Architectures
- Example: ARMv8 Architecture





# Objectives

---

- To provide a detailed description of various ways of **organizing memory** hardware
- To discuss various **memory-management** techniques,
- To provide a detailed description of the Intel **Pentium**, which supports both pure segmentation and segmentation with paging





# Background

---

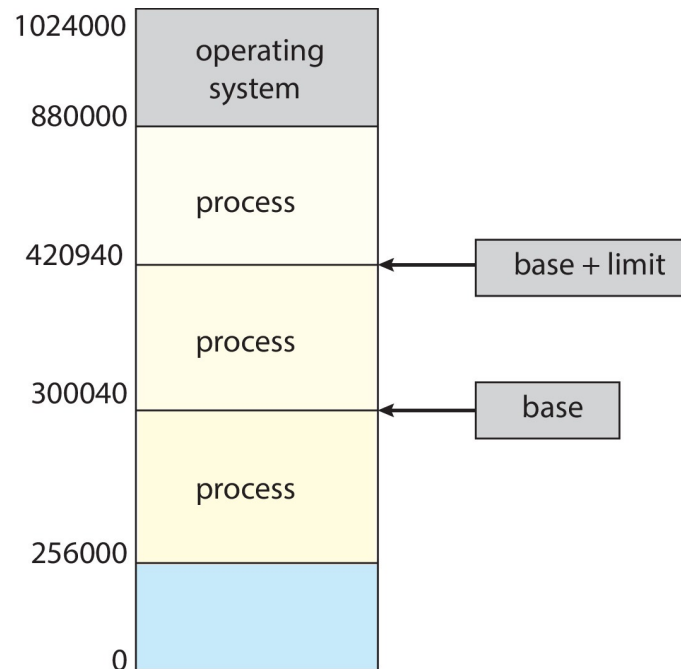
- Program must be brought (from disk) into memory and placed within a process for it to be run
- **Main memory** and **registers** are **only** storage **CPU** can **access directly**
- Memory unit only sees a stream of:
  - addresses + read requests, or
  - address + data and write requests
- Register access is done in one CPU clock (or less)
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation





# Protection

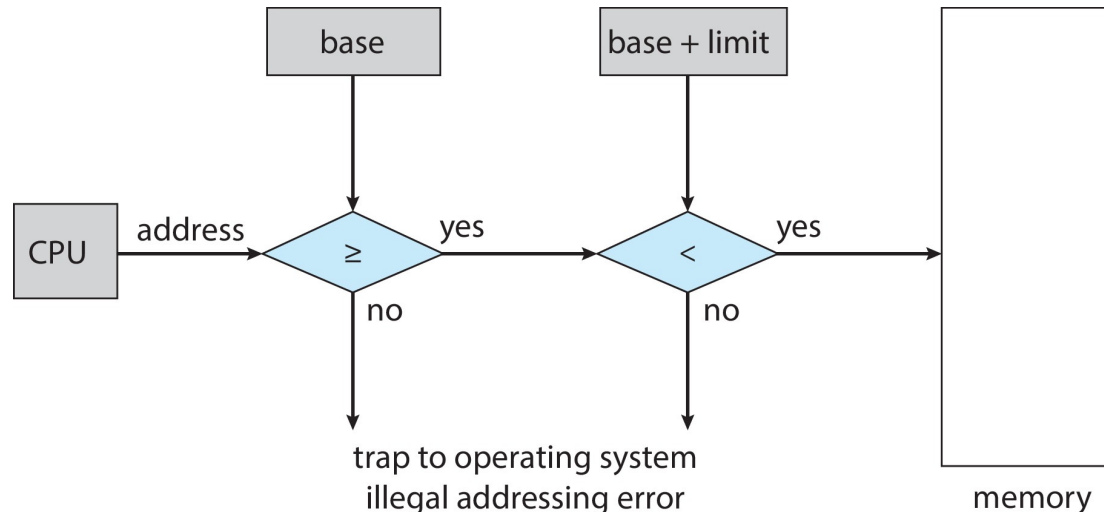
- Need to ensure that a process can access only those addresses in its address space.
- We can provide this protection by using a pair of **base** and **limit registers** to define the logical address space of a process





# Hardware Address Protection

- CPU must **check every memory access** generated in user mode to be sure it is between base and limit for that user



- the instructions to **loading the base and limit registers** are **privileged**

이 레지스터를 OS는 설정할 수 있지만  
사용자 프로그램은 변경할 수 없어야 한다.





# Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
  - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
  - How can it not be?
- Addresses represented in different ways at different stages of a program's life
  - **Source** code addresses usually **symbolic**
  - **Compiled** code addresses **bind** to **relocatable** addresses
    - ▶ i.e. “14 bytes from beginning of this module”
  - **Linker or loader** will bind relocatable addresses to **absolute** addresses
    - ▶ i.e. 74014
  - Each binding maps one address space to another





# Binding of Instructions and Data to Memory

- **Address binding** of instructions and data to memory addresses can happen at three different stages
  - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
  - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
  - **Execution time:** Binding **delayed until run time** if the process can be moved during its execution from one memory segment to another
    - ▶ **Need hardware support** for address maps (e.g., base and limit registers)

swapping, paging


가상메모리 개념 필요







# Multistep Processing of a User Program





# Logical vs. Physical Address Space


- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





# Memory-Management Unit (MMU)

- Hardware device that at run time **maps virtual to physical** address




- Many methods possible, covered in the rest of this chapter





# Memory-Management Unit (Cont.)

- Consider simple scheme. which is a generalization of the base-register scheme.
- The **base register** now called **relocation register**
- The value in the relocation register is **added** to every address generated by a user process at the time it is sent to memory



사용자 프로그램은 논리적인 주소를 사용할 뿐  
실제 물리적인 주소를 전혀 보지 못한다.





# Dynamic Loading

- The **entire** program does **not** need to be in memory to execute
- Routine is **not loaded until** it is **called**
- **Better** memory-space **utilization**; **unused** routine is **never** loaded
- All routines kept on disk in **relocatable** load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - Implemented through program design
  - OS can help by providing libraries to implement dynamic loading





# Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the **binary program image**
- **Dynamic linking** –linking postponed until execution time
- Small piece of code, **stub**, used to **locate** the appropriate **memory-resident library** routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
  - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**
- Consider applicability to patching system libraries
  - Versioning may be needed

사용된 버전에  
따라서 합당하게  
링크되어야 한다.

Dynamic loading과 달리 dynamic linking과  
shared libraries는 OS의 도움이 필요하다.





# Contiguous Allocation

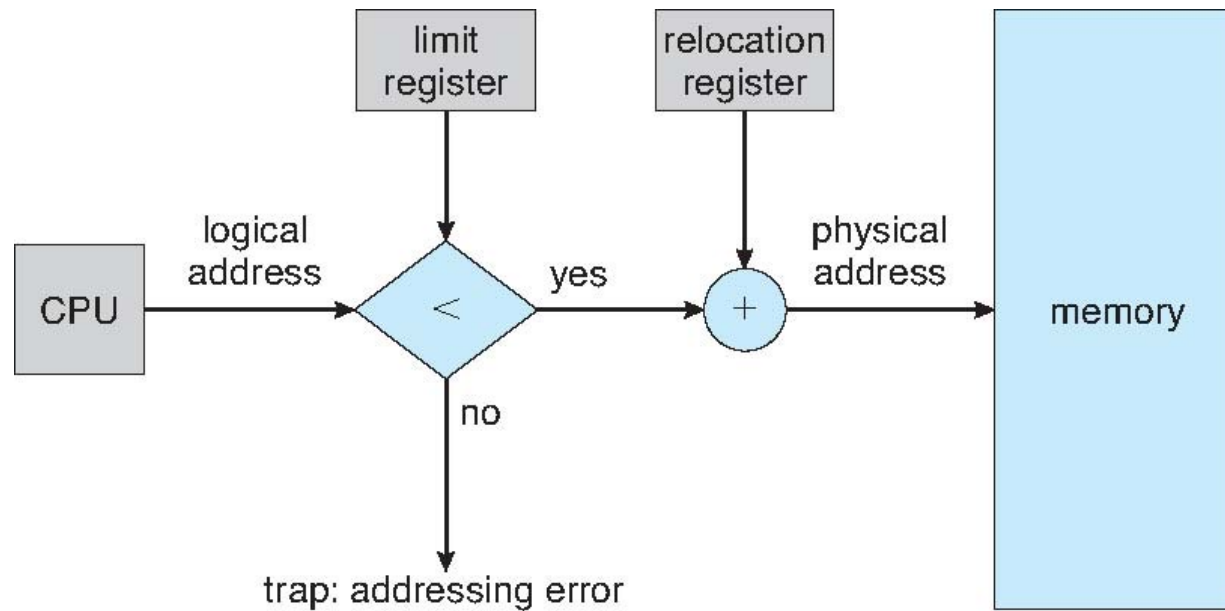
- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one **early** method
- Main memory usually into two **partitions**:
  - Resident **operating system**, usually held in **low** memory with interrupt vector
  - **User processes** then held in **high** memory
  - Each process contained in **single contiguous** section of memory

↓  
논리적으로는 연속적인 메모리 내에 있어야  
하지만 물리적으로는 그렇지 않을 수도 있다.





# Hardware Support for Relocation and Limit Registers








# Variable Partition

## ■ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole **large enough to accommodate it**
- Process exiting frees its partition, **adjacent** free partitions **combined**
- Operating system maintains information about:  
**a)** allocated partitions    **b)** free partitions (hole)





# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- **First-fit**: Allocate the **first** hole that is big enough
- **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit**: Allocate the **largest** hole; must also search entire list
  - Produces the largest leftover hole

**First-fit** and **best-fit** better than worst-fit in terms of speed and storage utilization

First-fit이 일반적으로 가장 빠름





# Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is **not contiguous**
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but **not being used**
- First fit analysis reveals that given  $N$  blocks allocated,  $0.5 N$  blocks lost to fragmentation
  - $1/3$  may be unusable -> **50-percent rule**

**First-fit**은  $N$ 개의 블록이 할당되었다면  $0.5N$ 개의 블록이 조각나서 사용할 수 없다는 분석이 있다. 이렇게 되면 전체에서  $0.5N/(N+0.5N) = 1/3$ 이 못쓰게 되는데 이것을 **50% 법칙**이라고 한다.





# Fragmentation (Cont.)

## ■ Reduce external fragmentation by **compaction**

- **Shuffle** memory contents to place all free memory **together** in **one large block**
- Compaction is possible *only* if relocation is dynamic, and is done at execution time
- I/O problem

조각모음은 실행중인 코드를 옮길 수 있어야 가능하다.

- ▶ Latch job in memory while it is involved in I/O
- ▶ Do I/O only into OS buffers

## ■ Now consider that **backing store** has same fragmentation problems

저장 공간은 fragmentation 문제를 가지고 있다

I/O를 실행중에 조각모음을 하게 되면 버퍼의 위치가 변경되서 문제가 발생함 → 해결책으로는 I/O를 OS 버퍼에서만 수행함





# Paging

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of **varying sized** memory chunks
- Divide **physical** memory into fixed-sized blocks called **frames**
  - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide **logical** memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size  **$N$**  pages, need to find  **$N$**  free frames and load program
- Set up a **page table** to **translate logical to physical** addresses
- Backing store likewise split into pages
- **Still** have **Internal** fragmentation





# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an **index** into a **page table** which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit

page number	page offset
$p$	$d$
$m - n$	$n$

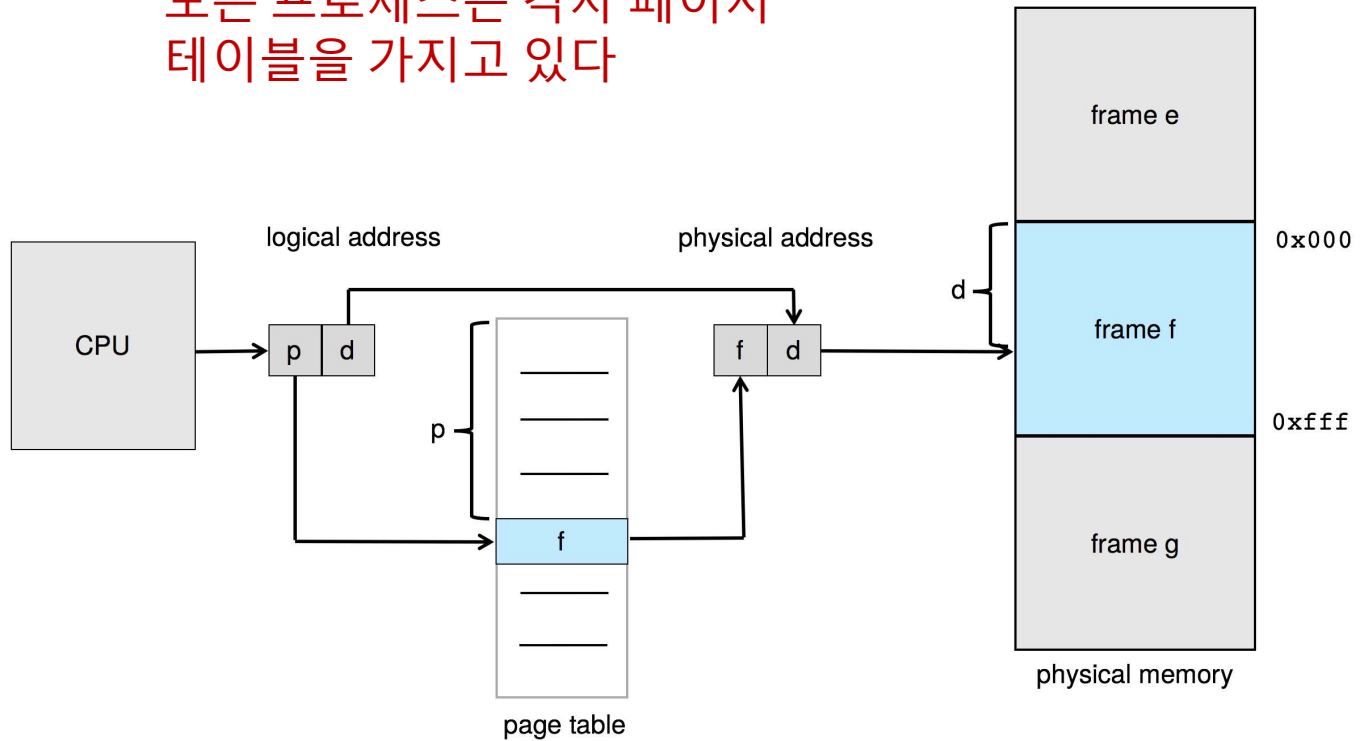
- For given logical **address space**  $2^m$  and **page size**  $2^n$





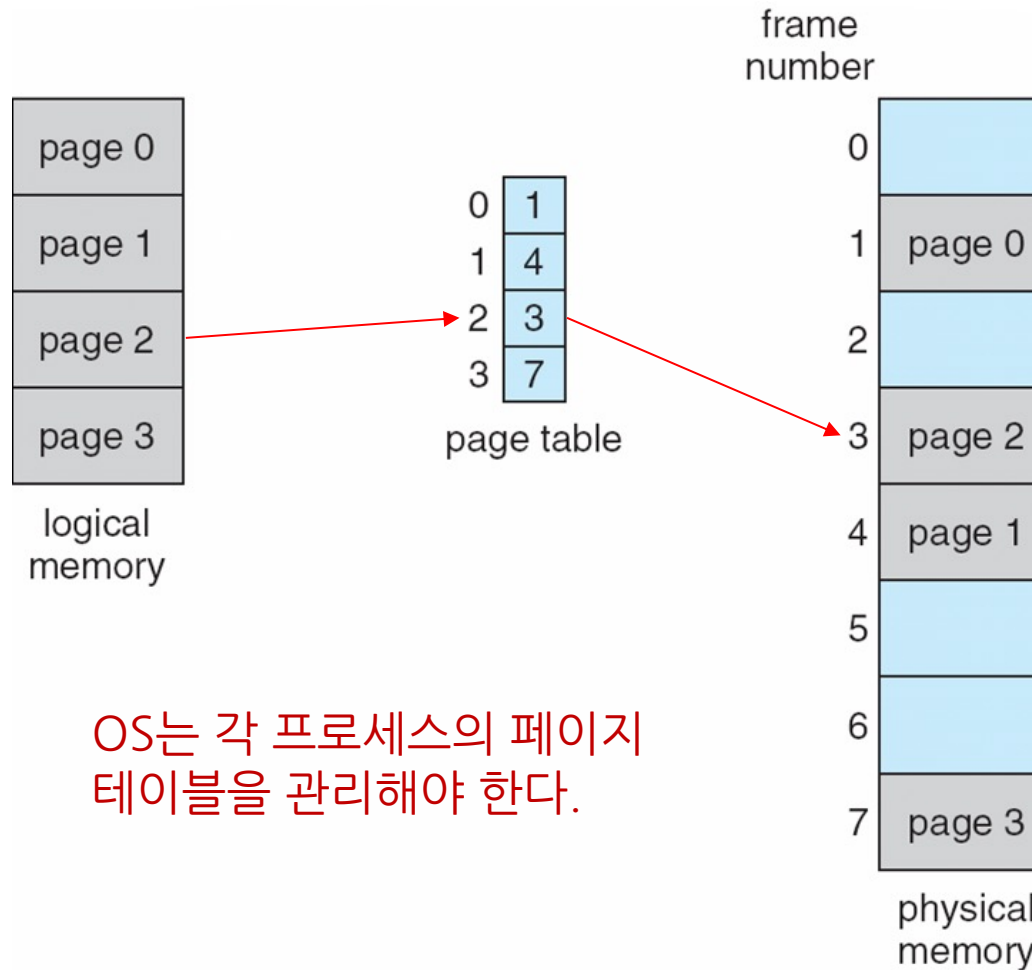
# Paging Hardware

모든 프로세스는 각자 페이지 테이블을 가지고 있다





# Paging Model of Logical and Physical Memory



OS는 각 프로세스의 페이지 테이블을 관리해야 한다.








# Paging Example

- Logical address:  $n = 2$  and  $m = 4$ . Using a page size of 4 bytes and a physical memory of 32 bytes (8 pages)



논리주소: 4 bits  
물리주소: 5 bits  
페이지 크기: 4 bytes





# Paging -- Calculating internal fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation =  $1 / 2$  frame size
- So small frame sizes desirable? → 페이지 테이블 엔트리 증가
- But each page table entry takes memory to track
- Page sizes growing over time → 일반적으로 페이지 크기 증가 추세
  - Solaris supports two page sizes – 8 KB and 4 MB

가장 인기 있는 크기는 4KB와 8KB이다


페이지 테이블에는 frame 값 이외에도 추가 정보를 저장할 공간이 충분히 있어야 한다.






# Free Frames

Frame table: 각 프레임이 어떤 프로세스의 어떤 페이지에 할당되었는지 정보가 있다.



(a)

Before allocation



(b)

After allocation





# Implementation of Page Table

- Page table is kept in main memory
  - **Page-table base register (PTBR)** points to the page table
  - **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires **two** memory accesses
  - One for the **page table** and one for the **data / instruction**
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **translation look-aside buffers (TLBs)** (also called **associative memory**).

content-addressable memory





# Translation Look-Aside Buffer

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – **uniquely identifies each process** to provide address-space protection for that process
  - **Otherwise** need to **flush** at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access

어떤 CPU는 instruction TLB와 data TLB를 별도로 제공하기도 한다.  
TLB는 계층구조를 가지기도 한다.

Intel Core i7 CPU: 128-entry L1  
instruction TLB + 64-entry L1  
data TLB + 512-entry L2 TLB

“제거할 수 없다”. 일반적으로  
주요한 커널 코드의 엔트리는  
wired down되어 있다.





# Hardware

## ■ Associative memory – parallel search

Page #	Frame #

## ■ Address translation (p, d)


- If p is in associative register, get frame # out
- Otherwise get frame # from page table in memory

TLB lookup은 instruction pipeline의 일부이기 때문에 성능 손실이 전혀 없다.





# Paging Hardware With TLB





# Effective Access Time

- Hit ratio – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
  - If we find the desired page in TLB then a mapped-memory access take 10 ns
  - Otherwise we need two memory access so it is 20 ns
- **Effective Access Time (EAT)**

$$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$

implying 20% slowdown in access time

- Consider a more realistic hit ratio of 99%,
$$\text{EAT} = 0.99 \times 10 + 0.01 \times 20 = 10.1\text{ns}$$
implying only 1% slowdown in access time.







# Memory Protection


- Memory protection implemented by associating **protection bit** with each frame to indicate if **read-only** or **read-write** access is allowed
  - Can also **add more bits** to indicate page **execute-only**, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel






# Valid (v) or Invalid (i) Bit In A Page Table

만일 14-bit 주소 공간을 사용한다면 그림처럼  
페이지 크기가 2K인 8개의 페이지 할당이 가능하다.



frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table



주어진 프로그램이 0~10,468 주소만 사용하는  
경우라면 (즉, 5개 페이지만 사용), 페이지 5번을 사용할  
수 없음에도 불구하고 그림처럼 valid로 나타나는 오류가  
발생할 수 있다. 우리는 페이지 5~7번을 "페이징으로  
인한 내부 조각 (internal fragmentation of  
paging)"이라고 한다. PTLR을 사용하면 해결이 가능함.





# Shared Pages

## ■ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., **text editors**, **compilers**, **shared libraries**)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

어떤 OS는 shared page를 통해 shared memory를 구현함


## ■ Private code and data

- Each process keeps a **separate copy** of the code and data
- The pages for the private code and data can appear anywhere in the logical address space





# Shared Pages Example



표준 C 라이브러리 공유 사례





# Structure of the Page Table


- Memory structures for paging can get **huge** using straight-forward methods
  - Consider a **32-bit** logical address space as on modern computers
  - Page size of **4 KB** ( $2^{12}$ )
  - Page table would have 1 million entries  $(2^{32} / 2^{12}) = 2^{20}$
  - If each entry is 4 bytes → each process 4 MB of physical address space for the page table alone =  $4 * 2^{20} = 4\text{MB}$  page table size
    - ▶ Don't want to allocate that contiguously in main memory
  - One simple solution is to divide the page table into smaller units
    - ▶ Hierarchical Paging
    - ▶ Hashed Page Tables
    - ▶ Inverted Page Tables





# Hierarchical Page Tables


- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table





# Two-Level Paging Example

- A logical address (on **32-bit** machine with **4K page size**) is divided into:
  - a **page number** consisting of **20** bits
  - a **page offset** consisting of **12** bits
- Since the page table is paged, the page number is further divided into:
  - a **10-bit** page number
  - a **10-bit** page offset
- Thus, a logical address is as follows:

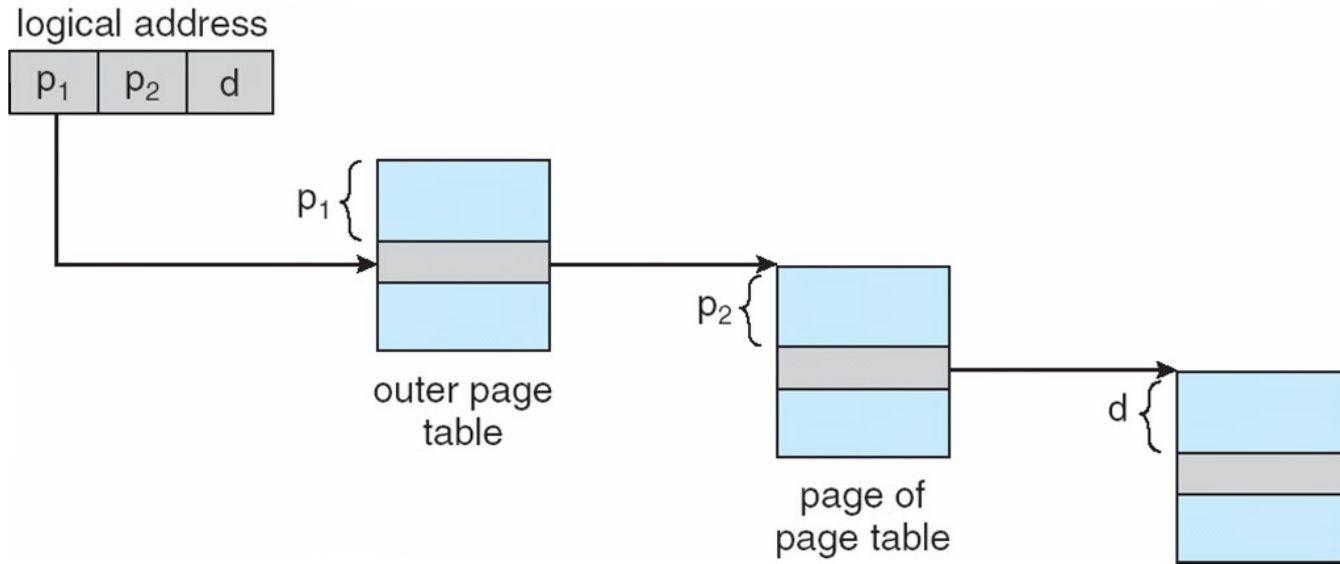


- where  $p_1$  is an index into the **outer page table**, and  $p_2$  is the displacement within the page of the **inner page table**
- Known as **forward-mapped page table**





# Address-Translation Scheme







# 64-bit Logical Address Space

- Even two-level paging scheme **not** sufficient
- If page size is **4 KB** ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes
- One solution is to add a 2<sup>nd</sup> outer page table
- But in the following example the 2<sup>nd</sup> outer page table is still  $2^{34}$  bytes in size
  - ▶ And possibly **4** memory access to get to **one** physical memory location





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12

64-bit 아키텍처에서 계층적 페이지 테이블을 사용하는 것은 계층이 늘어나서 어려움이 따른다. Intel x86-64의 경우 64 비트 중 48 비트만 논리주소로 사용하고 4계층 페이지 테이블을 사용한다.





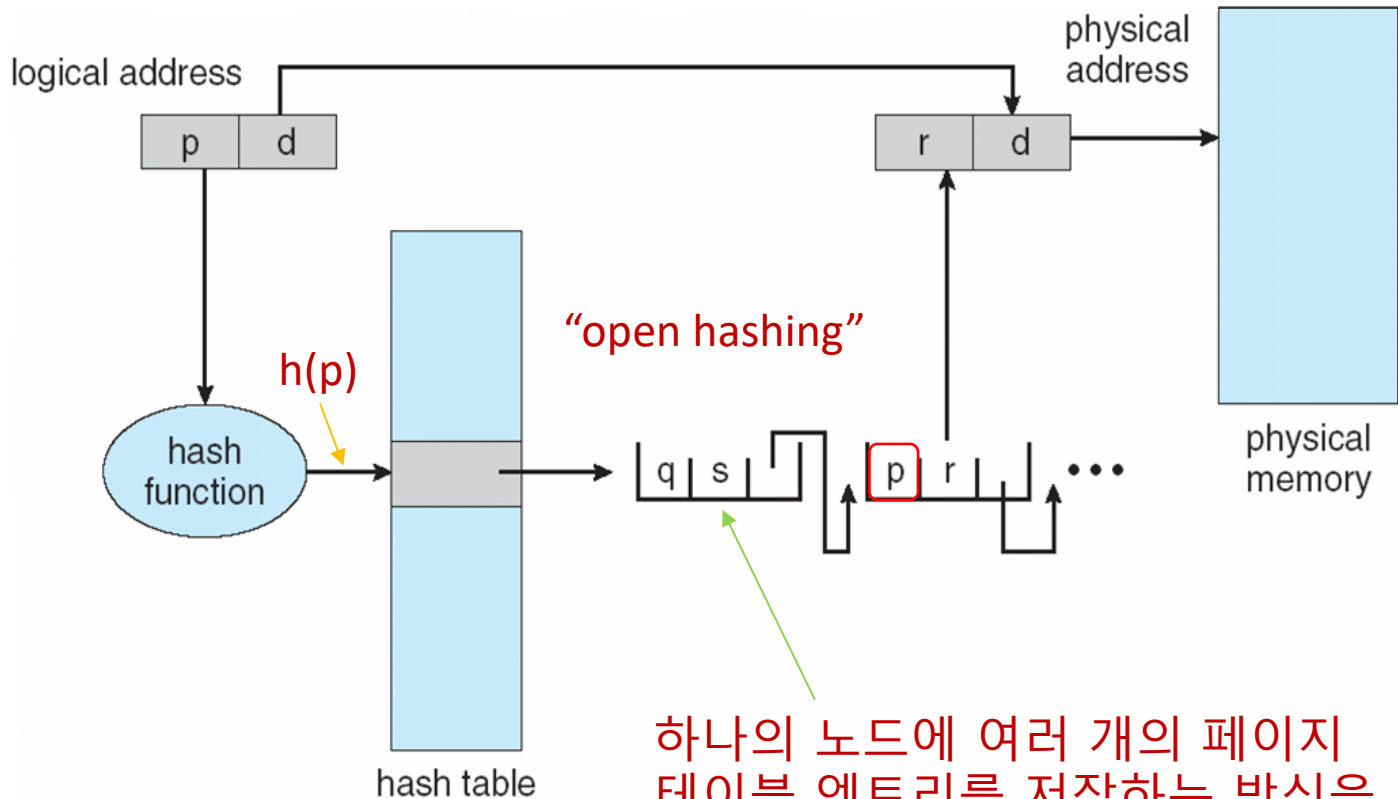
# Hashed Page Tables

- Common in address spaces  $> 32$  bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is clustered page tables
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for sparse address spaces (where memory references are non-contiguous and scattered)





# Hashed Page Table



하나의 노드에 여러 개의 페이지 테이블 엔트리를 저장하는 방식을 "clustered page table"이라고 한다.





# Inverted Page Table

프로세스마다 페이지 테이블을 가지지 말고  
물리적 페이지를 직접 관리하는 방법

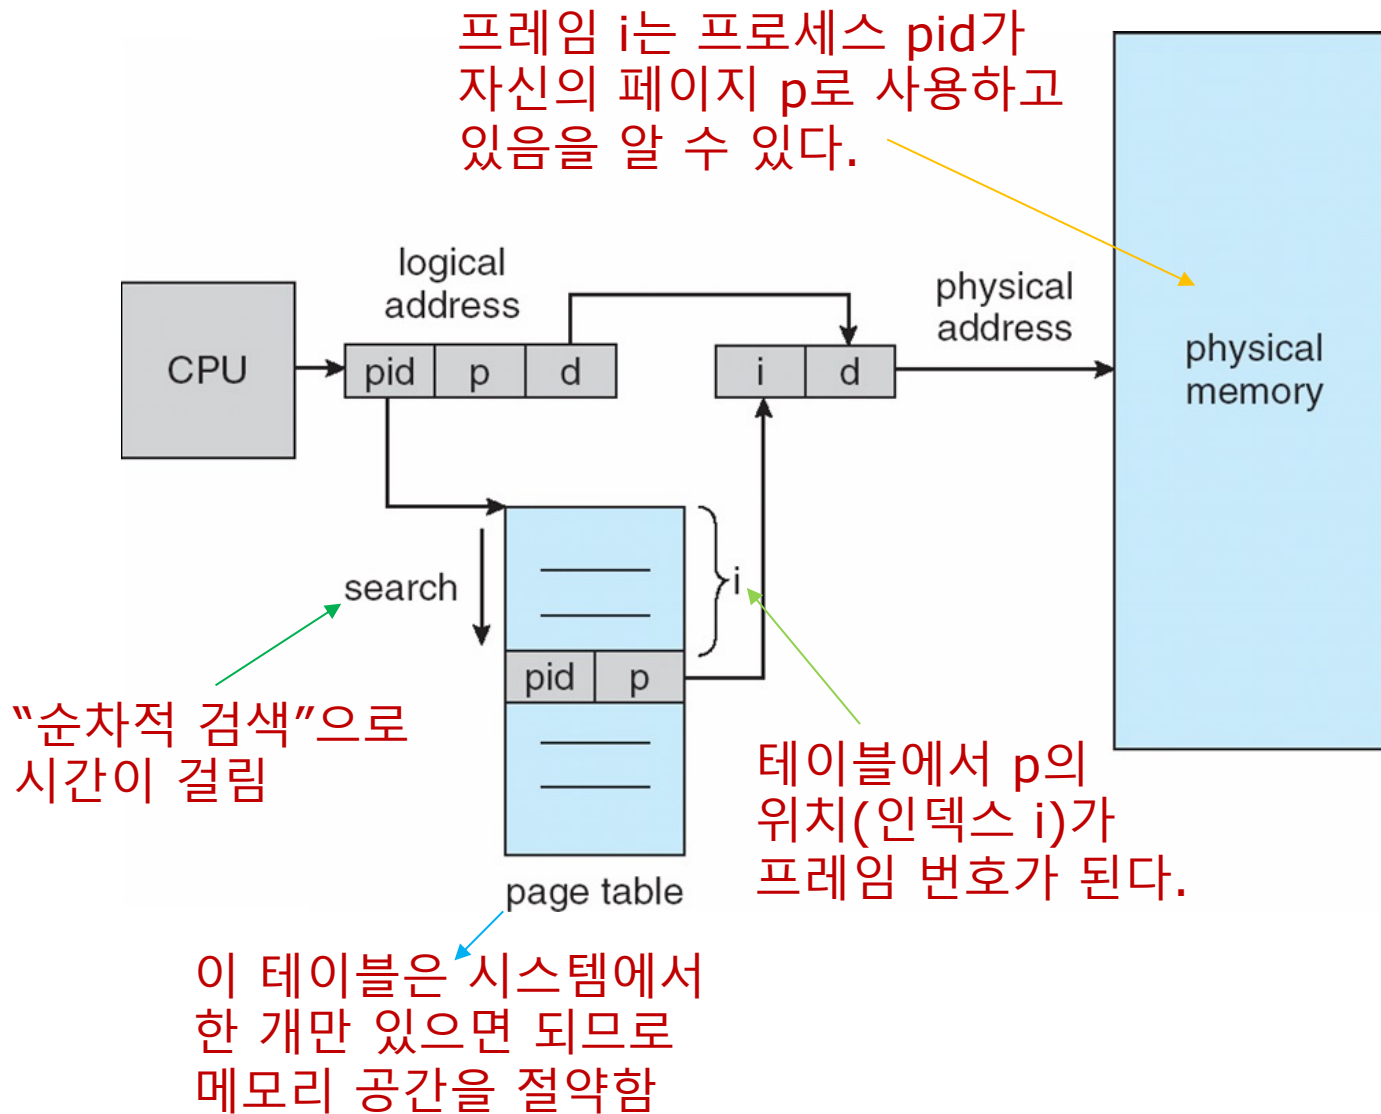
- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address

하나의 물리적 주소에 하나의 가상메모리 주소만 매핑되기 때문에 페이지 테이블과 같은 메모리 공유는 불가능하다. 공유메모리 구현이 어렵다는 단점이 있다.





# Inverted Page Table Architecture





# Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
  - Total physical memory space of processes can exceed physical memory
- **Backing store** – **fast disk** large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is **transfer time**; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Standard swapping = 프로세스 전체를 스와핑

Swapping = standard swapping

Paging = swapping with paging (페이지 단위로 스와핑)





# Swapping (Cont.)

- Does the swapped-out process need to swap back into **same physical addresses**?
- Depends on **address binding** method
  - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
  - Swapping **normally disabled**
  - Started if more than **threshold** amount of memory allocated
  - **Disabled again** once memory demand reduced below threshold


일반적으로 반드시 동일할 필요는 없으나, I/O 버퍼와 같이 주소 바인딩에 따라 요구될 수도 있다.







# Schematic View of Swapping





# Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to **swap out** a process and **swap in** target process
- Context switch time can then be **very high**
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
  - Swap out time of 2000 ms
  - Plus swap in of same sized process
  - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
  - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

스와핑에 필요한 실제 메모리 용량을 OS에 알려주면 불필요한 스와핑을 줄일 수 있다.





# Context Switch Time and Swapping (Cont.)

- Other constraints as well on swapping
  - **Pending I/O** – can't swap out as I/O would occur to wrong process
  - Or always transfer I/O to kernel space, then to I/O device
    - ▶ Known as **double buffering**, adds overhead
- **Standard swapping not used** in modern operating systems
  - But modified version common
    - ▶ Swap only when free memory **extremely low**

user buffer ↔ system buffer ↔ I/O device





# Swapping on Mobile Systems

## ■ Not typically supported

- Flash memory based

- ▶ Small amount of space
- ▶ Limited number of write cycles
- ▶ Poor throughput between flash memory and CPU on mobile platform

SSD는 overwrite 불가,  
별도의 erase cycle 필요

## ■ Instead use other methods to free memory if low

- iOS **asks** apps to voluntarily relinquish allocated memory

- ▶ Read-only data thrown out and reloaded from flash if needed
- ▶ Failure to free can result in termination

- Android terminates apps if low free memory, but first writes application state to flash for fast restart

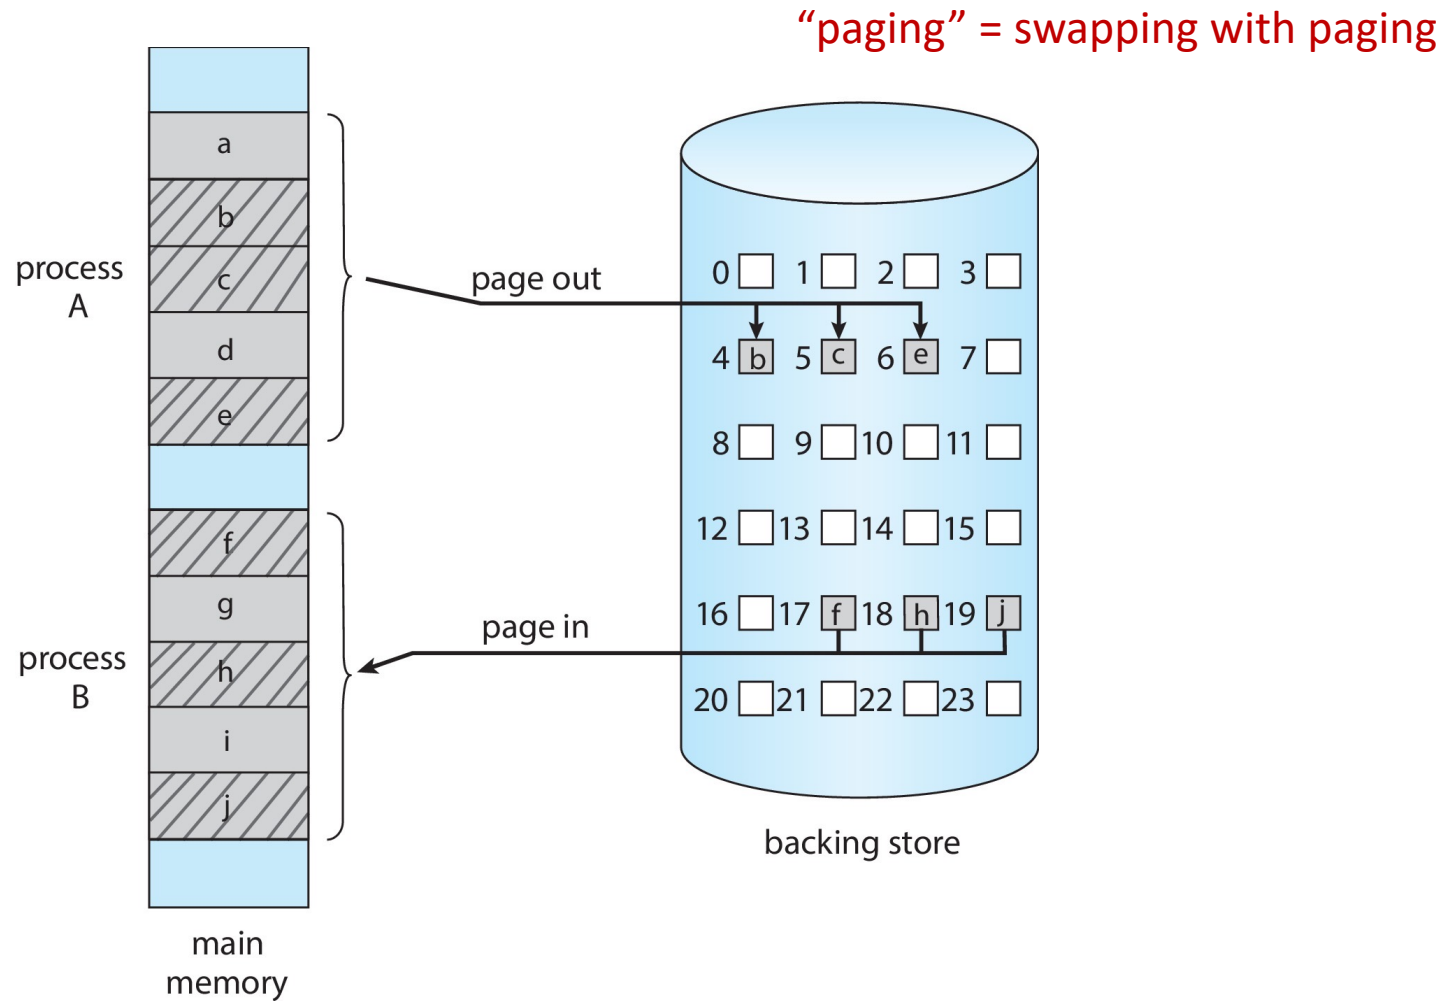
- Both OSes support paging as discussed later

- Use memory compression





# Swapping with Paging





## Example: The Intel 32 and 64-bit Architectures

---

- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
- Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here





# Example: The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging
  - Each segment can be 4 GB
  - Up to 16 K segments per process
  - Divided into two partitions
    - ▶ First partition of up to 8 K segments are private to process (kept in local descriptor table (LDT))
    - ▶ Second partition of up to 8K segments shared among all processes (kept in global descriptor table (GDT))

LTD, GTD의 각 엔트리는 8 바이트 segment descriptor로 구성, 여기에는 segment의 base location과 limit 정보가 포함됨






## Example: The Intel IA-32 Architecture (Cont.)

- CPU generates **logical address** = (selector, offset) 16 bits    32 bits

- Selector given to segmentation unit
  - ▶ Which produces linear addresses



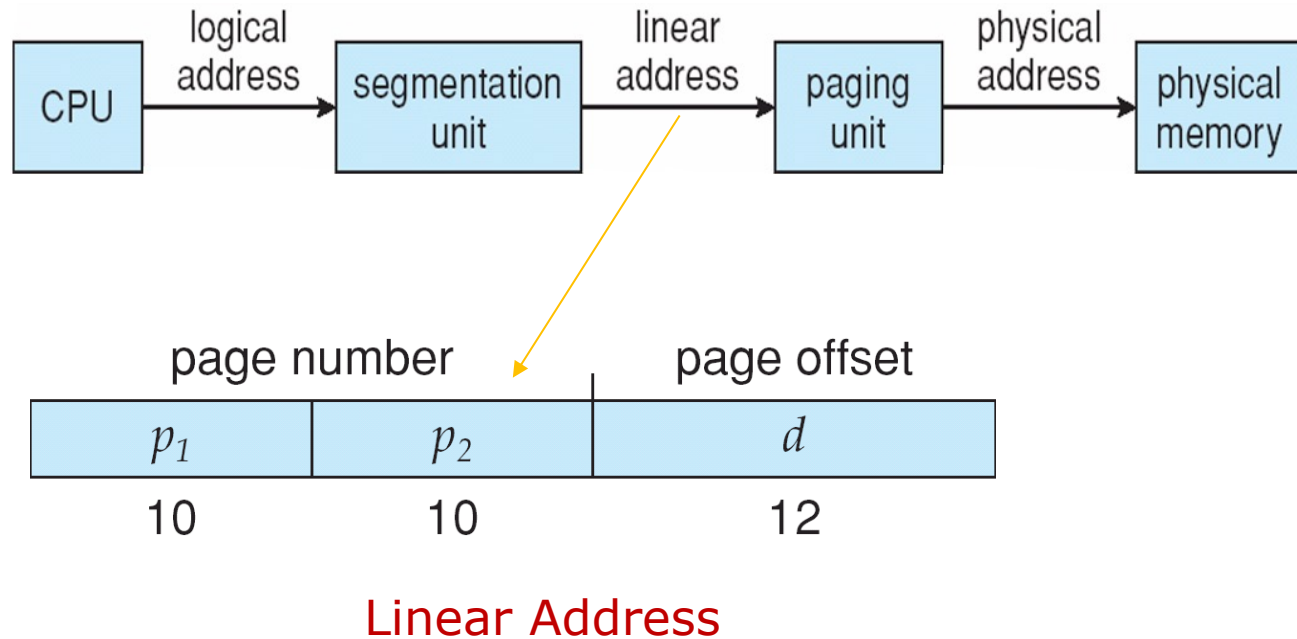
- **Linear address** given to paging unit
  - ▶ Which generates physical address in main memory
  - ▶ Paging units form equivalent of MMU
  - ▶ Pages sizes can be 4 KB or 4 MB





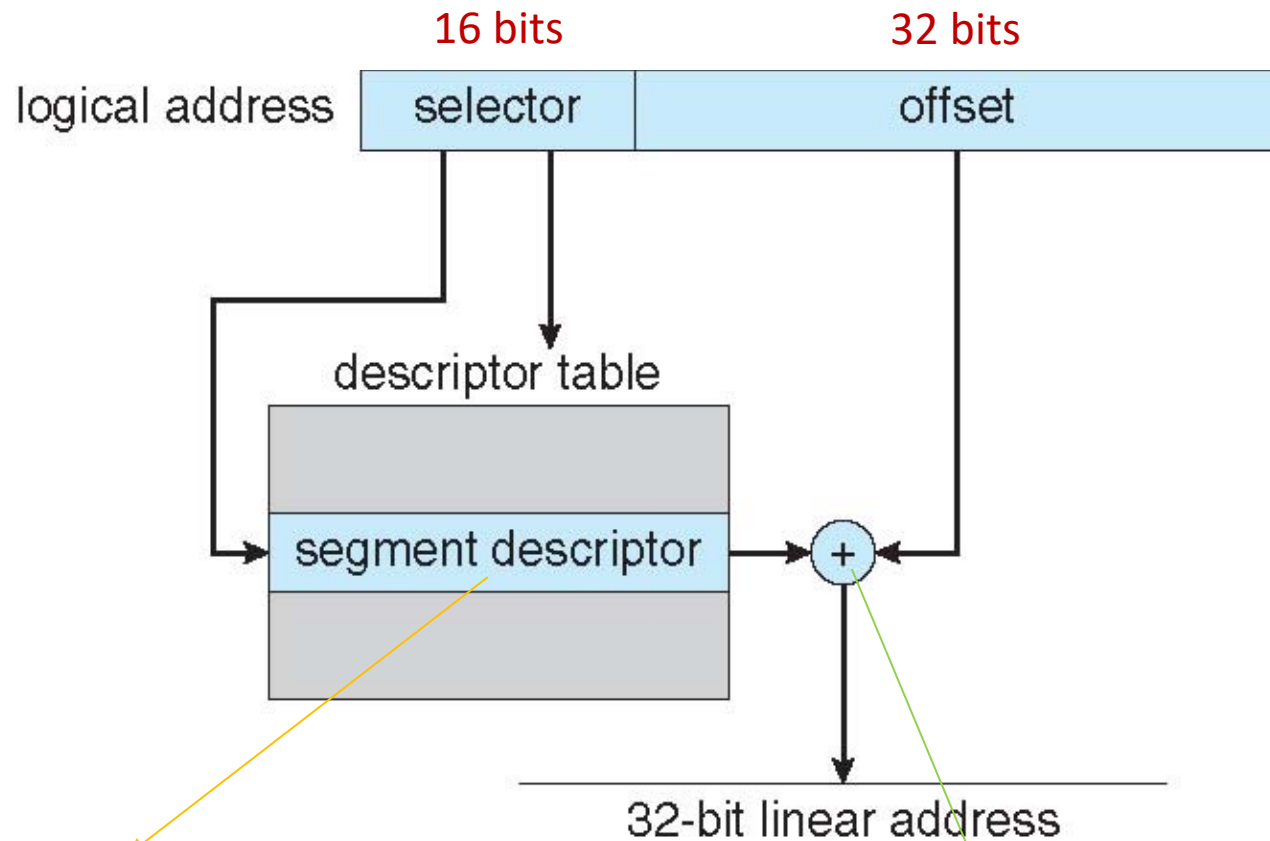


# Logical to Physical Address Translation in IA-32





# Intel IA-32 Segmentation



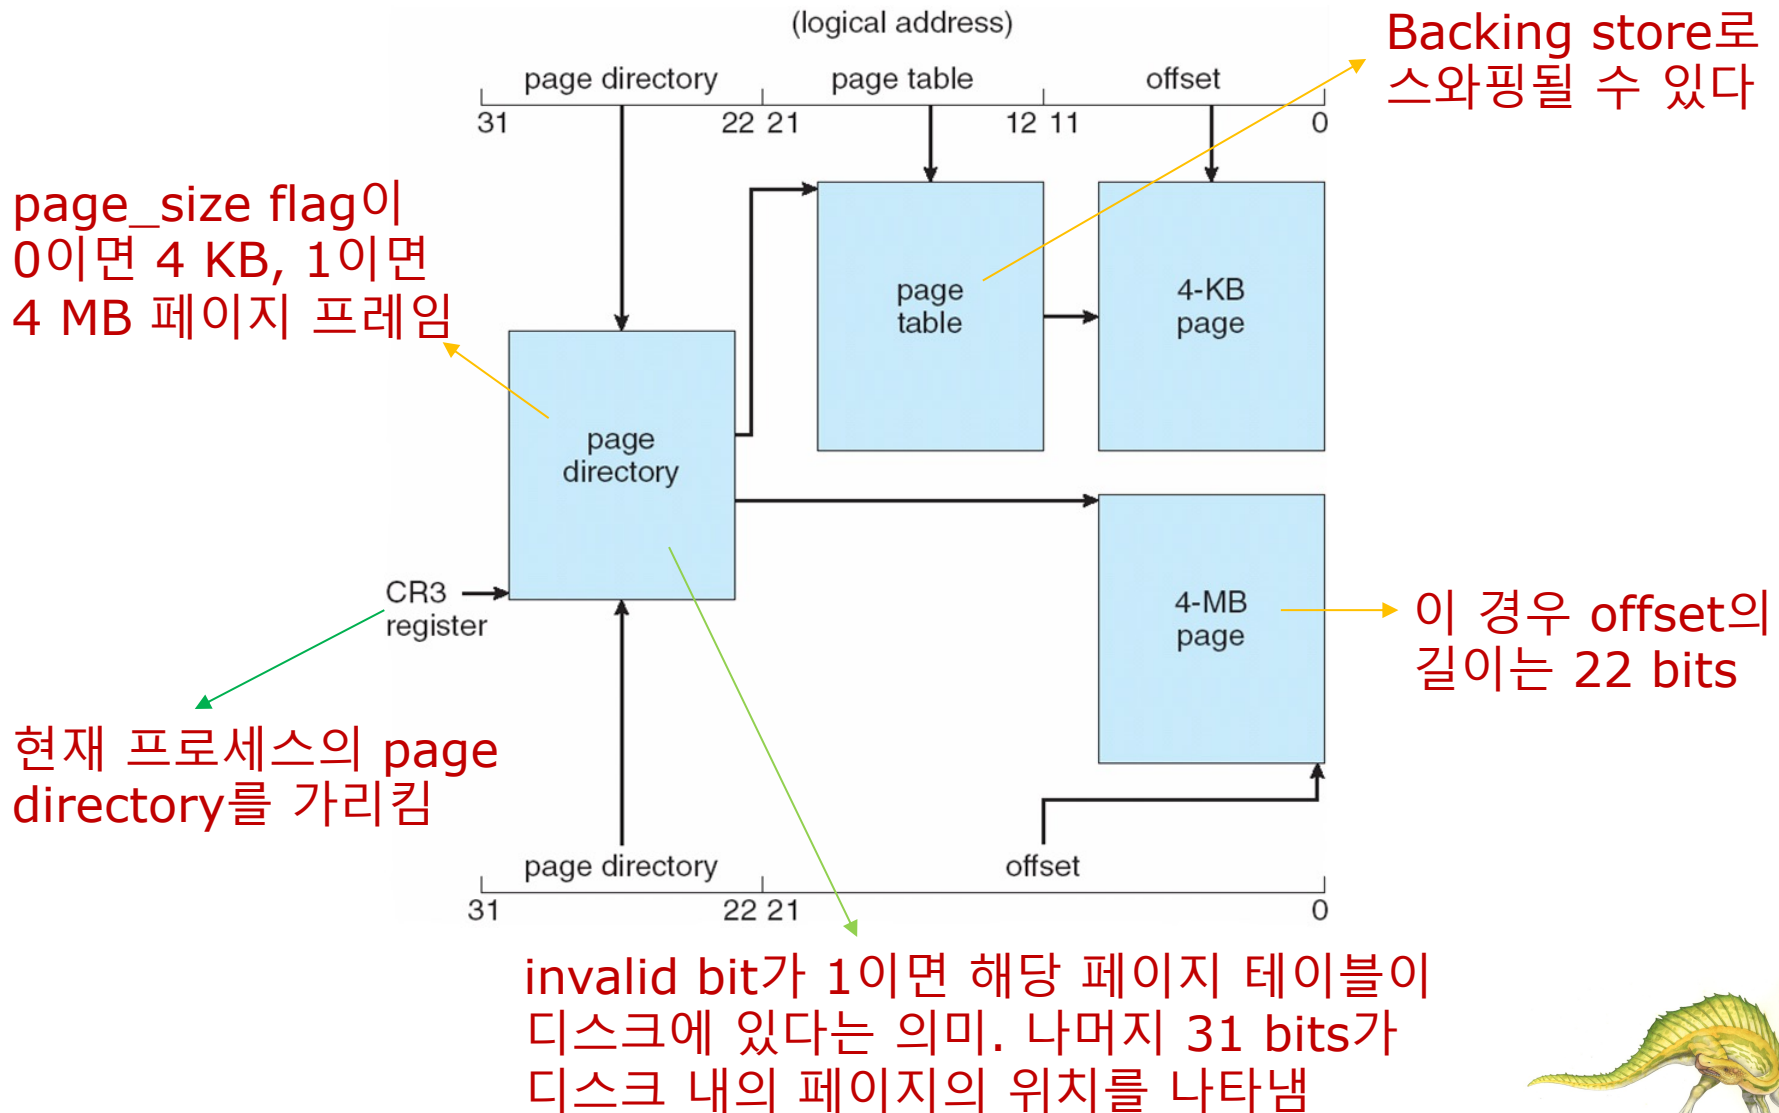
8 bytes로 구성. Segment의  
base location과 limit 정보 포함

offset이 limit 내에 있으면  
base location을 더해서 32-bit  
linear address 생성





# Intel IA-32 Paging Architecture






# Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to **more than 4GB** of memory space
  - Paging went to a **3-level** scheme
  - Top two bits refer to a **page directory pointer table**
  - Page-directory and page-table **entries** moved to **64-bits** in size
  - Net effect is increasing address space to 36 bits – **64GB** of physical memory

논리적인 메모리 공간은 여전히  $2^{32}$ 이지만 물리적인 메모리는  $2^{36}$ 까지 가능해서 메모리를 64GB까지 확장할 수 있다



PAE를 사용하려면 OS의 지원이 필요함. Linux, macOS는 PAE를 지원함. 32-bit Windows는 지원하지 않음.


각 엔트리가 32 bits에서 64 bits로 확장,  
base address를 20에서 24 bits로 확장





# Intel x86-64

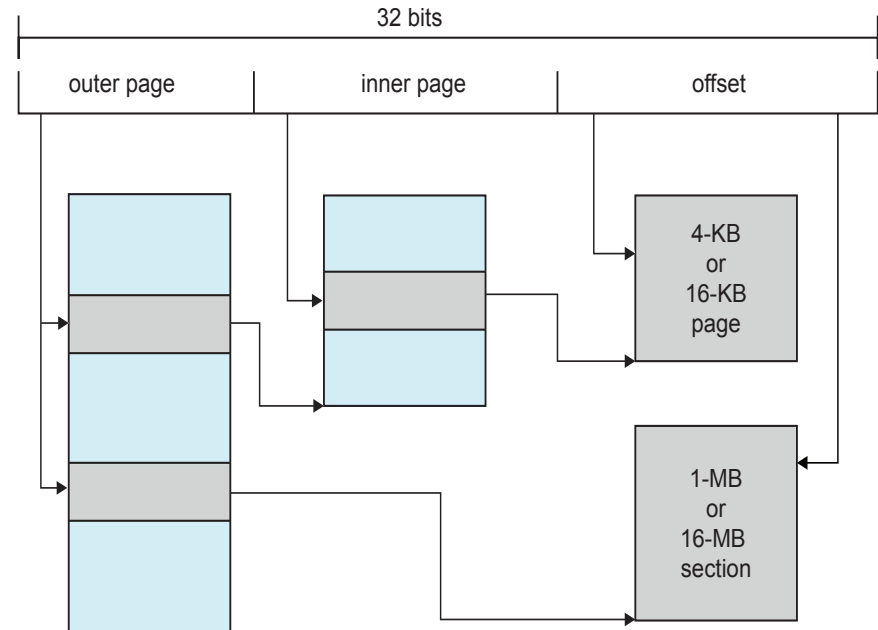
- Current generation Intel x86 architecture
- 64 bits is **ginormous** (> 16 exabytes)
- In practice only implement **48 bit** addressing
  - Page sizes of 4 KB, 2 MB, 1 GB
  - **Four levels** of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits (**256 TB**) and physical addresses are 52 bits (**4 PB**)





# Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
  - Outer level has two micro TLBs (one data, one instruction)
  - Inner is single main TLB
  - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



교과서에 있는 ARMv8 64-bit 아키텍처 참조. ARMv8은 48 비트 주소와 four-level hierarchical paging을 사용.



# End of Chapter 9

---

