

Computer Architecture (ENE1004)


Lec - 1: Computer Abstractions and Technology (Chapter 1)

What You Can Learn ...

- By the time we end this course, you should be able to answer:
 - How are programs written in a high-level languages translated into the language of the hardware? How does the hardware execute the resulting program?
 - What is the interface between the software and the hardware? How does software instruct the hardware to perform needed functions?
 - What determines the performance of a program? How can a programmer improve the performance?
- You can improve the performance of your program if you are aware of ...

Our interest	Hardware or software component	How this component affects performance	Where is this topic covered?
	Algorithm	Determines both the number of source-level statements and the number of I/O operations executed	Other books!
	Programming language, compiler, and architecture	Determines the number of computer instructions for each source-level statement	Chapters 2 and 3
	Processor and memory system	Determines how fast instructions can be executed	Chapters 4, 5, and 6
	I/O system (hardware and operating system)	Determines how fast I/O operations may be executed	Chapters 4, 5, and 6

Review: Layers of Software



- System software helps your application software to run on hardware
- Two kinds of key system software
- (1) Operation system
 - Allocating memory and storage
 - Handling input/output operations
 - Coordinating concurrent applications
- (2) Compiler
 - Translating a program written in a high-level language to instructions that the hardware can execute
 - Your program is based on a high-level language: $c = a + b$;
 - Hardware works based on binary values: 100111010

Review: from High-Level Language to Machine Language

High-level
language
program
(in C)

```
swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

- High-level language (e.g., C code)
- Text composed of words and algebraic notation

Compiler

Assembly
language
program
(for MIPS)

```
swap:
    multi $2, $5, 4
    add $2, $4, $2
    lw $15, 0($2)
    lw $16, 4($2)
    sw $16, 0($2)
    sw $15, 4($2)
    jr $31
```

- High-level language -> Assembly language
- Assembly language (e.g., MIPS code)
- A symbolic representation of machine instructions

Assembler

Binary machine
language
program
(for MIPS)

```
000000001010001000000000100011000
00000000100000100001000000100001
10001101111000100000000000000000
100011100001001000000000000000100
```

- Assembly language -> Machine language
- Machine language
- A binary representation of machine instructions

Review: from High-Level Language to Machine Language

- Regarding to the conversion, there are many terms
 - Compiling, linking, building, ... ; inclusion relationship, used interchangeably
 - Compilers can cut out the intermediate process and produce binary directly
- Recognition that a program can be written in a more powerful language and then translated into computer instructions was a great breakthrough
 - High-level languages offer several important benefits
 - It allows programmers to think in a more natural language, using English words and algebraic notation, resulting in programs that look much more like text
 - It improves programmer productivity; it takes less time to develop programs
 - It allows programs to be independent of the computer on which they were developed
 - Compiler has been evolving towards producing very efficient machine code optimized for the target machine
 - Study of compiler has a long history
 - CGO (<https://conf.researchr.org/program/cgo-2023/program-cgo-2023/>)
- Again, in this course, we are interested in how target machine executes the machine code

Example: Device Components



- Apple iPad 2
 - Display
 - Battery
 - Logic board



- Logic board of Apple iPad 2
 - A5 chip
 - Flash memory
 - Power controller chip
 - I/O controller chip

Example: Device Components



- Inside the A5 chip
 - Two ARM processors
 - GPU with four datapaths
 - DRAM interfaces
 - I/O controller chip
- Integrated circuit (chip, package)
- Processors (a.k.a., central processor unit, CPU)
 - (1) Datapath: performs arithmetic operations
 - (2) Control: tells datapath, memory, I/O device what to do
 - Chapter 4 will cover the two components in detail
- Interfaces & controllers
 - DRAM, I/O, etc
- Cache memory (SRAM)
 - Temporarily keeps data

Example: Device Components



H I E R A R C H Y

- Storage
 - Non-volatile: it keeps data after power loss
 - Hard disk drives (HDDs) or Solid state drives (SSDs) based on flash memory
 - Slow, low cost per bit
- Main memory
 - Volatile: data is lost when power is removed
 - Dynamic random access memory (DRAM)
 - Faster than storage, high cost per bit
- Cache memory
 - Volatile as well
 - Static random access memory (SRAM)
 - Faster than DRAM, more expensive than DRAM
- Chapter 5 will cover the memory systems in detail

Review: Two Performance Metrics

- We say one computer has better performance than another; what do we mean?
 - Performance of airplanes (cruising speed vs passenger throughput)

Airplane	Passenger capacity	Cruising range (miles)	Cruising speed (m.p.h.)	Passenger throughput (passengers × m.p.h.)
Boeing 777	375	4630	610	228,750
Boeing 747	470	4150	610	286,700
BAC/Sud Concorde	132	4000	1350	178,200
Douglas DC-8-50	146	8720	544	79,424

- Response time (execution time) - s, ms, us, ns, ...
 - The time between the start and the completion of a task
 - Important to individual users
 - Used in evaluating the performance of embedded and laptop computers
- Throughput (bandwidth) - MB/s, GB/s, ...
 - The total amount of work done in a given time
 - Important to datacenter managers
 - Used in evaluating the performance of servers

Review: Measuring Execution Time

- Relative Performance
 - If computer A runs a program in 10 seconds and computer B runs the same program in 15 seconds, how much faster is A than B?
 - A is n times faster than B if $\text{execution_time}_B / \text{execution_time}_A = n$
 - $15/10 = 1.5$; so, A is 1.5 times faster than B
- Two types of execution time
 - (1) Elapsed time
 - The total time to complete a task
 - It includes disk accesses, memory accesses, I/O activities, operating system overhead
 - (2) CPU time
 - The time the CPU spends computing for this task
 - It does not include time spent for other activities
- Once you can analyze both the elapsed time & CPU time of a task, you can find the performance bottleneck and thus where to optimize