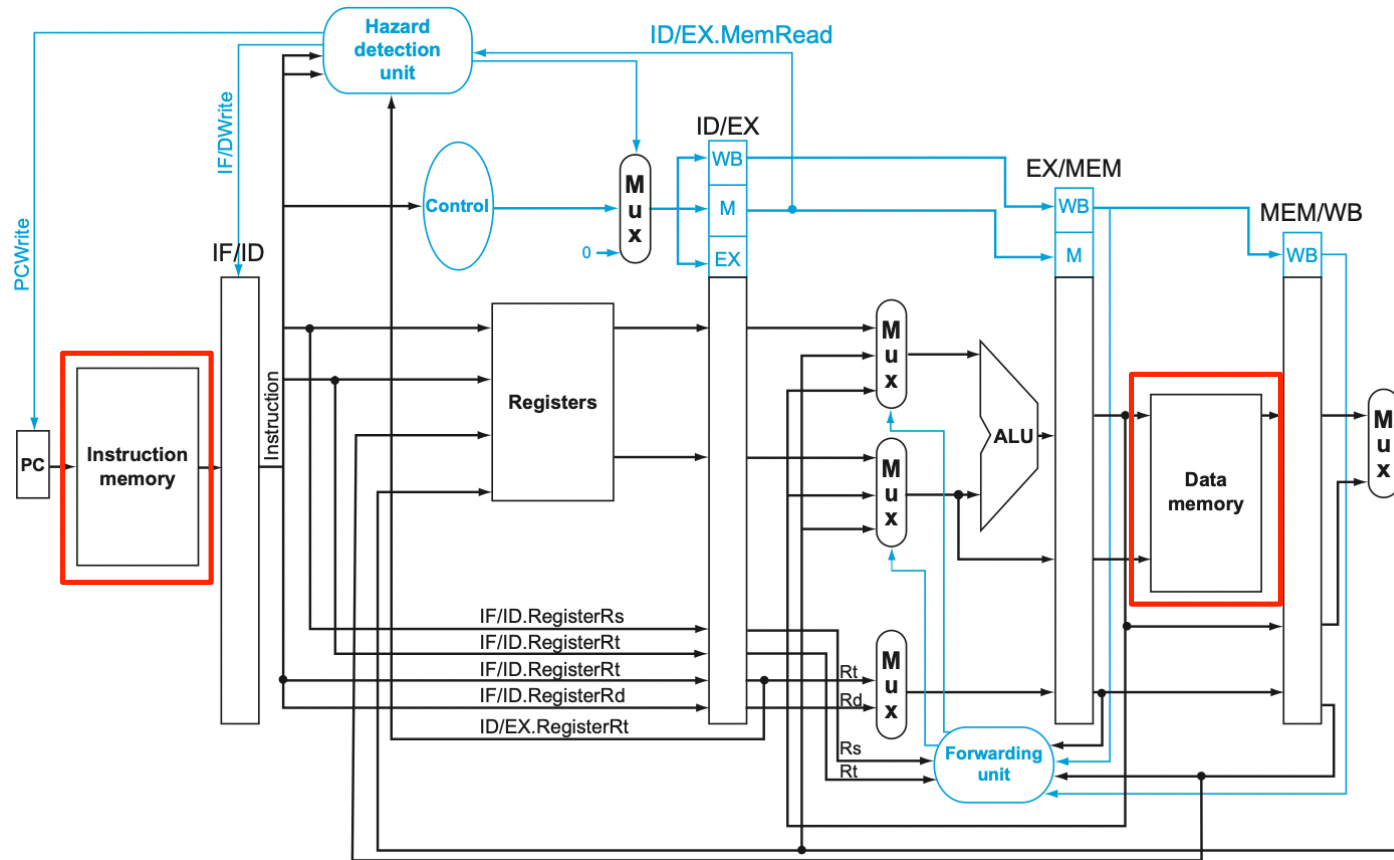


Computer Architecture (ENE1004)

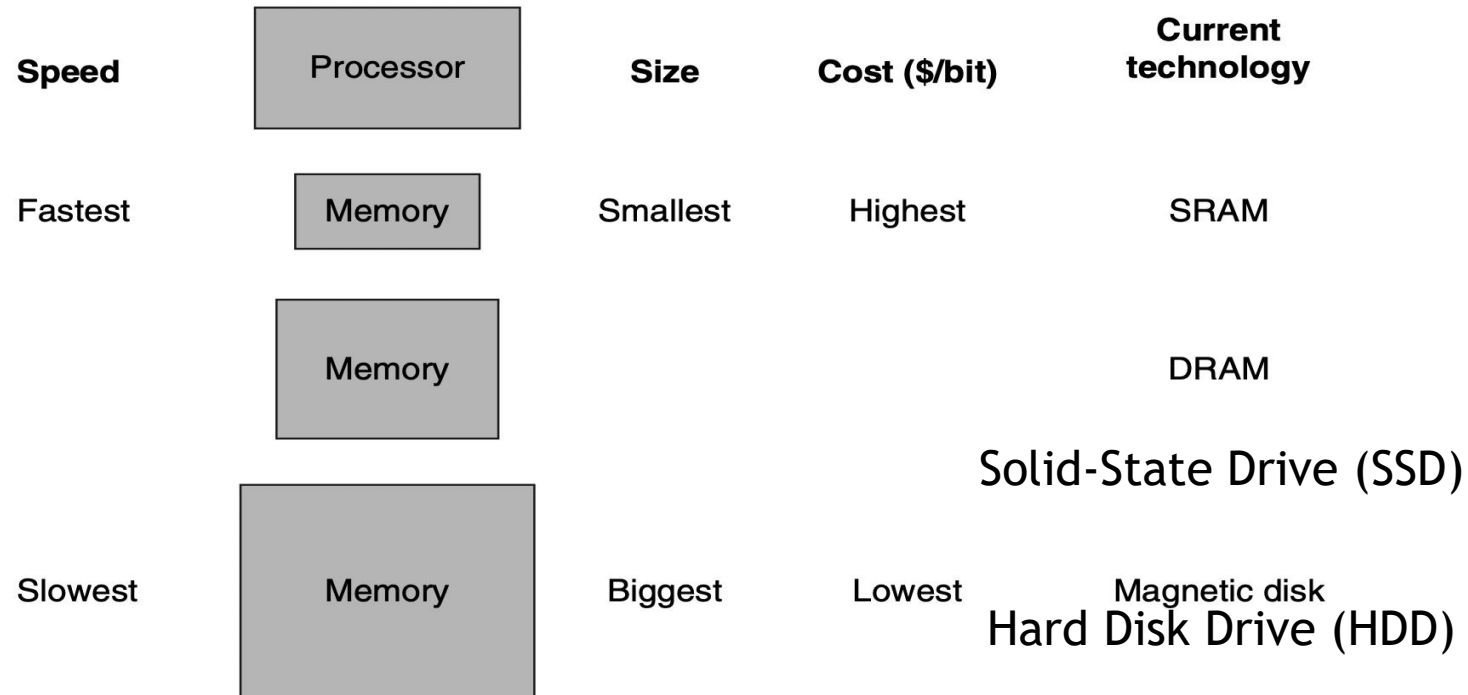
Lec - 19: Large and Fast: Exploiting Memory Hierarchy (Chapter 5) - 1

Ideal Memory System



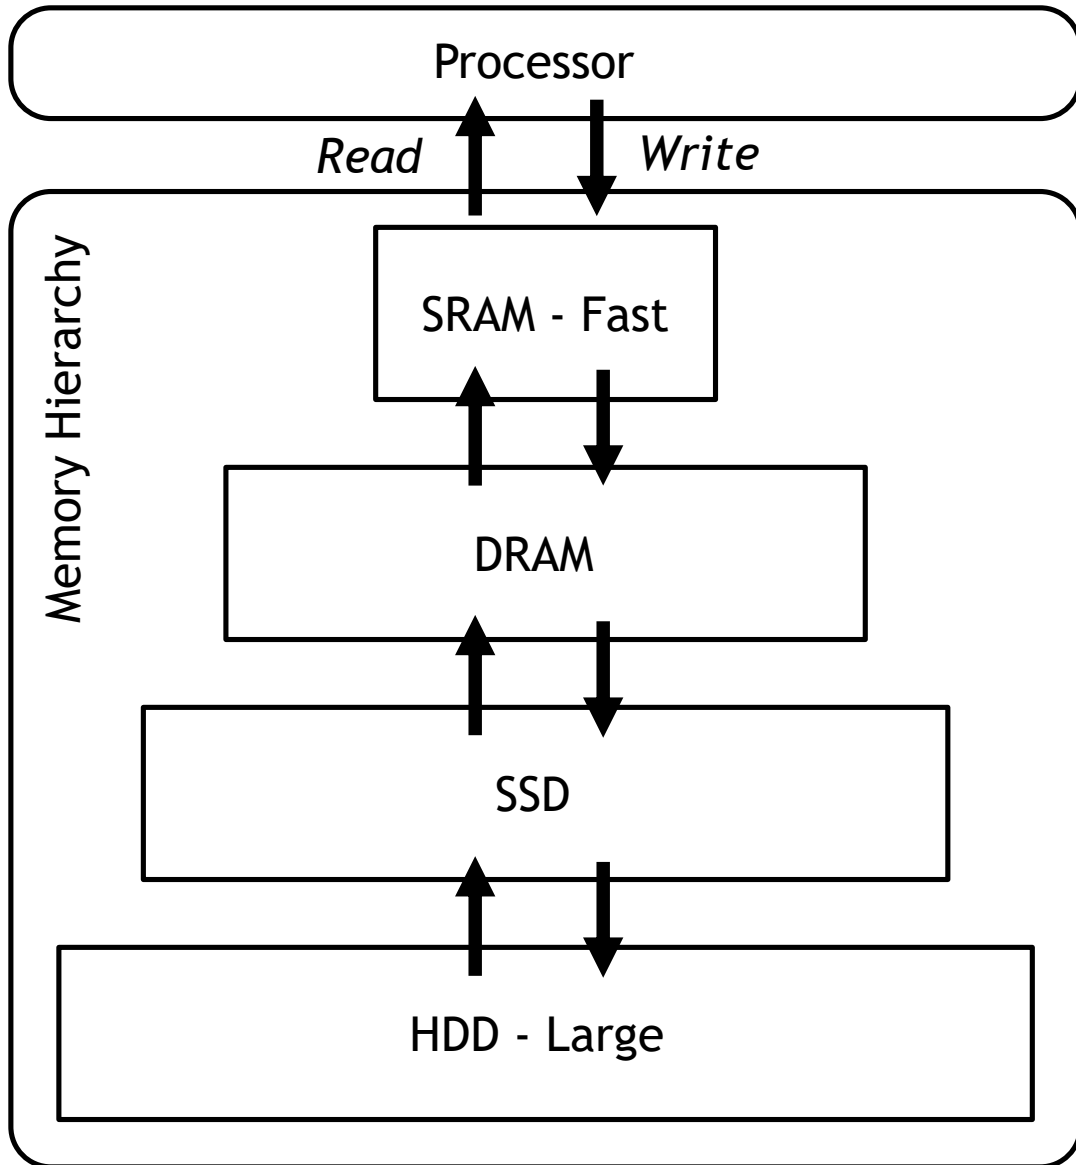
- All the data (program, data) are kept in the memory
 - Load and store instructions access (read from and write into) the memory
- From the earliest days, programmers have wanted “unlimited amounts of fast memory”
 - Fast memory (performance angle)
 - Large memory (capacity angle)

Current Memory Technologies



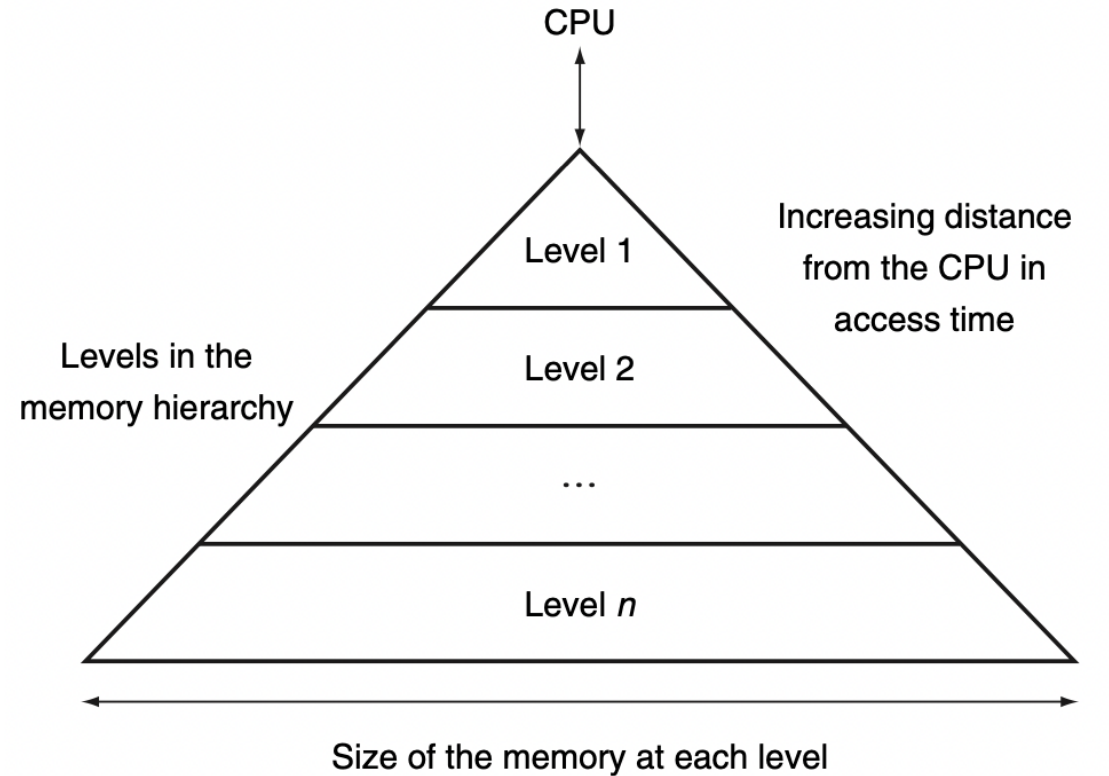
- Comparison: SRAM vs DRAM vs Solid-State Drive (SSD) vs Hard Disk Drive (HDD)
 - Cost (\$/bit) comparison: SRAM > DRAM > SSD > HDD
 - Performance (read/write speed) comparison: SRAM > DRAM > SSD > HDD
 - Capacity (size) comparison: SRAM < DRAM < SSD < HDD
- There is NO single memory technology that is both large and fast
 - However, modern computers provide an “illusion” of the large and fast memory

Memory Hierarchy (1)



- Memory system includes different memory technologies in a *hierarchical* fashion
 - At the lowest level, the largest/slowest memory
 - At the first level, the fastest/smallest memory
 - As the distance from the processor increases, both the size and the access time increase
- The data is similarly hierarchical
 - All the data is stored at the lowest level
 - A level closer to the processor is generally a “subset” of any level further away
 - The subset can change over time by sending data upwards/downwards between any two level
 - The smallest subset is stored at the first level
 - Processor can directly access only the subset data at the first level

Memory Hierarchy (2)

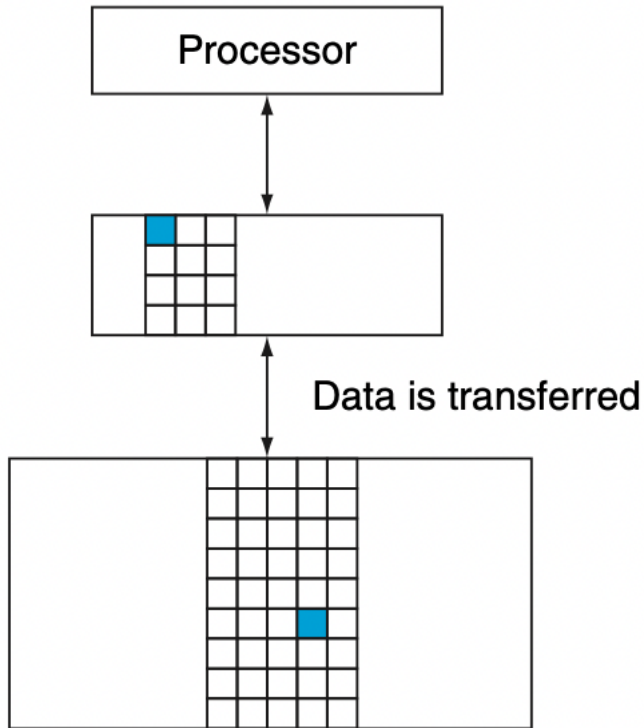


- Processor accesses only the level-1 memory
 - Processor *feels* like it works with the **fastest** memory
- Processor can access all the data set in the level-n memory
 - Any data can be sent upwards to the level-1 memory from level-n memory
 - Processor *feels* like it works with the **largest** memory

Memory Hierarchy (3)


- A level closer to the processor is generally a “subset” of any level further away
- This concept is based on the principle of locality in the memory accesses
 - Two different types of locality are observed in real-world applications
- Temporal locality (locality in time)
 - If a data item is referenced (accessed), it will tend to be referenced again soon
 - So, if such a data item is stored in an upper level, we can reduce the access times to it
 - In a for loop, the variable for iterations is referenced again in a short period of time
- Spatial locality (locality in space)
 - If an item is referenced, items whose addresses are close by will tend to be referenced soon
 - So, if data items in the neighborhood are stored in an upper level, we can reduce the access times to them
 - In an array, neighboring elements are usually referenced together in a short period of time

In Every Pair of Levels




- Every pair of levels in the hierarchy can be thought of as having an upper and lower level
 - Data is copied between only two adjacent levels at a time
- The minimum unit of information that can be either present or not present in the two levels is called “**block**” or “**line**”
- If the data requested by the processor appears in some block in the upper level, the request is called a “**hit**”
 - It can be serviced quickly from the upper-level faster memory
 - The time to access the upper level memory is “**hit time**”
- If the data is not found in the upper level, it is called “**miss**”
 - The data is copied from the lower level to the upper level; then is read
 - The total time is “**miss penalty**”
- “**Hit ratio**” is the fraction of accesses found in the upper level
 - The higher the hit ratio, the higher the performance
 - “**Miss ratio**” = 1 - “hit ratio”

Basics of Caches (1)



- In this course, we'll focus on the most upper level two memory technologies
 - SRAM in upper level is called “**cache**”
 - DRAM in lower level is called “main memory”
 - Processor always works with the cache
 - Data set in the cache changes by transferring data between the cache and the main memory

Basics of Caches (2)




a. Before the reference to X_n

b. After the reference to X_n


- An example of simple cache
 - The cache can hold up to 8 data items; currently, it holds 6 data items and two slots are empty
 - Given a new access to one of the six items, the data is provided from the cache to the processor
 - Given a new item X_n , it is brought from the main memory and accommodated in the cache
- Here, one may have the following important questions
 - (1) How do we know if a data item is in the cache?
 - (2) If it is, how do we find it?

Direct Mapped Cache (1)



- Answer: If each data item can go in exactly one place in the cache, then it is straightforward to find the data item if it is in the cache
- **Direct-mapped cache**: each memory location is mapped directly to **one location** in the cache
 - It assigns a location in the cache for each data item
 - The assignment of a cache location **based on the address of the data** item in memory
- Here, a cache location is assigned to multiple data items
 - 8 cache blocks should be assigned to 32 data items
 - 4 different data items can be placed in a single cache block


Direct Mapped Cache (2) - Cache Location



- Modulo operation
 - $7 \text{ module } 3 = 1$
 - $5 \text{ module } 3 = 2$
 - $10 \text{ module } 3 = 1$

- One simple way to assign a cache location based on the address of the data item
 - (Block address) modulo (Number of blocks in the cache)
- In the example, the number of blocks in the cache is 8 (2^3)
 - Memory address 1_{ten} (00001_{two}) is mapped to cache block 1_{ten} (001_{two}); $1 \text{ module } 8 = 1$
 - Memory addresses 9_{ten} (01001_{two}), 17_{ten} (10001_{two}), and 25_{ten} (11001_{two}) are also mapped to cache block 1_{ten} (001_{two}); $9 \text{ module } 8 = 1$, $17 \text{ module } 8 = 1$, and $25 \text{ module } 8 = 1$
 - Memory addresses 5_{ten} (00101_{two}), 13_{ten} (01101_{two}), 21_{ten} (10101_{two}), and 29_{ten} (11101_{two}) are all mapped to cache block 5_{ten} (101_{two})
- Simply, check the last **3** bits of the memory address; the size of cache (2^3)!

Direct Mapped Cache (3) - Tag




Data whose addresses are 00001, 01001, 10001, and 11001 can be placed in cache @ 001; indices are the same!

To identify one of the four, tags (00, 01, 10, 11) can be used; tags are all different!

- A cache location can contain the data items of different memory locations
 - Question: How do we know whether the requested data corresponds to the data in the cache?
 - We can answer the question by adding “tag” to each cache location
- Memory address is divided and used for two different purposes
 - The lower bits are used to identify the cache location (called “index”)
 - The upper bits are used to identify a specific data (called “tag”)
 - (b4 b3 b2 b1 b0) is divided; the lower 3 bits are “index”, which is used for a cache location; the upper 2 bits are “tag”, which is used for a specific data

Direct Mapped Cache (4) - Valid Bit



- A cache block does **not** always hold a valid data
 - The cache block can be empty, if none of the four data exist in the cache
 - When a processor starts up, the cache does not have any data
- We can address the above by adding a “**valid bit**” to each cache location
 - If it is set, the cache location has a valid data; then, refer to the tag if it is the data what you want
 - If it is not set, you do not have to check the tag, since any information in the cache location is not valid