

Computer Architecture

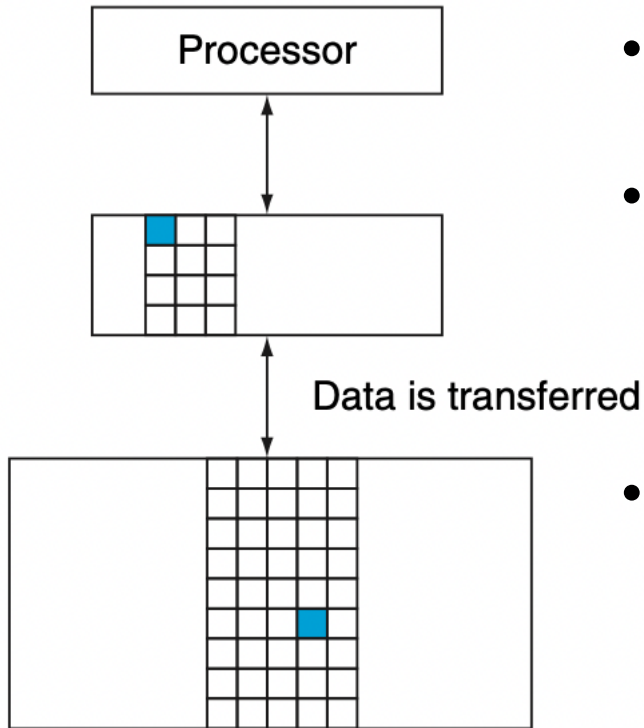
(ENE1004)

Lec - 20: Large and Fast: Exploiting Memory Hierarchy (Chapter 5) - 2

Schedule

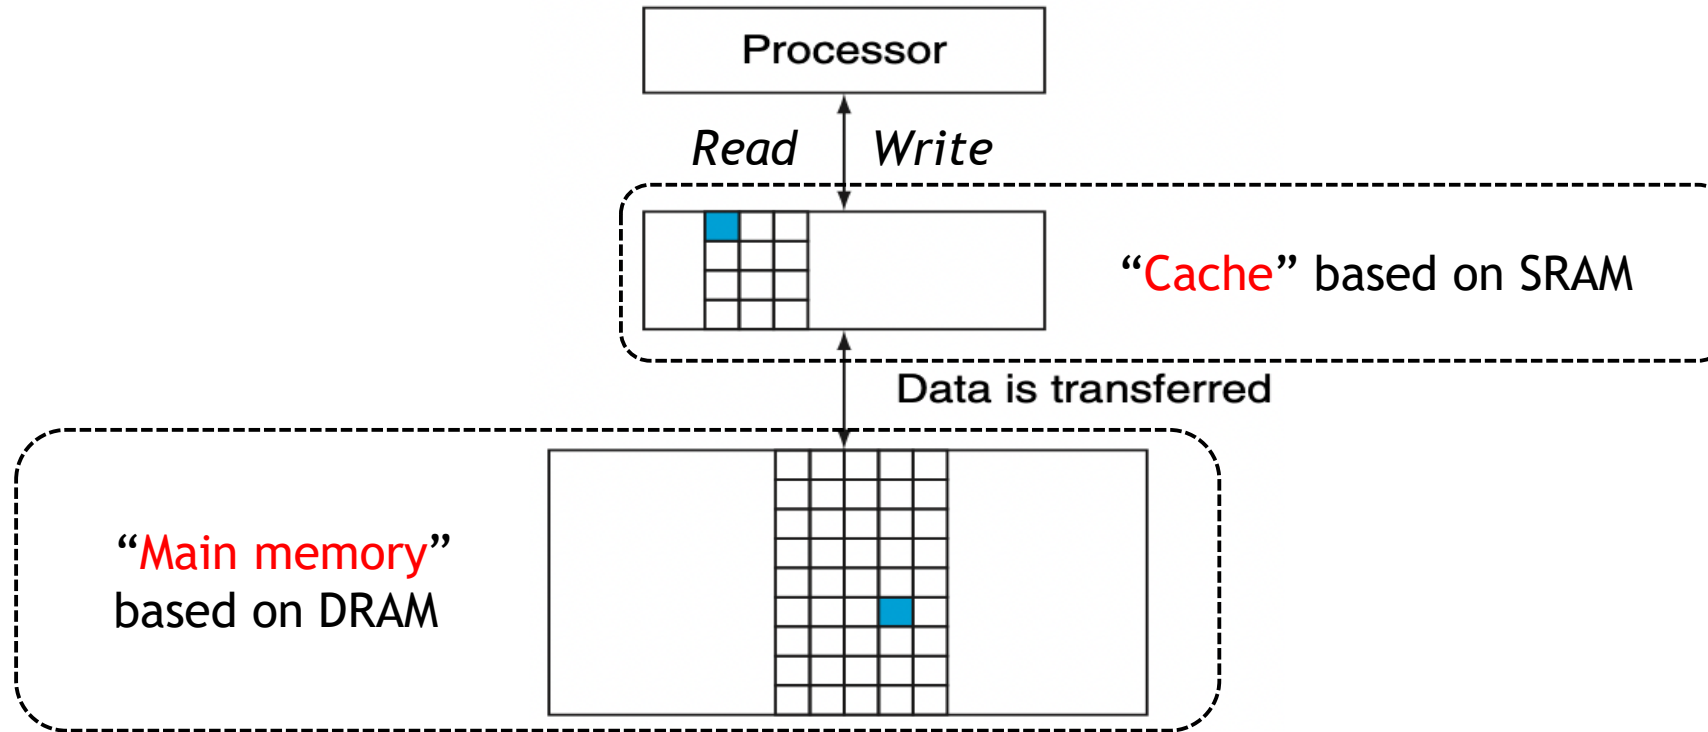
- Final exam: Jun. 19, Monday
 - (24334) 1:25~2:25pm
 - (23978) 2:30~3:30pm
 - Sample questions will be provided by Jun. 17 (Saturday)
- Assignment #2: Jun. 20, Tue. by midnight
 - It is put back as much as possible
- Remaining class days
 - 1(Thur), 5(Mon), 8(Thur), 12(Mon), 15 (Thur)

Memory Terminology



- Every pair of levels in the hierarchy can be thought of as having an upper and lower level
 - Data is copied between only two adjacent levels at a time
- “**Block (line)**” - The minimum unit of data item that can be either present or not present in the two levels is called “block”
- “**Hit**” - If the data requested by the processor appears in some block in the upper level, the request is called a “hit”
 - It can be serviced quickly from the upper-level faster memory
- “**Hit time**” - The time to access the upper level memory
- “**Miss**” - If the data is not found in the upper level (instead, it can be found in the lower level), it is called “miss”
 - The data is copied from the lower level to the upper level; then is read
 - The total time is “**miss penalty**”
- “**Hit ratio**” is the fraction of accesses found in the upper level
 - The higher the hit ratio, the higher the performance
 - “**Miss ratio**” = $1 - \text{“hit ratio”}$


Basics of Caches (1)



- In this course, we'll focus on the most upper level two memory technologies
 - SRAM in upper level is called "cache"
 - DRAM in lower level is called "main memory"
 - Processor always works with the cache
 - Data set in the cache changes by transferring data between the cache and the main memory
- Specifically, we'll see how a cache is designed and managed

Basics of Caches (2)

“Cache” with 8 blocks,
which can hold up to 8
data items at a time




a. Before the reference to X_n

b. After the reference to X_n

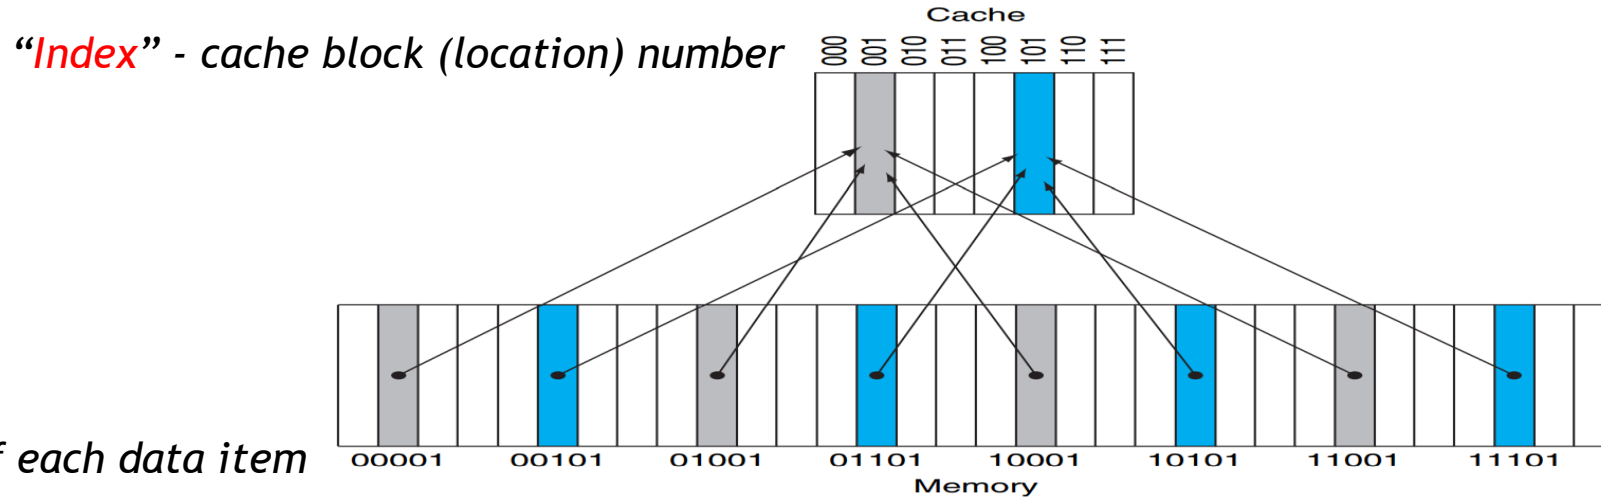
- An example of simple cache
 - The cache consists of 8 blocks (8 lines); currently, it holds 6 data items and two blocks are empty
 - Given a request to one of the 6 data items, it is “hit” and the item is provided from the cache
 - Given a request to data item (X_n), which does not exist in the cache, it is “miss” and the data item is brought from the main memory into the cache
- Here, one may have the following important questions
 - (1) How do we know if a data item is in the cache?
 - (2) If it is, how do we find it?

Direct Mapped Cache (1)



- One answer: If each data item can go in **exactly one place** in the cache, then it is straightforward to find the data item if it is in the cache
 - Given a request, we can just check the designated single place instead of scanning all the places
- **Direct-mapped cache**: each memory location is mapped directly to **one location** in the cache
 - It assigns a location in the cache to each data item
 - The assignment of a cache location is **based on “the address of the data item” in memory**
- Note that multiple data items should share a cache location
 - 8 cache blocks are assigned to 32 data items; 4 different data items should share a cache block
 - If a data item uses the block, the other three items cannot use the same block


Direct Mapped Cache (2) - Index



- Modulo operation
 - $7 \text{ module } 3 = 1$
 - $5 \text{ module } 3 = 2$
 - $10 \text{ module } 3 = 1$

- One simple way to assign a cache block is to use the **address** of the data item
 - We call the cache block number "**index**"
 - **Index = (Address) modulo (Number of blocks in the cache)**
- In the example, the number of blocks in the cache is 8 (2^3)
 - Memory address 1_{ten} (00001_{two}) is mapped to cache block 1_{ten} (001_{two}); $1 \text{ module } 8 = 1$
 - Memory addresses 9_{ten} (01001_{two}), 17_{ten} (10001_{two}), and 25_{ten} (11001_{two}) are also mapped to cache block 1_{ten} (001_{two}); $9 \text{ module } 8 = 1$, $17 \text{ module } 8 = 1$, and $25 \text{ module } 8 = 1$
 - Memory addresses 5_{ten} (00101_{two}), 13_{ten} (01101_{two}), 21_{ten} (10101_{two}), and 29_{ten} (11101_{two}) are all mapped to cache block 5_{ten} (101_{two})
- Simply, check the last **3** bits of the memory address; the size of cache (2^3)!
 - $00\mathbf{001}_{\text{two}}$, $01\mathbf{001}_{\text{two}}$, $10\mathbf{001}_{\text{two}}$, $11\mathbf{001}_{\text{two}}$; the index of these four data items is $\mathbf{001}_{\text{two}}$
 - $00\mathbf{101}_{\text{two}}$, $01\mathbf{101}_{\text{two}}$, $10\mathbf{101}_{\text{two}}$, $11\mathbf{101}_{\text{two}}$; the index of these four data items is $\mathbf{101}_{\text{two}}$

Direct Mapped Cache (3) - Tag




Data whose addresses are **00001**, **01001**, **10001**, and **11001** can be placed in cache block **001**; they have the same index!

To identify one of the four, tags (**00**, **01**, **10**, **11**) can be used; tags are all different!

- A cache block can contain the different data items
 - **Question: How do we know whether the requested data corresponds to the data in the cache?**
 - We can answer the question by adding “**tag**” to each cache block
- The address of data is divided and used for two different purposes
 - The lower bits are used for “index” (to identify its cache block)
 - **The upper bits are used for “tag”** (to identify the specific data)
 - (**b4 b3 b2 b1 b0**) is divided; the lower 3 bits are “**index**”, which is used for a cache block; the upper 2 bits are “**tag**”, which is used to indicate the specific data

Direct Mapped Cache (4) - Valid Bit



- A cache block does **not** always hold a valid data
 - The cache block can be empty, if none of the four data exist in the cache
 - When a processor starts up, the cache does not have any data
- We can address the above by adding a “**valid bit**” to each cache block
 - If it is set, the cache block has a valid data; then, you need to check the tag whether the valid data is what you want or not
 - If it is not set, you do not have to check the tag, since any information in the cache block is not valid

Access a Direct Mapped Cache (1)

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

- Each entry of the cache consists of “index” + “valid bit” + “tag” + “data”
 - There are 8 entries (blocks) in the cache, each of which is identified based on the 3-bit “index”
 - If the 1-bit “valid bit” tells whether the information in the entry is valid or not
 - If a data is placed in its corresponding entry, its 2-bit “tag” is set to indicate the specific data
- At first, the cache is empty (it does not hold any data) right after power-on
 - The valid bit of each entry is set to “N”
 - The tag and data are also empty

Access a Direct Mapped Cache (2)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (5.6b)	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	miss (5.6c)	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
22	10110_{two}	hit	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	hit	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	miss (5.6d)	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
3	00011_{two}	miss (5.6e)	$(00\mathbf{011}_{\text{two}} \bmod 8) = \mathbf{011}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
18	10010_{two}	miss (5.6f)	$(10\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

b. After handling a miss of address (10110_{two})

- An example scenario where a series of requests (references) are given to the cache
 - If the referenced data is in the cache, it is “hit” and the data is serviced from the cache
 - If the referenced data is not in the cache, it is “miss” and the existing data is replaced with the referenced data, which is copied from the main memory
- Data 22 ($10\mathbf{110}_{\text{two}}$) is requested
 - It can be placed in the cache block whose “index” is 110_{two} (the lower 3 bits $10\mathbf{110}_{\text{two}}$)
 - It is “miss” because the valid bit is set to N (Data 22 does not exist in the cache block)
 - Data 22 is brought from the main memory and stored in the cache block whose index is 110_{two}
 - Tag is set to 10_{two} to indicate data 22 (the upper 2 bits of $10\mathbf{110}_{\text{two}}$); valid bit is set to Y

Access a Direct Mapped Cache (3)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110 _{two}	miss (5.6b)	(10110 _{two} mod 8) = 110 _{two}
26	11010 _{two}	miss (5.6c)	(11010 _{two} mod 8) = 010 _{two}
22	10110 _{two}	hit	(10110 _{two} mod 8) = 110 _{two}
26	11010 _{two}	hit	(11010 _{two} mod 8) = 010 _{two}
16	10000 _{two}	miss (5.6d)	(10000 _{two} mod 8) = 000 _{two}
3	00011 _{two}	miss (5.6e)	(00011 _{two} mod 8) = 011 _{two}
16	10000 _{two}	hit	(10000 _{two} mod 8) = 000 _{two}
18	10010 _{two}	miss (5.6f)	(10010 _{two} mod 8) = 010 _{two}
16	10000 _{two}	hit	(10000 _{two} mod 8) = 000 _{two}

Index	V	Tag	Data
000	N		
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

c. After handling a miss of address (11010_{two})

- Data 26 (11010_{two}) is requested
 - It can be placed in the cache block whose “index” is 010_{two} (the lower 3 bits 11010_{two})
 - It is “miss” because the valid bit is set to N (Data 26 does not exist in the cache block)
 - Data 26 is brought from the main memory and stored in the cache block whose index is 010_{two}
 - Tag is set to 11_{two} to indicate data 26 (the upper 2 bits of 11010_{two}); valid bit is set to Y
- Data 22 (10110_{two}) is requested again
 - Go to the index 110_{two}; first, check its valid bit Y; then, check the tag 10_{two}; it is “hit”
- Data 26 (11010_{two}) is requested again
 - Go to the index 010_{two}; first, check its valid bit Y; then, check the tag 11_{two}; it is “hit”

Access a Direct Mapped Cache (4)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (5.6b)	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	miss (5.6c)	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
22	10110_{two}	hit	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	hit	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	miss (5.6d)	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
3	00011_{two}	miss (5.6e)	$(00\mathbf{011}_{\text{two}} \bmod 8) = \mathbf{011}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
18	10010_{two}	miss (5.6f)	$(10\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

d. After handling a miss of address (10000_{two})

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	Y	00_{two}	Memory (00011_{two})
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

e. After handling a miss of address (00011_{two})

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	10_{two}	Memory (10010_{two})
011	Y	00_{two}	Memory (00011_{two})
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

f. After handling a miss of address (10010_{two})