

Computer Architecture

(ENE1004)

Lec - 21: Large and Fast: Exploiting Memory Hierarchy (Chapter 5) - 3

Schedule

- Final exam: Jun. 19, Monday
 - (24334) 1:25~2:25pm
 - (23978) 2:30~3:30pm
 - Sample questions will be provided by Jun. 17 (Saturday)
- Assignment #2: Jun. 20, Tue. by midnight
 - It is put back as much as possible
- Remaining class days
 - 5(Mon), 8(Thur), 12(Mon), 15 (Thur)

Access a Direct Mapped Cache (1)

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

- Each entry of the cache consists of “index” + “valid bit” + “tag” + “data”
 - There are 8 entries (blocks) in the cache, each of which is identified based on the 3-bit “index”
 - The 1-bit “valid bit” tells whether the information in the entry is valid or not
 - If a data is placed in its corresponding entry, its 2-bit “tag” is set to indicate the specific data
- At first, the cache is empty (it does not hold any data) right after power-on
 - The valid bit of each entry is set to “N”
 - The tag and data are also empty

Access a Direct Mapped Cache (2)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (5.6b)	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	miss (5.6c)	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
22	10110_{two}	hit	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	hit	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	miss (5.6d)	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
3	00011_{two}	miss (5.6e)	$(00\mathbf{011}_{\text{two}} \bmod 8) = \mathbf{011}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
18	10010_{two}	miss (5.6f)	$(10\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

b. After handling a miss of address (10110_{two})

- An example scenario where a series of requests (references) are given to the cache
 - If the referenced data is in the cache, it is “hit” and the data is serviced from the cache
 - If the referenced data is not in the cache, it is “miss” and the existing data is replaced with the referenced data, which is copied from the main memory
- Data 22 ($10\mathbf{110}_{\text{two}}$) is requested
 - It can be placed in the cache block whose “index” is $\mathbf{110}_{\text{two}}$ (the lower 3 bits $10\mathbf{110}_{\text{two}}$)
 - It is “miss” because the valid bit is set to N (Data 22 does not exist in the cache block)
 - Data 22 is brought from the main memory and stored in the cache block whose index is 110_{two}
 - Tag is set to 10_{two} to indicate data 22 (the upper 2 bits of $10\mathbf{110}_{\text{two}}$); valid bit is set to Y

Access a Direct Mapped Cache (3)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110 _{two}	miss (5.6b)	(10110 _{two} mod 8) = 110 _{two}
26	11010 _{two}	miss (5.6c)	(11010 _{two} mod 8) = 010 _{two}
22	10110 _{two}	hit	(10110 _{two} mod 8) = 110 _{two}
26	11010 _{two}	hit	(11010 _{two} mod 8) = 010 _{two}
16	10000 _{two}	miss (5.6d)	(10000 _{two} mod 8) = 000 _{two}
3	00011 _{two}	miss (5.6e)	(00011 _{two} mod 8) = 011 _{two}
16	10000 _{two}	hit	(10000 _{two} mod 8) = 000 _{two}
18	10010 _{two}	miss (5.6f)	(10010 _{two} mod 8) = 010 _{two}
16	10000 _{two}	hit	(10000 _{two} mod 8) = 000 _{two}

Index	V	Tag	Data
000	N		
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

c. After handling a miss of address (11010_{two})

- Data 26 (11010_{two}) is requested
 - It can be placed in the cache block whose “index” is 010_{two} (the lower 3 bits 11010_{two})
 - It is “miss” because the valid bit is set to N (Data 26 does not exist in the cache block)
 - Data 26 is brought from the main memory and stored in the cache block whose index is 010_{two}
 - Tag is set to 11_{two} to indicate data 26 (the upper 2 bits of 11010_{two}); valid bit is set to Y
- Data 22 (10110_{two}) is requested again
 - Go to the index 110_{two}; first, check its valid bit Y; then, check the tag 10_{two}; it is “hit”
- Data 26 (11010_{two}) is requested again
 - Go to the index 010_{two}; first, check its valid bit Y; then, check the tag 11_{two}; it is “hit”

Access a Direct Mapped Cache (4)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110 _{two}	miss (5.6b)	$(10\mathbf{110}_{two} \bmod 8) = \mathbf{110}_{two}$
26	11010 _{two}	miss (5.6c)	$(11\mathbf{010}_{two} \bmod 8) = \mathbf{010}_{two}$
22	10110 _{two}	hit	$(10\mathbf{110}_{two} \bmod 8) = \mathbf{110}_{two}$
26	11010 _{two}	hit	$(11\mathbf{010}_{two} \bmod 8) = \mathbf{010}_{two}$
16	10000 _{two}	miss (5.6d)	$(10\mathbf{000}_{two} \bmod 8) = \mathbf{000}_{two}$
3	00011 _{two}	miss (5.6e)	$(00\mathbf{011}_{two} \bmod 8) = \mathbf{011}_{two}$
16	10000 _{two}	hit	$(10\mathbf{000}_{two} \bmod 8) = \mathbf{000}_{two}$
18	10010 _{two}	miss (5.6f)	$(10\mathbf{010}_{two} \bmod 8) = \mathbf{010}_{two}$
16	10000 _{two}	hit	$(10\mathbf{000}_{two} \bmod 8) = \mathbf{000}_{two}$

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

d. After handling a miss of address (10000_{two})

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

e. After handling a miss of address (00011_{two})

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	10 _{two}	Memory (10010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

f. After handling a miss of address (10010_{two})

• 26(11010_{two}) in cache block (index 010) is replaced by 18(10010_{two})

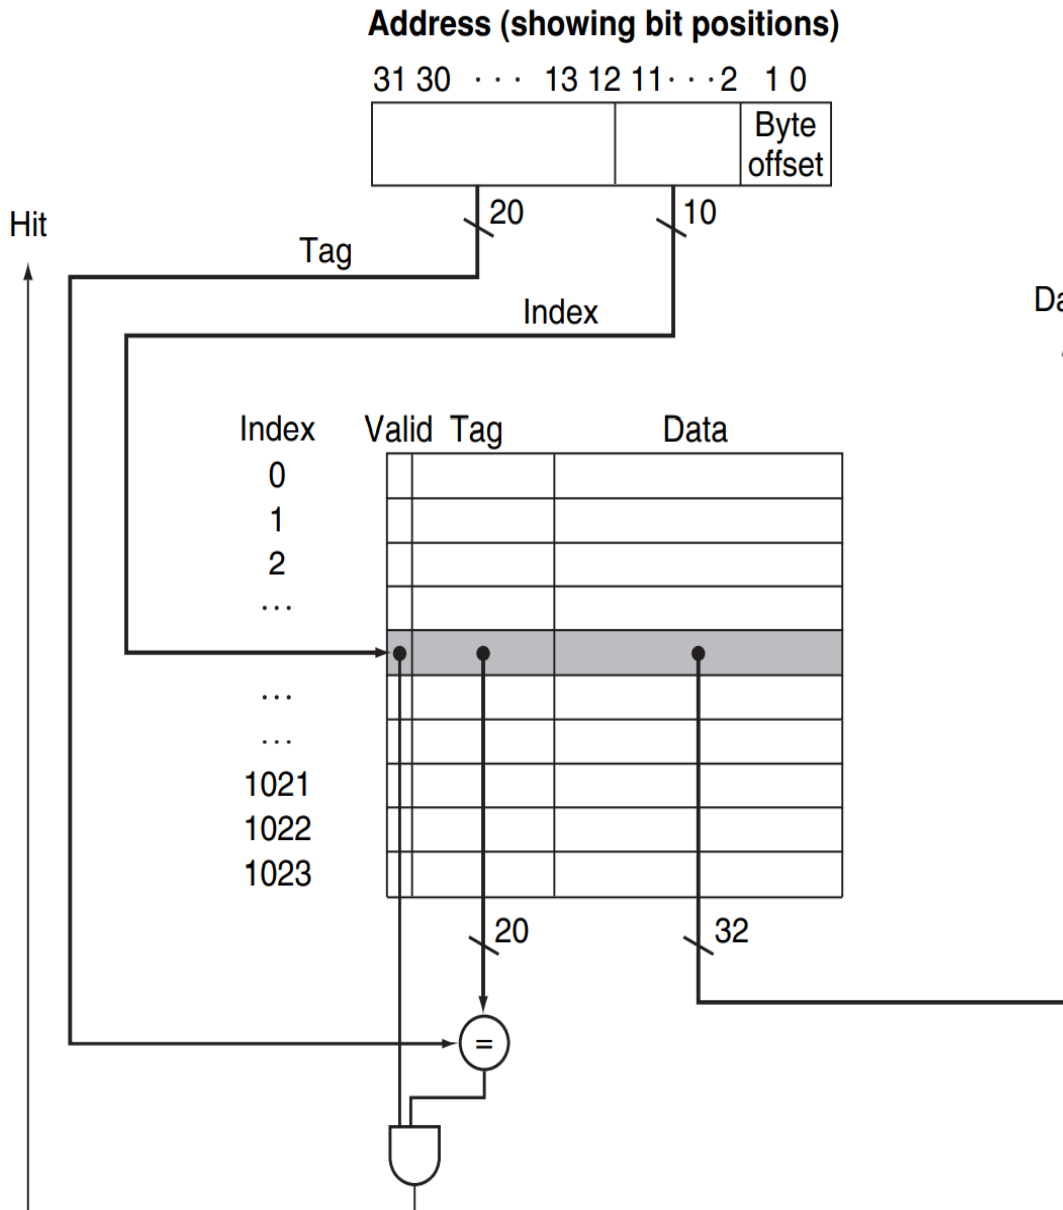
- They are mapped to the same cache block (index 010)
- The recently referenced data item replaces less recently referenced data
- This behavior allows a cache to take advantage of “temporal locality”

Access a Direct Mapped Cache (5)

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (5.6b)	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	miss (5.6c)	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
22	10110_{two}	hit	$(10\mathbf{110}_{\text{two}} \bmod 8) = \mathbf{110}_{\text{two}}$
26	11010_{two}	hit	$(11\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	miss (5.6d)	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
3	00011_{two}	miss (5.6e)	$(00\mathbf{011}_{\text{two}} \bmod 8) = \mathbf{011}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$
18	10010_{two}	miss (5.6f)	$(10\mathbf{010}_{\text{two}} \bmod 8) = \mathbf{010}_{\text{two}}$
16	10000_{two}	hit	$(10\mathbf{000}_{\text{two}} \bmod 8) = \mathbf{000}_{\text{two}}$

- In this example, how many cache accesses are there?
 - 9 cache accesses
- Among them, how many hits and misses are there?
 - 4 hits and 5 misses
- What is the cache hit ratio (miss ratio)?
 - 4/9 (5/9)
- An ideal cache achieves the hit ratio of “1”

Direct Mapped Cache for the Real World (1)



- Assumption

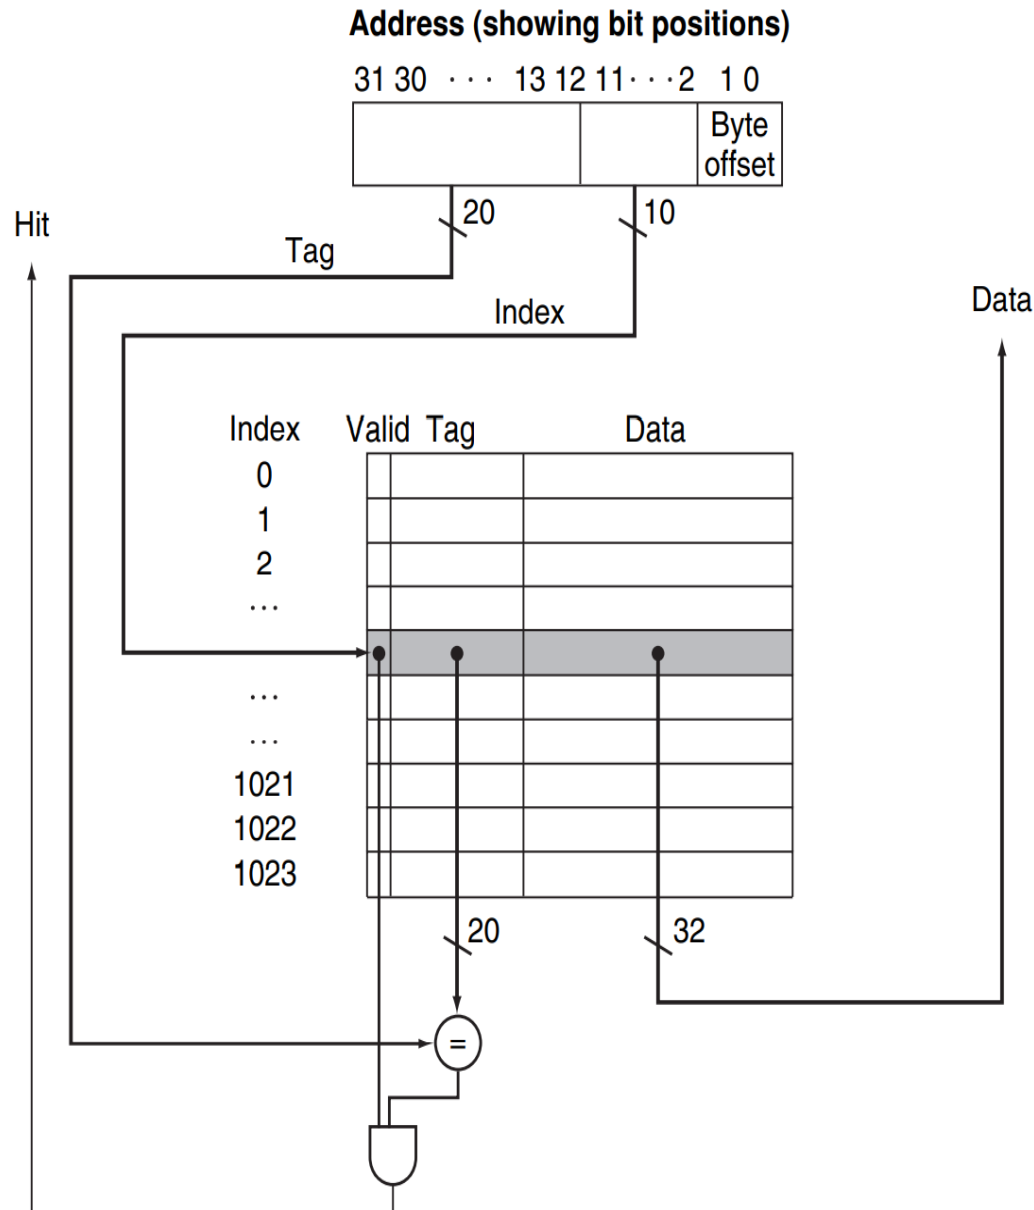
- 32-bit address (vs 5-bit address)
- The cache size is 2^{10} blocks (vs 8 blocks)
- The block size is 2^2 bytes

- Now, let us determine how an address of data is used for a cache access

- Byte offset

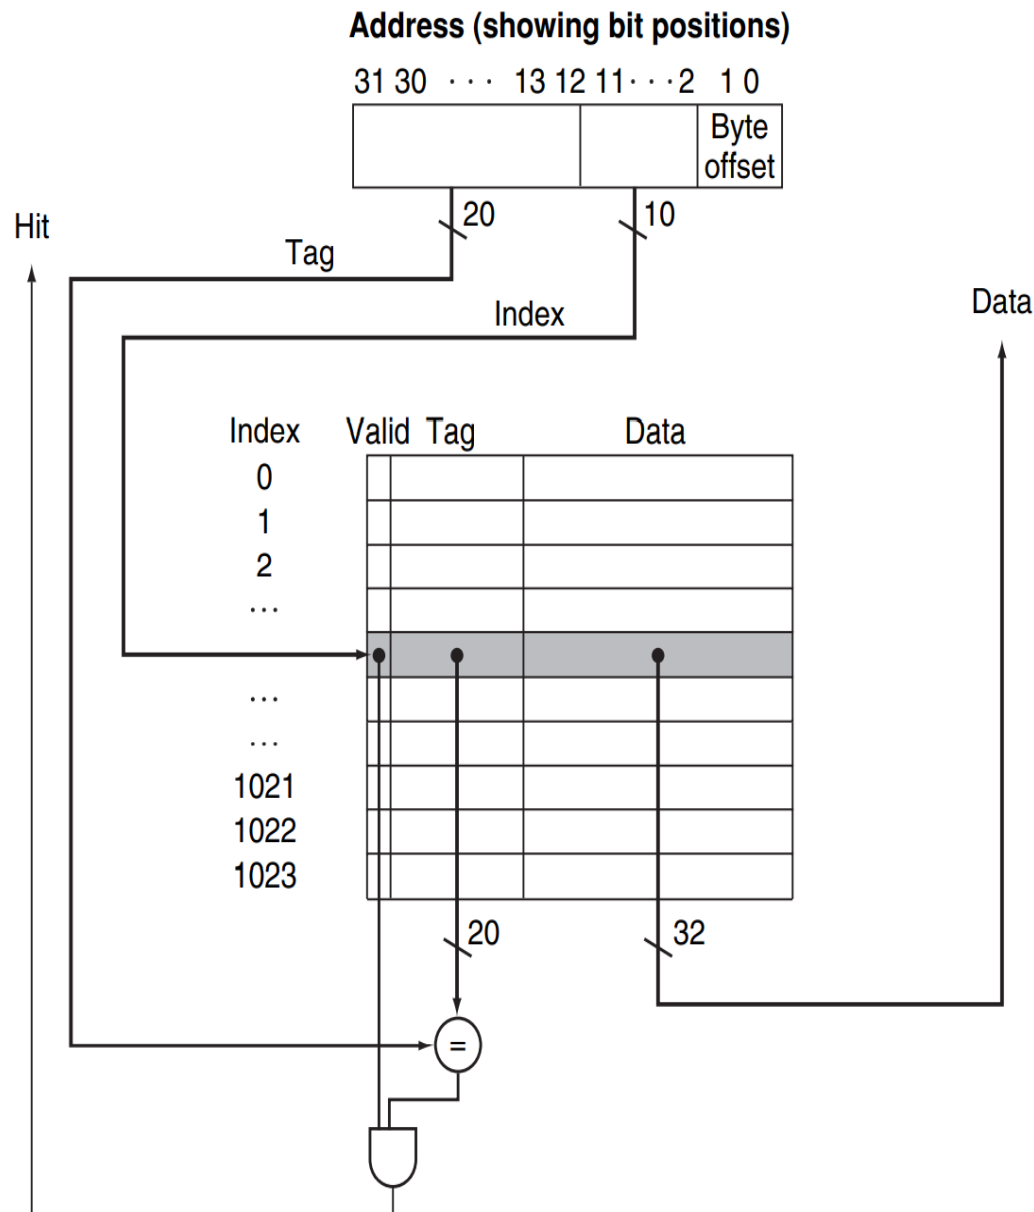
- Each cache block can contain 2^2 bytes of data
- The base unit of a data item is a “byte”
- We need to identify a byte of data within cache block
- The last 2 bits are used for byte offset
- A byte of data is referenced, 2^2 bytes of data including the referenced and neighboring data are brought from the memory to the cache
- This behavior allows a cache to take advantage of “spatial locality”

Direct Mapped Cache for the Real World (2)



- Assumption
 - 32-bit address (vs 5-bit address)
 - The cache size is 2^{10} blocks (vs 8 blocks)
 - The block size is 2^2 bytes
- Index (to select a cache block)
 - There are 2^{10} blocks in the cache
 - We need “10 bits” to determine the index of the data
- Tag (to identify a specific data item)
 - The remaining “20 bits” are used for the tag
 - $20 = 32 - 10$ (for index) - 2 (for byte offset)
- Cache access
 - Given an address, the target cache block is accessed using the 10-bit (b11-b2) index value
 - Then, the 20-bit (b31-b12) tag value is compared to the tag field of the cache block
 - If they are equal (“hit”), the entire 4-byte data is read
 - If they are not equal (“miss”), the referenced 4-byte data replaces the existing 4-byte data in the cache

Generalization of Direct Mapped Cache (1)



- Assumption

- 32-bit address
- The cache size is 2^n blocks
- The block size is 2^m words (2^{m+2} bytes)

- Byte offset

- Each cache block can contain 2^{m+2} bytes of data
- “m bits” are used to identify a word within a block
- “2 bits” are used to identify a byte within a word
- The last $m+2$ bits are used for byte offset


- Index (to select a cache block)

- There are 2^n blocks in the cache
- We need “n bits” to determine the index of the data

- Tag (to identify a specific data item)

- The remaining bits are used for the tag
- The size of the tag field is $32 - (n + m + 2)$

Generalization of Direct Mapped Cache (2)

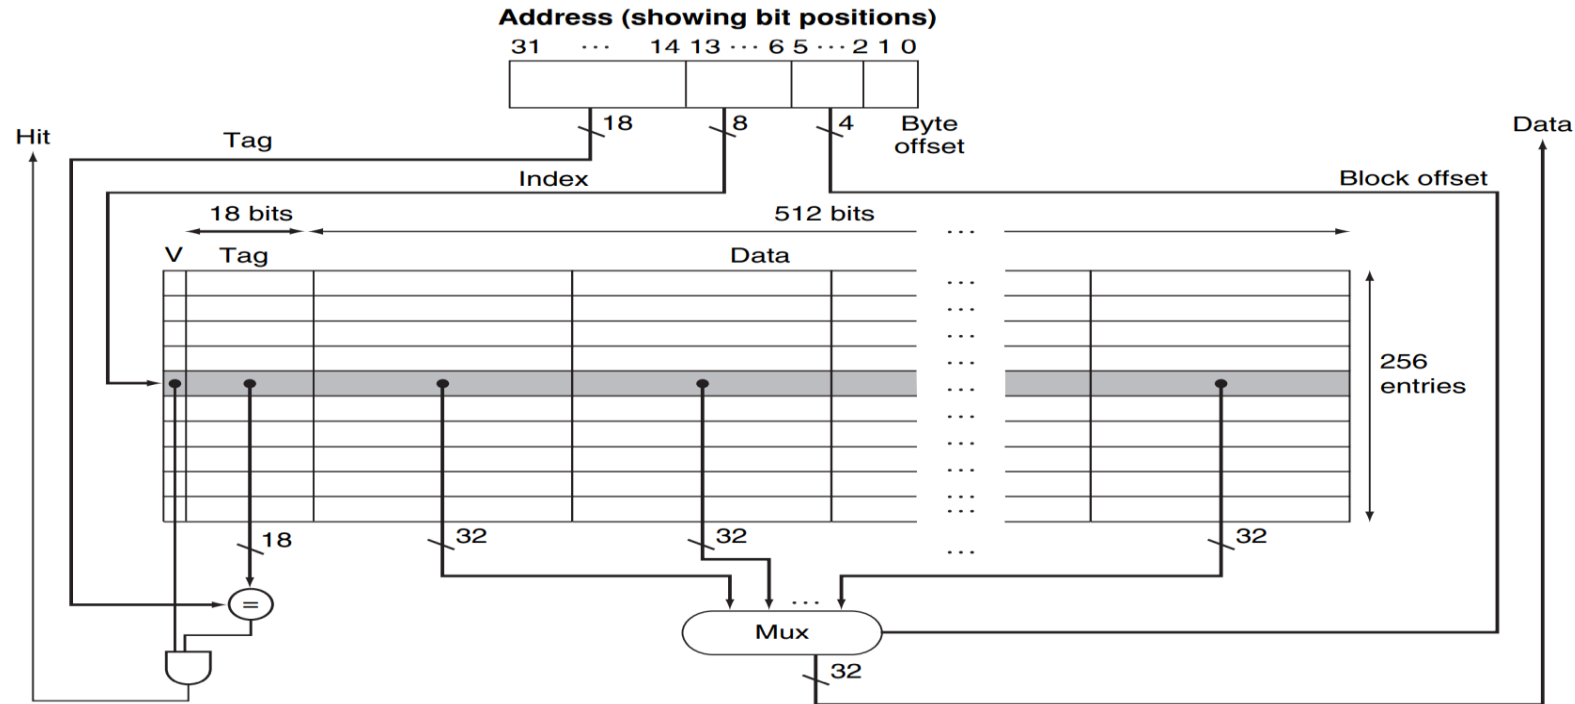


- Assumption
 - 32-bit address
 - The cache size is 2^n blocks
 - The block size is 2^m words (2^{m+2} bytes)
- Address fields
 - Byte offset: $m+2$ bits
 - Index: n bits
 - Tag: $32 - (n + m + 2)$ bits
- The total number of bits in a cache
 - # of blocks x (block size + tag size + valid bit size)
 - $2^n \times (2^{m+2} \text{ bytes} + (32-n-m-2) \text{ bits} + 1 \text{ bit})$
 - $2^n \times (2^{m+5} + (32-n-m-2) + 1) \text{ bits}$
- Due to the tag and valid bit fields, the actual size of the cache is larger than the total amount of data ($2^n \times 2^{m+2}$ bytes)

Bits in Cache

- Question: How many total bits are required for a direct mapped cache with 16 KB of data and 4-word blocks, assuming a 32-bit address?
- Hint: We need to calculate “total # of blocks (indices)” and “# bits for tag”
 - Cache size = total # blocks x (block size + tag size + valid bit size) for each block
- The total number of cache blocks
 - The entire size of blocks = 16 KB (2^{14} Bytes = 2^{12} words)
 - The size of each block = 4 words (2^4 Bytes = 2^2 words)
 - The total number of blocks = entire size / a block size = $2^{10} = 1024$
 - The number of bits for “index” is 10 bits
- The number of bits for “tag”
 - # bits for tag = total # address bits - (# bits for index + # bits for byte offset)
 - # bits for byte offset = 4, because the size of a cache block is of 4-word (2^4 Bytes)
 - # bits for tag = $32 - (10 + 4) = 18$
- Cache size = total # blocks x (block size + tag bit size + valid bit size) for each block
 - Cache size = $1024 \times (4 \text{ words} + 18 \text{ bits} + 1 \text{ bit}) = 1024 \times (2^7 \text{ bits} + 18 \text{ bits} + 1 \text{ bit}) = 2^{10} \times 147 \text{ bits}$
 - 147 Kbits

Bits in Cache (2)



- Cache block size: 512 bits = 64 (2^6) Bytes = 16 words
 - Byte offset = 6 bits (b5-b0)
 - You may want to access the data in the unit of “word”; then, you can use upper bits in the byte offset bits; in this example, upper 4 bits (b5-b2) is used for word offset (called “block offset”)
- Total number of cache blocks is 256 (2^8); # bits for index is 8 (b13-b6)
- # bits for tag = $32 - (8 + 6) = 18$ bits
- The cache size = $256 \times (512 \text{ bits for data} + 18 \text{ bits for tag} + 1 \text{ bit for valid bit})$